

Multicore Communications API Specification V2.000

Document ID: MCAPI API Specification
Document Version: 2.000
Status: Release
Distribution: General

Name	Role
MCAPI WG Chair, Sven Brehmer	State of doc assessment
MCAPI WG members	Content, coverage and grammar
President of the Multicore Association, Markus Levy	Posting and Publication

Page Intentionally Left Blank

Table of Contents

1	History	6
2	Introduction	8
3	MCAPI Overview	9
3.1	MCAPI Domains	9
3.2	MCAPI Nodes	9
3.3	MCAPI Endpoints	9
3.4	MCAPI Channels	10
3.4.1	MCAPI Communications	10
3.5	Data Delivery	11
3.5.1	Data Sending	11
3.5.2	Data Routing	12
3.5.3	Data Consumption	12
3.5.4	Waiting for non-blocking operations	12
3.6	MCAPI Messages	13
3.7	MCAPI Packet Channels	14
3.8	MCAPI Scalar Channels	14
3.9	MCAPI Zero Copy	15
3.10	MCAPI Error Handling Philosophy	15
3.11	MCAPI Buffer Management Philosophy	16
3.12	Timeout/Cancellation Philosophy	17
3.13	Backpressure mechanism	17
3.14	MCAPI Implementation Concerns	17
3.14.1	Link Management	17
3.14.2	Thread Safe Implementations	17
3.15	MCAPI Potential Future Extensions	17
3.15.1	Zero Copy	18
3.15.2	Multicast	18
3.15.3	Debug, Statistics and Status functions	18
3.16	MCAPI Data Types	19
3.16.1	mcapi_domain_t	19
3.16.2	mcapi_node_t	19
3.16.3	mcapi_port_t	19
3.16.4	mcapi_endpoint_t	19
3.16.5	mcapi_pktchan_rcv_hndl_t	19
3.16.6	mcapi_pktchan_send_hndl_t	20
3.16.7	mcapi_scnchan_rcv_hndl_t	20
3.16.8	mcapi_scnchan_send_hndl_t	20
3.16.9	mcapi_uint64_t, mcapi_uint32_t, mcapi_uint16_t & mcapi_uint8_t,	20
3.16.10	mcapi_request_t	20
3.16.11	mcapi_status_t	21
3.16.12	mcapi_timeout_t	21
3.16.13	MCAPI endpoint attributes	21
3.16.14	Other MCAPI data types	23
3.16.15	MCAPI Error and Status Codes	23
3.16.16	API parameter readability	24
3.17	What is new in MCAPI 2.000	25
3.17.1	MCAPI to MCA types	25
3.17.2	MCAPI Domains	25
3.17.3	MCAPI endpoint attributes	25
3.17.4	MCAPI header files	25
3.17.5	New functionality and functions	25
3.17.6	API Consistency	27
3.17.7	Definitions and constants	27
3.17.8	Clarifications	27
3.18	MCAPI 2.000 to 1.0 backwards compatibility	28
3.18.1	Affected functions	28

4	MCAPI API	29
4.1	General	29
4.1.1	MCAPI_INITIALIZE	30
4.1.2	MCAPI_FINALIZE	32
4.1.3	MCAPI_DOMAIN_ID_GET	33
4.1.4	MCAPI_NODE_ID_GET	34
4.1.5	MCAPI_NODE_INIT_ATTRIBUTES	35
4.1.6	MCAPI_NODE_SET_ATTRIBUTE	36
4.1.7	MCAPI_NODE_GET_ATTRIBUTE	38
4.2	Endpoints	40
4.2.1	MCAPI_ENDPOINT_CREATE	41
4.2.2	MCAPI_ENDPOINT_GET_I	43
4.2.3	MCAPI_ENDPOINT_GET	45
4.2.4	MCAPI_ENDPOINT_DELETE	47
4.2.5	MCAPI_ENDPOINT_GET_ATTRIBUTE	48
4.2.6	MCAPI_ENDPOINT_SET_ATTRIBUTE	51
4.3	MCAPI Messages	54
4.3.1	MCAPI_MSG_SEND_I	55
4.3.2	MCAPI_MSG_SEND	57
4.3.3	MCAPI_MSG_RECV_I	59
4.3.4	MCAPI_MSG_RECV	61
4.3.5	MCAPI_MSG_AVAILABLE	63
4.4	MCAPI Packet Channels	64
4.4.1	MCAPI_PKTCHAN_CONNECT_I	65
4.4.2	MCAPI_PKTCHAN_RECV_OPEN_I	67
4.4.3	MCAPI_PKTCHAN_SEND_OPEN_I	69
4.4.4	MCAPI_PKTCHAN_SEND_I	71
4.4.5	MCAPI_PKTCHAN_SEND	73
4.4.6	MCAPI_PKTCHAN_RECV_I	75
4.4.7	MCAPI_PKTCHAN_RECV	78
4.4.8	MCAPI_PKTCHAN_AVAILABLE	80
4.4.9	MCAPI_PKTCHAN_RELEASE	81
4.4.10	MCAPI_PKTCHAN_RELEASE_TEST	82
4.4.11	MCAPI_PKTCHAN_RECV_CLOSE_I	83
4.4.12	MCAPI_PKTCHAN_SEND_CLOSE_I	85
4.5	MCAPI Scalar Channels	87
4.5.1	MCAPI_SCLCHAN_CONNECT_I	88
4.5.2	MCAPI_SCLCHAN_RECV_OPEN_I	90
4.5.3	MCAPI_SCLCHAN_SEND_OPEN_I	92
4.5.4	MCAPI_SCLCHAN_SEND_UINT64	94
4.5.5	MCAPI_SCLCHAN_SEND_UINT32	95
4.5.6	MCAPI_SCLCHAN_SEND_UINT16	96
4.5.7	MCAPI_SCLCHAN_SEND_UINT8	97
4.5.8	MCAPI_SCLCHAN_RECV_UINT64	98
4.5.9	MCAPI_SCLCHAN_RECV_UINT32	99
4.5.10	MCAPI_SCLCHAN_RECV_UINT16	100
4.5.11	MCAPI_SCLCHAN_RECV_UINT8	101
4.5.12	MCAPI_SCLCHAN_AVAILABLE	102
4.5.13	MCAPI_SCLCHAN_RECV_CLOSE_I	103
4.5.14	MCAPI_SCLCHAN_SEND_CLOSE_I	105
4.6	MCAPI non-blocking operations	107
4.6.1	MCAPI_TEST	108
4.6.2	MCAPI_WAIT	110
4.6.3	MCAPI_WAIT_ANY	112
4.6.4	MCAPI_CANCEL	114
4.7	MCAPI support functions	115
4.7.1	MCAPI_DISPLAY_STATUS	116
5	Header files	117
6	FAQ	118

7	Use Cases.....	121
7.1	Example Usage of Static Naming for Initialization	121
7.2	Example Initialization (Discovery and Bootstrapping) of Dynamic Endpoints.....	121
7.3	Automotive Use Case.....	122
7.3.1	Characteristics.....	122
7.3.2	Key functionality requirements.....	122
7.3.3	Context/Constraints	123
7.3.4	Metrics	123
7.3.5	Possible Factorings	123
7.3.6	MCAPI Requirements Implications.....	123
7.3.7	Mental Models	124
7.3.8	Applying the API draft to the pseudo-code	125
7.4	Multimedia Processing Use Cases	131
7.4.1	Characteristics.....	131
7.4.2	Key Functionality Requirements.....	136
7.4.3	Pseudo-Code Example:	137
7.5	Packet Processing	143
7.5.1	Packet Processing Code	144

1 History

Although multiprocessor designs have been around for decades, only recently have semiconductor vendors turned to multicore processors as a solution for the so-called Moore's Gap¹ effect. Simply put, modern multicore CPU design enables CPU performance to track Moore's law with attractive CPU power consumption properties. This trend is causing many system designers that previously would use single core designs to design multicore processors into hardware in order to meet their performance and/or power consumption requirements. This has a huge impact on the software architecture for those systems, which must now consider inter-processor communication (IPC) issues to pass data between the cores.

The Multicore Association's Communication API (MCAPI) traces its heritage to communications APIs such as MPI (Message Passing Interface) and sockets. Both MPI and sockets were developed primarily with inter-computer communication in mind, while MCAPI is targeted primarily towards inter-core communication in a multicore chip. Accordingly, a principal design goal of MCAPI was to serve as a low-latency interface leveraging efficient on-chip interconnect in a multicore chip. MCAPI is thus a light-weight API whose communications latencies and memory footprint are expected to be significantly lower than that of MPI or sockets. However, because of the more limited scope of multicore communications and its goal of low latency, MCAPI is less flexible than MPI or sockets.

MCAPI can also be distinguished from POSIX (portable Unix) threads, pthreads. MCAPI is a communications API for messaging and streaming, while pthreads is an API for parallelizing or threading applications. Pthreads offers a programming model that relies on shared memory for communication and locks for synchronization. In multicore processors with caches, efficient pthreads implementations require support for cache coherence. Because they lack modularity, lock-based parallel programs are hard to compose together. MCAPI, with parallelism specified using standard means such as processes or threads, offers an alternative and complementary parallel programming approach that is modular and composable. Because it works both in multicore processors with only private memory or with shared memory, MCAPI allows simpler embedded multicore implementations.

IPC for device software utilizing multicore processors have some new requirements beyond those imposed on multiprocessor designs that utilized discrete processors. Within a multicore chip, the shorter distance for electrical signals enables data to be passed much more quickly, energy-efficiently, and reliably than between discrete processors on a board or over a backplane. In addition, the bandwidth available on chip is orders of magnitude greater. The low latency and high bandwidth offer the potential of ASIC-like (application specific integrated circuit) high-speed communication capabilities between cores. This means that the focus of IPC for multicore processors must prioritize performance in terms of both latency and throughput. Also, some multicore processors will have many cores that rely on chip-internal memory or cache to execute code (to avoid the von Neumann bottleneck which multicore processors can suffer due to the multiple cores contending for external memory accesses). Such cores that execute from chip-internal memory will require an IPC implementation that has a small memory footprint. For these reasons, the two primary goals for a multicore IPC API design are that the implementation of it can achieve extremely high performance and low memory footprint.

In order for a multicore IPC implementation to achieve high performance and tiny footprint, these two goals need to be immutable when in conflict with other desirable IPC API attributes such as physical or logical connectivity topology flexibility, ease of use, OS integration or compatibility with existing IPC programming models or APIs.

The Multicore Association's Communication API (MCAPI) endeavors to meet these goals. Before MCAPI, no IPC API has been designed with these two immutable goals targeting multicore processors.

¹ Moore's Law has hit a wall where increased transistors per sequential CPU chip is no longer resulting in a linear increase in performance of that CPU chip.

The MCAPI API was designed to be sufficiently complete for many multicore application programming tasks and also to serve as a solid foundation for sophisticated multicore software services. Therefore MCAPI can be used to implement distributed services such as name services and distributed work queues, as well as one sided services such as DMA.

2 Introduction

This document is intended to assist software developers who are either implementing MCAPI or writing applications that use MCAPI.

The MCAPI specification is both an API and communications semantic specification. It does not define which link management, device model or wire protocol is used underneath it. As such, by defining a standard API, it is intended to provide source code compatibility for application code to be ported from one operating environment to another.

MCAPI defines three fundamental communications types. These are:

1. Messages – connection-less datagrams.
2. Packet channels – connection-oriented, uni-directional, FIFO packet streams.
3. Scalar channels – connection-oriented single word uni-directional, FIFO scalar streams.

Each of these communications types have their own API calls and are desirable for certain types of systems. Messages are the most flexible form of communication in MCAPI, and are useful when senders and receivers and per message priorities are dynamically changing. These are commonly used for synchronization, initialization and load balancing.

Packet and scalar channels provide light-weight socket-like stream communication mechanisms for senders and receivers with static communication graphs. In a multicore, MCAPI's channel APIs provide an extremely low-overhead ASIC-like uni-directional FIFO communications capability. Channels are commonly set up once during initialization, during which the MCAPI runtime system attempts to perform as much of the work involved in communications (such as name lookup, route determination, and buffer allocation) between a specific pair of endpoints as possible. Subsequent channel sends and receives thus incur just the minimal overhead of physically transferring the data. Packet channels support streaming communication of multiword data buffers, while scalar channels are optimized for sequences of scalar values. Channel API calls are simple and statically typed, thereby minimizing dynamic software overhead, which allows applications to access the underlying multicore hardware with extremely low latency and energy cost.

MCAPI's objective is to provide a limited number of calls with sufficient communication functionality while keeping it simple enough to allow efficient implementations. Additional functionality can be layered on top of the API set. The calls are exemplifying functionality and are not mapped to any particular existing implementation.

3 MCAPI Overview

The major MCAPI concepts are covered in the following sections.

Editors note: Request that technical writer includes MCAPI header file information such as attributes and other function relevant information directly in the function sections.

Also section references have to be validated]

3.1 MCAPI Domains

An MCAPI domain is comprised of one or more MCAPI nodes in a multicore topology and used for routing purposes. The scope of a domain is implementation defined and its scope could for example be a single chip with multiple cores or multiple processor chips on a board. The domain id is specified once at node initialization. A domain can contain multiple nodes. Some example potential uses for domains are topologies that may change dynamically, include non-MCAPI sub-topologies, require separation between different transports or have open and secure areas.

3.2 MCAPI Nodes

An MCAPI node is a logical abstraction that can be mapped to many entities, including but not limited to: a process, a thread, an instance of an operating system, a hardware accelerator, or a processor core. The node id is specified at the call to `mcapi_initialize()`. A given MCAPI implementation will specify what is a node (i.e., what thread of control – process, thread, or other -- is a node) for that implementation. A thread and process are just two examples of threads of control, and there could be others.

The MCAPI standard aims to be implementable on both process-oriented and thread-oriented systems. MCAPI defines a transport layer between 'nodes', which could be processes, threads, or some other thread-of-control abstraction. The standard explicitly avoids having any dependence on a particular shared memory model or process protection model. The communication operations defined by MCAPI should be implementable with identical semantics on either thread- or process- based systems. Thus, programs that use only the MCAPI APIs should be portable between MCAPI implementations. On the other hand, programs that make use of APIs outside of MCAPI, for example pthreads, will only be portable between systems that include those extra APIs.

THINGS TO REMEMBER:

- The MCAPI domain and node numbering plan is established when the MCAPI communication topology is configured at design-time, so programmers don't have to worry about this at run-time.
- It is up to the MCAPI implementation to configure communication topology interfaces to enable communication with other nodes in the communication topology.

3.3 MCAPI Endpoints

MCAPI endpoints are socket-like communication termination points. MCAPI endpoints are created with the `mcapi_endpoint_create()` function. An MCAPI node can have multiple endpoints. Thus, an MCAPI endpoint is a topology-global unique identifier. Endpoints are identified by a `<domain_id, node_id, port_id>` tuple. An endpoint is created by implicitly referring to the `node_id` of the node on which the endpoint is being created, and explicitly specifying a `port_id`. Alternatively, MCAPI allows the creation of an endpoint, by requesting the next available endpoint. . Static names allow the programmer to

Comment [ML1]: Do you want to include a Definitions section similar to what is found in MRAPI Overview section 3.1?

Comment [S2]: yes, makes sense.

Comment [S3]: Some of the functions have attributes that are listed in the header files. They should also be included in the function description, similar to what we have in the MRAPI spec.

Comment [S4]: For tech editor.

define a network topology at compile time and to embed the topology in the source code in a straightforward fashion. It also enables the use of external tools to generate topologies automatically and it facilitates simple initialization. Endpoints can also be created by requesting the next available endpoint which allows for topology flexibility.

It is up to the implementation to ensure that the resulting endpoint reference is unique in the system. MCAPAPI allows a reference to endpoints to be obtained by a node other than the node whose `node_id` is associated with the endpoint, using the `mcapi_endpoint_get()` function, and this endpoint reference can be used as a destination for messages or to specify the opposite end of a connection. Endpoint references created by requesting the next available endpoint must be passed by the creating node to other nodes to facilitate communication. Only the creating node is allowed to receive from an endpoint.

There are three types of endpoints, message, packet channel and scalar channel. The endpoint type defines certain aspects of the endpoint's behavior. The type can for example be used to manage avoidance of messages being sent to connected endpoints. The type is set with `mcapi_endpoint_set_attribute()`. An endpoint is always created as a message type, which is the default type. Message type endpoints can have multiple outstanding endpoint get endpoint requests from other nodes, whereas channel type endpoints can have only one outstanding get.

Endpoints have a set of attributes. These attributes may be related to Quality of Service (QoS), buffers, timeouts, etc. Currently MCAPAPI defines a basic set of attribute numbers and their associated structure or data type, found in the `mcapi.h` file. Implementations may request a range of vendor specific attribute numbers from the Multicore Association to add additional attributes (numbers and data types). MCAPAPI does define API functions to get and set the attributes of endpoints. Attributes can only be set on node local endpoints but be gotten from local and remote endpoints. Furthermore, it is an error to attempt a connection between endpoints whose attributes are set in an incompatible way (for now, whether attributes are compatible or not is implementation defined). It is also an error to attempt to change the attributes of endpoints that are connected.

3.4 MCAPAPI Channels

Channels provide point-to-point FIFO connections between a pair of endpoints. MCAPAPI channels are unidirectional. There are two types of channels: packet channels and scalar channels. There is no separate channel object. Instead, channels can be referenced through the endpoints to which they are bound. Since channels are point-to-point connections between a pair of endpoints, either endpoint can be used to refer to the channel.

3.4.1 MCAPAPI Communications

Applications in an MCAPAPI topology communicate with one another by sending data between communication endpoints. Three modes of communication are supported: connectionless messages, connection-oriented packet channels, or connection-oriented scalar channels. Each communication action in all three methods (for example, a message send, a packet channel send, or a scalar channel send) requires the specification of a pair of endpoints – a send endpoint and a receive endpoint, explicitly for messages and implicitly for channels using a handle. Messages and packet channels provide both blocking and non-blocking send and receive functions. The MCAPAPI runtime system implements each of these logical communication actions over some physical medium that is supported by MCAPAPI (such as on-chip interconnect, bus, shared memory, or Ethernet) by automatically establishing a “link” to enable communication with the other node in the network, and takes care of routing traffic over the appropriate link.

Messages and packet channels communicate using data buffers, while scalar channels transfer scalar values such as four-byte words.

From an application's perspective, a data buffer to be communicated is a byte string from 0 up to some maximum length determined by the amount of memory in the system, whose internal structure is

determined by the application. This specification does not define a communication byte order because MCAPi is an API specification. Since MCAPi does not currently define a wire protocol (there may be a future Multicore Association specification which does specify that) then endian-ness and alignment issues etc are not relevant to this specification.

Connectionless communication allows an endpoint to send or receive messages with one or more endpoints elsewhere in the communication topology. A given message can be sent to a single endpoint (unicast). Multicast and broadcast messaging modes are not supported by MCAPi, but can be built on top of it.

Connection-oriented communication allows an endpoint to establish a socket-like streaming connection to a peer endpoint elsewhere in the communication topology, and then send data to that peer. A connection is established using an explicit handshake mechanism prior to sending or receiving any application data. Once a connection has been established, it remains active for highly efficient data transfers until it is terminated by one of the sides, or until the communication path between the endpoints is severed (for example, by the failure of the node or link which one of the endpoints is utilizing).

Connection-oriented communication over MCAPi channels is designed to be reliable, in that an application can send data over a connection and assume that the data will be delivered to the specified destination as long as that destination is reachable.

THINGS TO REMEMBER:

- In an MCAPi communication topology where different nodes may be running on different CPU types and/or operating systems, applications must ensure that the internal structure of a message is well-defined, and account for any differences in message content endian-ness, field size and field alignment.

3.5 Data Delivery

On the surface, whether the communications use messages, packet channels or scalar channels, the data delivery in MCAPi is a simple series of steps: a sender creates and sends the data (either a buffer for messages and packet channels, or a scalar value for scalar channels), the MCAPi implementation carries the data to the specified destination, and the receiver receives and then consumes the data. In practice, this is exactly what happens most of the time. However, there are a number of places along the way where things can get complicated, and in these cases it is important for application designers to understand exactly what MCAPi will do.

The sections that follow describe the various steps performed by MCAPi during the transfer of a data packet.

3.5.1 Data Sending

The first step in sending data is to create it. The user application then sends the data using one of several mechanisms. The application supplies the data in a user buffer for messages and packet channels. The data is supplied directly as a scalar variable for scalar channels.

The most common reason MCAPi is unable to send a piece of data is because the sender passes in one or more invalid arguments to the send routine. The term "invalid" refers both to values that are never acceptable under any circumstances (such as specifying a data buffer size exceeding the maximum size allowed for a specific implementation) and to values that are not acceptable for the current sender (such as an invalid endpoint handle, or an unconnected channel).

In all of these cases the send operation will return a failure code indicating that the intended data was not sent. If the data is sent successfully, the send operation returns a success indication. Success means that the entire buffer has been sent.

THINGS TO REMEMBER:

- If the sender specifies a destination address for a connectionless message, that does not currently exist within the MCAPI communication topology, MCAPI does **NOT** treat this as an invalid send request (i.e. it's not the sender's fault that the destination doesn't exist). Implementations may however treat this as an error condition, and if so has to document the specifics. Instead MCAPI creates the message and then sends it. The return value for the send operation will indicate success since the message was successfully created and processed by MCAPI. MCAPI cannot return an error since the connectivity failure could be occurring elsewhere in the MCAPI communication topology.

3.5.2 Data Routing

Once a send of some data has been requested, MCAPI then determines what node the data should be sent to. The endpoint address indicates the destination node to which the packet should be sent. The data is then either handed off to the destination endpoint directly if it is on the same node as the sender, or passed to a link for off-node transmission to the destination node.

One problem that can arise during the routing phase of packet delivery is that no working link to the specified destination node can be found. In this case, the packet is discarded. Implementations may optionally report this is an error condition.

3.5.3 Data Consumption

When the data reaches the destination endpoint, it is added to an endpoint receive queue. The data typically remains in the endpoint's receive queue until it is received by the application that owns the endpoint. Queued data items are consumed by the application in a FIFO manner for channels, and FIFO order per priority level for messages. For messages, the user application must supply a buffer into which the MCAPI runtime system fills in the data. For packet channels, on the other hand, the MCAPI runtime system supplies the buffer containing the data to the user. For messages, the application can use its data buffer after it is filled in by MCAPI in any way it chooses. Specifically, the application can re-supply the data buffer to MCAPI to receive the next message. Data buffers supplied by MCAPI during packet channel receives can also be used by the application in any way it chooses. MCAPI reuses a data buffer it has supplied only after the user application has specifically freed the buffer using the `mcapi_pktchan_release()` function.

If an application terminates access to the before all data in the receive queue is consumed, all unconsumed data items are considered undeliverable and are discarded.

It is very important that MCAPI applications be engineered to consume their incoming data items at a rate that prevents them from accumulating in large numbers in any endpoint receive queue.

For the packet and scalar channels interfaces, a FIFO model is used. FIFOs have limited storage, and when the storage is used up, an implementation can choose to block until more storage is available or return an error code.

Non-blocking send operations will still return in a timely fashion, but data transfer will not occur if the FIFO is full. The sender will block upon calling service routines to check for completion of a non-blocking operation, and will continue to block until a receive operation is performed or a timeout occurs.

The user can use the `mcapi_endpoint_get/set_attribute()` APIs to query or to control the amount of receive buffering that is provided for a given endpoint. Implementations may define a static receive buffer size according to the amount of buffering provided in drivers and/or hardware. On the other hand, an implementation could be allowed to dynamically allocate storage until the system runs out of memory.

3.5.4 Waiting for non-blocking operations

The connectionless message and packet channel API functions have both blocking and non-blocking variants. The non-blocking variants of these MCAPI functions have “_i” appended to the function name to indicate that the function will return *immediately* and will complete in a non-blocking manner.

The non-blocking versions fill in a `mcapi_request_t` object and return control to the user before the communication operation is completed. The user can then use the `mcapi_test()`, `mcapi_wait()`, and `mcapi_wait_any()` functions to query the status of the non-blocking operation. These functions are non-destructive, meaning that no message, packet or scalar is removed from the endpoint queue by the functions. The `mcapi_test()` function is non-blocking whereas the `mcapi_wait()` and `mcapi_wait_any()` functions will block until the requested operation completes or a timeout occurs. Multiple threads should not be waiting the same request and attempting to do so will result in an error.

If a buffer of data is passed to a non-blocking operation (for example, to `mcapi_msg_send_i()`, `mcapi_msg_recv_i()`, or to `mcapi_pktchan_send_i()`), that buffer should not be accessed by the user application for the duration of the non-blocking operation. That is, once a buffer has been passed to a non-blocking operation, the program may not read or write the buffer until `mcapi_test()`, `mcapi_wait()`, or `mcapi_wait_any()` have indicated completion, or until `mcapi_cancel()` has canceled the operation.

The MCAPI scalar channels API provides only blocking send and receive methods. Scalar channels are intended to provide a very low overhead interface for moving a stream of values. Non-blocking operations add overhead. The sort of streaming algorithms that take advantage of scalar channels should not require a non-blocking send or receive method; each process should simply receive a value to work on, do its work, send the result out on a channel, and repeat. Applications that require non-blocking semantics should use packet channels instead of scalar channels.

3.6 MCAPI Messages

MCAPI messages provide a flexible method to transmit data between endpoints without first establishing a connection. The buffers on both sender and receiver sides must be provided by the user application. MCAPI messages may be sent with different priorities, on a per message basis. It is not allowed to send a message to a connected endpoint. Implementations may choose to prevent messages from being sent to connected endpoint or to leave it up to the application to manage this. Functionality for this may be added in a future version of MCAPI in it is therefore recommended that implementations preventing messages from being sent to connected endpoint use `MCAPI_ERR_GENERAL` to report an error. The behavior should be documented.




Figure 1 - Connectionless messages

3.7 MCAPI Packet Channels

MCAPI packet channels provide a method to transmit data between endpoints by first establishing a connection, thus potentially removing or reducing the message header and route discovery overhead. Packet channels are unidirectional and deliver data in a FIFO (first in first out) manner. The buffers are provided by the MCAPI implementation on the receive side, and by the user application on the send side. Packet channels provide per endpoint priority. It is not allowed to send a message to a connected endpoint.

3.8 MCAPI Scalar Channels

MCAPI scalar channels provide a method to transmit scalars very efficiently between endpoints by first establishing a connection. Like packet channels, scalar channels are unidirectional and deliver data in a FIFO (first in first out) manner. The scalar functions come in 8-bit, 16-bit, 32-bit and 64-bit variants. The scalar receives must be of the same size as the scalar sends. A mismatch in size results in an error. Scalar channels provide per endpoint priority.




Figure 2 - Connected channels

3.9 MCAPID Zero Copy

Zero copy means that data is passed by reference instead of a physical copy. This method is useful when multiple cores have shared memory, as one core can operate on a buffer of data and simply pass a pointer to the buffer to the next core, and save the time and energy of moving the data physically. Zero copy requires specific buffer management and is intended to be orthogonal to other API operations, to avoid complexity (simplifying code migration) and performance penalties for non-zero copy operations. With the addition of the `mcapi_pktchan_release_test()` function (Version 1.1) MCAPID is capable of supporting Zero Copy, directly through the API functions, provided the availability of shared memory and support in the underlying implementation. The application is responsible for allocation of shared memory.

3.10 MCAPID Error Handling Philosophy

Error handling is a fundamental part of the specification, however, some accommodations have been made to allow for trading off completeness for efficiency of implementation. For example, a few API functions allow implementations to optionally handle errors². Consistent and efficient coding styles also govern the design of the error handling. In general, function calls include an error code parameter used by the API function to indicate detailed status. In addition, the return value of several API functions indicates success or failure which enables efficient coding practice. A value of type `mcapi_status_t`, indicates success or failure states of API calls. `MCAPID_NULL` is a valid return value for `mcapi_status_t` (can be used for implementation optimization) and implementations should state when this is the case.

² The `mcapi_selchan_unit_send_*` and `mcapi_selchan_uint_recv_*` routines define error arguments, but implementation are free to assume the routines always succeed. The message on a packet channel send and receive functions can optionally report transmission errors.

If a process or thread attached to a node fails it is generally up to the application to recover from this failure. MCAPi provides timeouts, as endpoint attributes as well as for the `mcapi_wait()` and `mcapi_wait_any()` functions and provides the `mcapi_cancel()` function to clear outstanding non-blocking requests (at the non-failing side of the communication). It is also possible to reinitialize a failed node, by first calling `mcapi_finalize()`.

3.11 MCAPi Buffer Management Philosophy

MCAPi provides two APIs for transferring buffers of data between endpoints: MCAPi messages and MCAPi packet channels. Both APIs allow the programmer to send a buffer of data on one node and receive it on another node. The APIs differ in that messaging is connectionless, allowing any node to communicate with any other node at any time, with `pr` message, priority. Packet channels, on the other hand, perform repeated communication via a connection between two endpoints.

The APIs for sending a buffer via messaging or packet channels are very similar. Both APIs have a send method that allows the programmer to specify a buffer of data with two parameters: `(void*, size_t)`. The programmer may send any data buffer they choose. There is no requirement that the buffer be allocated by MCAPi or returned to MCAPi after the send call completes.

The messaging and packet channels APIs differ in their handling of received buffers. The messaging API provides a “user-specified buffer” communications interface – the programmer specifies an empty buffer to be filled with incoming data on the receive side. The programmer can specify any buffer they choose and there is no requirement to allocate or release the buffer via an MCAPi call.

On the other hand, packet channels provide a “system-specified buffer” interface – the receive method returns a buffer of data at an address chosen by the runtime system. Since the receive buffer for a packet channel is allocated by the system, the programmer must return the buffer to the system by calling `mcapi_pktchan_release()`.

MCAPi data buffers may have arbitrary alignment. However, MCAPi does define two macros to help provide buffer alignment in performance-critical cases. Use of these macros is optional and can be defined with no value (“no-op”); implementers are required to handle send and receive buffers of arbitrary alignment, but can optimize for aligned buffers.

MCAPi implementations define the following two macros:

1. `MCAPi_DECL_ALIGNED`: a macro placed on static declarations in order to force the compiler to make the declared object aligned. For example:

```
mcapi_int_t MCAPi_DECL_ALIGNED my_int_to_send; /* See mcapi.h */
```

2. `MCAPi_BUF_ALIGN`: a macro that evaluates to the number of bytes to which dynamically allocated buffers should be aligned. For example:

```
memalign(MCAPi_BUF_ALIGN, sizeof(my_message));
```

MCAPi does not provide for any maximum buffer size. Applications may use any message size they choose, but implementations may specify a maximum buffer size and may fail with `MCAPi_ERR_MEM_LIMIT` if the system runs out of allocatable memory. Thus, the maximum size of a message in any particular implementation is limited by an implementation specified size or the amount of system memory available to that implementation.

3.12 Timeout/Cancellation Philosophy

The MCAPI API provides timeout functionality for its non-blocking calls through the timeout parameters of the `mcapi_wait()` and `mcapi_wait_any()` functions. For blocking functions implementations can optionally provide timeout capability through the use of endpoint attributes.

3.13 Backpressure mechanism

In many systems, the sender and receiver must coordinate to ensure that the messages are handled without overwhelming the receiver. Without such coordination, the sender must implement elaborate recovery mechanisms in case a message is dropped by the receiver. The overhead to handle these error cases is higher than pausing for some time and continuing later. The sender handles such scenarios by inquiring the receiver status before sending a message. The receiver returns the status based on the number of available buffers or buffer size.

The sender throttles the messages using the `mcapi_endpoint_attribute_get()` API and the `MCAPI_ATTR_NUM_RECV_BUFFERS_AVAILABLE` attribute. If required, the sender and receiver can reserve some buffers for higher priority messages. Further, a more sophisticated mechanism can be build on top of this API. For example, zones (green, yellow and red) can be defined using the number of available messages. The sender can use the zones to throttle the messages.

3.14 MCAPI Implementation Concerns

3.14.1 Link Management

The process for link configuration and link management are not specified within the MCAPI API spec. However, it is important to note that MCAPI does not define APIs or services (such as membership services) related to dynamic link state detection, node or service discovery and node or service state management. This is because MCAPI is designed to address multicore and multiprocessor systems where the number of nodes and their connectivity topology is known when the MCAPI system is configured. Since the MCAPI API specification does not specify any APIs that deal with network interfaces, an MCAPI implementation is free to manage links underneath the interface as it sees fit. In particular, an MCAPI implementation is able to restrict the topologies supported or the number of links between nodes in an MCAPI communication topology.

3.14.2 Thread Safe Implementations

MCAPI implementations are assumed to be reentrant (thread safe). Essentially, if an MCAPI implementation is available in a threaded environment, then it needs to be thread safe. MCAPI implementations can also be available in non-threaded environments, but the provider of such implementations will need to clearly indicate that the implementation is not thread safe.

3.15 MCAPI Potential Future Extensions

With the goals of implementing MCAPI efficiently, the APIs are kept simple with potential for adding more functionality on top of MCAPI later. Some specific areas include additional zero copy functionality, multicast, and informational functions for debugging, statistics (optimization) and status. These areas are strong candidates for future extensions and they are briefly described in the following subsections.

3.15.1 Zero Copy

As mentioned above zero copy can be very useful in systems where multiple nodes have shared memory. While, zero copy buffer management support can relatively simply be implemented on top of MCAPAPI providing specific buffer management API support may be useful, and something to consider for the future.

3.15.2 Multicast

MCAPAPI currently supports point-to-point communication. Communication with one sender and multiple receivers, often referred to as multi-cast, can be layered on top of MCAPAPI. For transports with (hardware) support for one-to-multiple communication, specific API functions supporting this type of communication would allow for more efficient implementations than is possible through layering.

3.15.3 Debug, Statistics and Status functions

Support functions providing information for debugging, optimization and system status are useful in most systems, and would be a valuable addition to MCAPAPI as well, and are also worth future consideration.

3.16 MCAPI Data Types

MCAPI uses predefined data types for maximum portability. The predefined MCAPI data types are defined in the following subsections. To simplify the use of multiple MCA (Multicore Association) standard API's some MCAPI data types have MCA equivalents and some MCAPI functions will have MCA equivalent functions that can be used for multiple MCA API's. An MCAPI implementation is not required to provided MCA equivalent functions.

3.16.1 mcapi_domain_t

The `mcapi_domain_t` type is used for MCAPI domains. Domain numbering is implementation defined. For application portability we recommend using symbolic constants in your code. The `mcapi_domain_t` has an `mca_domain_t` equivalent.

3.16.2 mcapi_node_t

The `mcapi_node_t` type is used for MCAPI nodes. The node numbering is implementation defined. For application portability we recommend using symbolic constants in your code. The `mcapi_node_t` has an `mca_node_t` equivalent.

3.16.3 mcapi_port_t

The `mcapi_port_t` type is used for MCAPI ports. The port numbering is implementation defined. For application portability we recommend using symbolic constants in your code.

3.16.4 mcapi_endpoint_t

The `mcapi_endpoint_t` type is used for creating and managing endpoints for sending and receiving data (see Section 2). MCAPI routines for creating and managing endpoints are described in Section 4.24.2, packet channel creation is covered in Section 4.4.14.4.1, and scalar channel creation is covered in Section 00. The `mcapi_endpoint_t` type is an opaque data type whose exact definition is implementation defined. The endpoint identifier is globally unique to an MCAPI topology.

NOTE: The MCAPI API user should not attempt to examine the contents of this data type as this can result in non-portable application code.

Implementation advice: The endpoint type must provide for topology global uniqueness, be returnable by a function, be passable as a parameter to a function and should allow simple arithmetic equality comparison (`a == b`), such as a 32- bit scalar or pointer.

3.16.5 mcapi_pktchan_rcv_hdl_t

The `mcapi_pktchan_rcv_hdl_t` type is used to receive packets from a connected packet channel (see Section 22). MCAPI routines for creating and using the `mcapi_pktchan_rcv_hdl_t` type are covered in Section 4.44.4. The `mcapi_pktchan_rcv_hdl_t` is an opaque data type whose exact definition is implementation defined.

NOTE: The MCAPI API user should not attempt to examine the contents of this data type as this can result in non-portable application code.

Implementation advice: The handle must be passable as a parameter to a function and should allow simple arithmetic equality comparison (`a == b`), such as a 32- bit scalar or pointer.

Formatted: Font:

Formatted: Font:

Formatted: Font:

Formatted: Font:

Formatted: Font:

3.16.6 mcapi_pktchan_send_hdl_t

The `mcapi_pktchan_send_hdl_t` type is used to send packets to a connected packet channel (see Section 22). MCAPi routines for creating and using the `mcapi_pktchan_send_hdl_t` type are covered in Section 4.44.4. The `mcapi_pktchan_send_hdl_t` is an opaque data type whose exact definition is implementation defined.

NOTE: The MCAPi API user should not attempt to examine the contents of this data type as this can result in non-portable application code.

Implementation advice: The handle must be passable as a parameter to a function and should allow simple arithmetic equality comparison (`a == b`), such as a 32-bit scalar or pointer.

3.16.7 mcapi_sc1chan_rcv_hdl_t

The `mcapi_sc1chan_rcv_hdl_t` type is used to receive scalars from a connected scalar channel (see Section 2). MCAPi routines for creating and using the `mcapi_sc1chan_rcv_hdl_t` type are covered in Section 4.54.5. The `mcapi_sc1chan_rcv_hdl_t` is an opaque data type whose exact definition is implementation defined.

NOTE: The MCAPi API user should not attempt to examine the contents of this data type as this can result in non-portable application code.

Implementation advice: The handle must be passable as a parameter to a function and should allow simple arithmetic equality comparison (`a == b`), such as a 32-bit scalar or pointer.

3.16.8 mcapi_sc1chan_send_hdl_t

The `mcapi_sc1chan_send_hdl_t` type is used to send scalars to a connected scalar channel (see Section 2). MCAPi routines for creating and using the `mcapi_sc1chan_send_hdl_t` type are covered in Section 4.54.5. The `mcapi_sc1chan_send_hdl_t` is an opaque data type whose exact definition is implementation defined.

NOTE: The MCAPi API user should not attempt to examine the contents of this data type as this can result in non-portable application code.

Implementation advice: The handle must be passable as a parameter to a function and should allow simple arithmetic equality comparison (`a == b`), such as a 32-bit scalar or pointer.

3.16.9 mcapi_uint64_t, mcapi_uint32_t, mcapi_uint16_t & mcapi_uint8_t,

The `mcapi_uint64_t`, `mcapi_uint32_t`, `mcapi_uint16_t`, and `mcapi_uint8_t` types are used for 64-, 32-, 16, and 8-bit scalars.

3.16.10 mcapi_request_t

The `mcapi_request_t` type is used to record the state of a pending non-blocking MCAPi transaction (see Section 3.5.43.5.4). Non-blocking MCAPi routines exist for message send and receive (see Section 4.34.3) and packet send and receive (see Section 4.44.4). The MCAPi request can only be used by the node it was created on. The `mcapi_request_t` has an `mca_request_t` equivalent.

NOTE: The MCAPi API user should not attempt to examine the contents of this data type as this can result in non-portable application code.

Formatted: Font:

Formatted: Font:

Formatted: Font:

Formatted: Font:

Formatted: Font:

Formatted: Font:

Formatted: Font:

Implementation advice: The request should allow simple arithmetic equality comparison ($a == b$), such as a 32-bit scalar or pointer.

3.16.11 `mcapi_status_t`

The `mcapi_status_t` type is an enumerated type used to record the result of a MCAP API call. If a status can be returned by an API call, the associated MCAP API call will allow a `mcapi_status_t` to be passed by reference. The API call will fill in the status code and the API user may examine the `mcapi_status_t` variable to determine the result of the call. The function return values are valid only when `mcapi_status` returns MCAP_SUCCESS. The `mcapi_status_t` has an `mca_status_t` equivalent.

3.16.12 `mcapi_timeout_t`

The `mcapi_timeout_t` type is a scalar type used to indicate the duration blocking API routine will block before reporting a timeout. The units of the `mcapi_timeout_t` data type are implementation defined since mechanisms for time keeping vary from system to system. Applications should therefore not rely on this feature for satisfaction of absolute real time constraints as its usage will not guarantee application portability across MCAP implementations. The `mcapi_timeout_t` data type is intended primarily to allow for error detection and recovery and the application developer must take appropriate action if used for any other purpose. The `mcapi_timeout_t` has an `mca_timeout_t` equivalent. A value of MCAP_TIMEOUT_IMMEDIATE (0) for the `timeout` parameter indicates that the function will return without blocking with failure or success and a value of MCAP_INFINITE (~0) for the `timeout` parameter indicates no timeout is requested, i.e. the function will block until it is unblocked because of failure or success.

3.16.13 MCAP endpoint attributes

MCAP endpoint attributes provide access to endpoint type, characteristics and state, see header files for detailed information. Implementations may designate endpoint attributes as read-only. Some endpoint attributes have to be compatible for a successful channel, as noted below.

3.16.13.1 `mcapi_endp_attr_max_payload_size_t`

This attribute defines the maximum payload size. This is a channel compatibility attribute, meaning that a channel connection requires that this attribute value is the same for both channel endpoints. Default value is implementation defined.

3.16.13.2 `mcapi_endp_attr_buffer_type_t`

This attribute defines the endpoint buffer type, at the present only FIFO type exists. It means that the order of transmission is FIFO for channels and FIFO per priority level for messages. This is a channel compatibility attribute, meaning that a channel connection requires that this attribute value is the same for both channel endpoints. Default value is implementation defined.

Buffer types: MCAP_ENDP_ATTR_FIFO_BUFFER Default

3.16.13.3 mcap_i_endp_attr_memory_type_t

This attribute defines the memory type both the memory's locality, local, shared and remote . This is a channel compatibility attribute, meaning that a channel connection requires that this attribute value is the same for both channel endpoints.

Memory locality:

MCAP_I_ENDP_ATTR_LOCAL_MEMORY	Default
MCAP_I_ENDP_ATTR_SHARED_MEMORY	
MCAP_I_ENDP_ATTR_REMOTE_MEMORY	

3.16.13.4 mcap_i_endp_attr_num_priorities_t

This attribute defines the number of endpoint priorities. This is a channel compatibility attribute, meaning that a channel connection requires that this attribute value is the same for both channel endpoints. Default value is implementation defined.

3.16.13.5 mcap_i_endp_attr_priority_t

This attribute defines the endpoint priority, applied to a channel at the time the channel is connected. This is a channel compatibility attribute, meaning that a channel connection requires that this attribute value is the same for both channel endpoints. A lower number means higher priority. A value of MCAP_I_MAX_PRIORITY (0) denotes the highest priority.

3.16.13.6 mcap_i_endp_attr_num_send_buffers_t

This attribute contains the number of send buffers at the current endpoint priority level. Default value is implementation defined.

3.16.13.7 mcap_i_endp_attr_num_rcv_buffers_t

This attribute contains the number of receive buffers available. This can for example be used for throttling. Implementation defined default value

3.16.13.8 mcap_i_endp_attr_status_t

This attribute contains endpoint status flags. Flags are used to query the status of an endpoint, e.g. if it is connected and if so what type of channel, direction, etc.

Note: The lower 16 bits are defined in mcap_i.h whereas the upper 16 bits are reserved for implementation specific purposes and if used must be defined in implementation_spec.h. It is therefore recommended that the upper 16 bits are masked off at the application level.

```
0x00000000
      ---- mcap_i.h
      ---- implementation_spec.h
```

Default = 0x00000000

Standard status flags:

MCAP_I_ENDP_ATTR_STATUS_CONNECTED	/* The endpoint is one end of a connected channel*/
MCAP_I_ENDP_ATTR_STATUS_OPEN	/* A channels is open on this endpoint*/
MCAP_I_ENDP_ATTR_STATUS_OPEN_PENDING	/* A channel open is pending */
MCAP_I_ENDP_ATTR_STATUS_CLOSE_PENDING	/* A channel close is pending */
MCAP_I_ENDP_ATTR_STATUS_PKTCHAN	/* Packet channel */
MCAP_I_ENDP_ATTR_STATUS_SCLCHAN	/* Scalar channel */
MCAP_I_ENDP_ATTR_STATUS_SEND	/* Send side */
MCAP_I_ENDP_ATTR_STATUS_RECEIVE	/* Receive side */

3.16.13.9 mcapl_endp_attr_timeout_t

This attribute contains the timeout value for blocking send and receive functions. a value of MCAPI_TIMEOUT_IMMEDIATE (or 0) means that the function will return "immediately", with success or failure. MCAPI_TIMEOUT_INFINITE means that the function will block until it completes with success or failure. Default = MCAPI_TIMEOUT_INFINITE.

3.16.14 Other MCAPI data types

MCAPI also defines its own integer, Boolean and other types, some of which have MCA equivalents. See the header files in this document for specifics on these data types.

3.16.15 MCAPI Error and Status Codes

Status code	Description
MCAPI_SUCCESS	Indicates operation was successful
MCAPI_PENDING	Indicates operation is pending without errors
MCAPI_TIMEOUT	The operation timed out
MCAPI_ERR_PARAMETER	Incorrect parameter
MCAPI_ERR_DOMAIN_INVALID	The parameter is not a valid domain
MCAPI_ERR_NODE_INVALID	The parameter is not a valid node
MCAPI_ERR_NODE_INITFAILED	The MCAPI node could not be initialized
MCAPI_ERR_NODE_INITIALIZED	MCAPI node is already initialized
MCAPI_ERR_NODE_NOTINIT	The MCAPI node is not initialized
MCAPI_ERR_NODE_FINALFAILED	The MCAPI could not be finalized
MCAPI_ERR_PORT_INVALID	The parameter is not a valid port
MCAPI_ERR_ENDP_INVALID	The parameter is not a valid endpoint descriptor
MCAPI_ERR_ENDP_EXISTS	The endpoint is already created
MCAPI_ERR_ENDP_GET_LIMIT	The endpoint get reference count is too high
MCAPI_ERR_ENDP_NOTOWNER	An endpoint can only be deleted by its creator
MCAPI_ERR_ENDP_REMOTE	Certain operations are only allowed on the node local endpoints
MCAPI_ERR_ATTR_INCOMPATIBLE	Connection of endpoints with incompatible attributes not allowed
MCAPI_ERR_ATTR_SIZE	Incorrect attribute size
MCAPI_ERR_ATTR_NUM	Incorrect attribute number
MCAPI_ERR_ATTR_VALUE	Incorrect attribute value
MCAPI_ERR_ATTR_NOTSUPPORTED	Attribute not supported by the implementation
MCAPI_ERR_ATTR_READONLY	Attribute is read only
MCAPI_ERR_MSG_SIZE	The message size exceeds the maximum size allowed by the MCAPI implementation
MCAPI_ERR_MSG_TRUNCATED	The message size exceeds the buffer size
MCAPI_ERR_CHAN_OPEN	A channel is open, certain operations are not allowed
MCAPI_ERR_CHAN_TYPE	Attempt to open a packet/scalar channel on an endpoint that has been connected with a different channel type
MCAPI_ERR_CHAN_DIRECTION	Attempt to open a send handle on a port that was connected as a receiver, or vice versa
MCAPI_ERR_CHAN_CONNECTED	A channel connection has already been established for one or both of the specified endpoints

MCAPI_ERR_CHAN_OPENPENDING	An open request is pending
MCAPI_ERR_CHAN_CLOSEPENDING	A close request is pending.
MCAPI_ERR_CHAN_NOTOPEN	The channel is not open (cannot be closed)
MCAPI_ERR_CHAN_INVALID	Argument is not a channel handle
MCAPI_ERR_PKT_SIZE	The packet size exceeds the maximum size allowed by the MCAPI implementation
MCAPI_ERR_TRANSMISSION	Transmission failure
MCAPI_ERR_PRIORITY	Incorrect priority level
MCAPI_ERR_BUF_INVALID	Not a valid buffer descriptor
MCAPI_ERR_MEM_LIMIT	Out of memory
MCAPI_ERR_REQUEST_INVALID	Argument is not a valid request handle
MCAPI_ERR_REQUEST_LIMIT	Out of request handles
MCAPI_ERR_REQUEST_CANCELLED	The request was already canceled
MCAPI_ERR_WAIT_PENDING	A wait is pending
MCAPI_ERR_GENERAL	To be used by implementations for error conditions not covered by the other status codes
MCAPI_STATUSCODE_END	This should always be last

3.16.16 API parameter readability

MCAPI_IN and MCAPI_OUT are used only for clarity to distinguish between input and output parameters.

3.17 What is new in MCAPI 2.000

The changes from version 1.0 to 2.000 are in the areas of improved API consistency, bug fixes, enhancements and added functionality.

3.17.1 MCAPI to MCA types

Some MCAPI data types and definitions were promoted to MCA data types and definitions to simplify interoperability between MCAPI and other MCA standards and to allow implementations to provide MCA versions of some functions that could be used for multiple MCA standards. For this purpose an `mca.h` header file was added.

3.17.2 MCAPI Domains

An MCAPI domain is comprised of one or more MCAPI nodes in a multicore topology and used for routing purposes. The scope of a domain is implementation defined and its scope could for example be a single chip with multiple cores or multiple processor chips on a board. The domain id is specified once at node initialization. A domain can contain multiple nodes. Some example potential uses for domains are topologies that may change dynamically, include non-MCAPI sub-topologies, require separation between different transports or have open and secure areas.

3.17.3 MCAPI endpoint attributes

A few attributes were added, for details see the attribute section for details. A certain number of endpoint attributes were reserved for MCA. Vendors desiring to add additional endpoint attributes in their implementation can request to be assigned a range of vendor specific endpoint attribute numbers from the MCA. Attribute ranges are defined in `mca.h`.

3.17.4 MCAPI header files

The header files have been restructured to better separate standard and implementation specific definitions.

3.17.5 New functionality and functions

3.17.5.1 Initialization parameters and information

Initialization parameters were added to allow implementations to configure the MCAPI at run time. A parameter was also added to allow implementations to provide information about the MCAPI runtime with both MCAPI specified and implementation specific information. For specifics on MCAPI information see below. The topology information can for example be used for basic initial system discovery by using the `number_of_domains` and `number_of_nodes` parameters to establish communications for more thorough discovery.

3.17.5.1.1 mcapi_param_t

Initialization parameters will vary by implementation, and may include specifications of the amount of resources to be used for a specific implementation or configuration, such as for example the maximum number of outstanding requests, etc. This simply means that each implementations has to provide documentation and/or sample code for the initialization. Some initialization parameters may be standardized in future versions of MCAPI.

3.17.5.1.2 mcapi_info_t

The informational parameters include MCAPI specified information as outlined below, as well as implementation specific information. Implementation specific information must be documented by the implementer.

MCAPI defined initialization information:

mcapi_version	MCAPI version, the three last (rightmost) hex digits are the minor number and those left of minor the major number.
organization_id	Implementation vendor/organization id. Assigned by MCA.
implementation_version	Implementation version, represented by a 32 bit scalar, the specific format is implementation defined.
number_of_domains	Number of domains in the topology.
number_of_nodes	Number of nodes in the domain, can be used for basic per domain topology discovery. Note: This information may not be available in all implementations. The functionality must be documented.
number_of_ports	Number of available ports on the local node.

3.17.5.2 MCAPI_domains

An MCAPI domain is comprised of one or more MCAPI nodes in a multicore topology and used for routing purposes. The scope of a domain is implementation defined and could for example be a single chip with multiple cores or multiple processor chips on a board. The domain id is specified once at node initialization. A domain can contain multiple nodes. Some example potential uses for domains are topologies that may change dynamically, include non-MCAPI sub-topologies or have open and secure areas.

3.17.5.3 mcapi_domain_id_get()

This function was added to allow the application to find its MCAPI domain id. Other functions affected are: mcapi_initialize(), mcapi_endpoint_get_i() and mcapi_endpoint_get().

3.17.5.4 mcapi_endpoint_get() functionality

To avoid the potential for blocking indefinitely on this function a timeout parameter was added.

3.17.5.5 mcapi_node_init_attribute()

This function was added to allow initialization of a node's attributes structure for setting of non-default node attributes to be used by mcapi_initialize().

3.17.5.6 mcapi_node_set_attribute()

This function was added to allow setting of non-default node attributes to be used by mcapi_initialize().

3.17.5.7 mcapi_node_get_attribute()

This function was added to allow the application to query other nodes about their attributes, e.g. type of node, characteristics, etc. This functionality takes basic topology discovery a step further and simplifies applying MCAPi to a multitude of diverse multicore platforms.

3.17.5.8 mcapi_pktchan_release_test()

This function was added to allow the sender side of a packet channel to find out if a buffer has been released by the receiver for reuse. This may for example be useful for zero copy operations.

3.17.5.9 mcapi_display_status()

This function was added to enable displaying an MCAPi status code in text format, for convenience.

3.17.6 API Consistency

3.17.6.1 API function nomenclature

Some API function names were modified to adhere to a consistent nomenclature. The nomenclature is mcapi_<functional area>_<action>, e.g. mcapi_msg_send().

Editors note: Request that technical writer compile a list of changed functions.]

Comment [S5]: For tech editor.

3.17.6.2 Error and status codes

Error and status codes were modified for consistency and others were added to allow for more precise error and status reporting.

3.17.7 Definitions and constants

Some definitions and constants have been more clearly defined and specified.

3.17.8 Clarifications

Clarifications were added throughout the document to improve the understanding of the MCAPI specification.

3.18 MCAPI 2.000 to 1.0 backwards compatibility

To address backwards compatibility as much as possible, a new header file `mcapi_v1000_to_v2000.h` and a compatibility function were added. A few notes on the backwards compatibility follow.

3.18.1 Affected functions

3.18.1.1 `mcapi_initialize()`

The `mcapi_initialize()` function has the additional parameters `domain_id` and `init_parameters` and modified `mcapi_info` parameter. An `mcapi_1000_initialize()` function is added in the header file showing how to map to the v 2.000 `mcapi_initialize()` function. MCAPI 1.0 nodes will always be in domain 0.

3.18.1.2 `mcapi_endpoint_get_i()` and `mcapi_endpoint_get()`

The `mcapi_endpoint_get_i()` and `mcapi_endpoint_get()` functions have the additional `domain_id` parameter and `mcapi_endpoint_get()` a new timeout parameter. `mcapi_1000_endpoint_get_i()` and `mcapi_1000_endpoint_get()` functions are added in the header file showing how to map to the v 2.000 functions. MCAPI 1.0 endpoints will always be in domain 0.

3.18.1.3 Modified function names and status/error codes

Functions with modified names and unchanged parameters are handled with defines. Modified status/error codes are handled with defines.

4 MCAPI API

4.1 General

This section describes initialization, finalization and introspection functions.

4.1.1 MCAP1_INITIALIZE

NAME

`mcapi_initialize` – Initializes an MCAP1 node.

Comment [S6]: I suggest we ask the tech writer for advice on this.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_initialize(
    MCAP1_IN mcapi_domain_t domain_id,
    MCAP1_IN mcapi_node_t node_id,
    MCAP1_IN mcapi_node_attributes_t* mcapi_node_attributes,
    MCAP1_IN mcapi_param_t* mcapi_parameters,
    MCAP1_OUT mcapi_info_t* mcapi_info,
    MCAP1_OUT mcapi_status_t* mcapi_status
);
```

Comment [ML7]: Do you want to keep this brief description? Note difference to MRAP1 doc.

DESCRIPTION

`mcapi_initialize()` initializes the MCAP1 environment on a given MCAP1 node in a given MCAP1 domain. It has to be called by each node using MCAP1. `mcapi_node_attributes` is used to pass MCAP1 defined initialization parameters. `mcapi_parameters` is used to pass implementation specific initialization parameters. `mcapi_info` is used to obtain information from the MCAP1 implementation, including MCAP1 standardized information, see below and the header files and vendor specific implementation information. A node is a process, a thread, or a processor (or core) with an independent program counter running a piece of code. In other words, an MCAP1 node is an independent thread of control. An MCAP1 node can call `mcapi_initialize()` once per node, and it is an error to call `mcapi_initialize()` multiple times from a given node, unless `mcapi_finalize()` is called in between. A given MCAP1 implementation will specify what is a node (i.e., what thread of control – process, thread, or other -- is a node) in that implementation. A thread and process are just two examples of threads of control, and there could be others.

Initialization parameters will vary by implementation, and may include specifications of the amount of resources to be used for a specific implementation, such as the maximum number of outstanding requests, etc. An `mcapi_node_attributes` structure is passed in by reference with the `mcapi_parameters`. A NULL value for the `mcapi_node_attributes` pointer indicates that default values should be set.

MCAP1 defined node attributes:

`MCAP1_NODE_ATTR_TYPE_REGULAR` The node is regular. Default

The informational parameters include MCAP1 specified information as outlined below, as well as implementation specific information. Implementation specific information shall be documented by the implementer.

MCAP1 defined initialization information:

`mcapi_version` MCAP1 version, the three last (rightmost) hex digits are the minor number and those left of minor the major number.

<code>organization_id</code>	Implementation vendor/organization id.
<code>implementation_version</code>	Vendor version, the three last (rightmost) hex digits are the minor number and those left of minor the major number.
<code>implementation_version</code>	Vendor version, the three last (rightmost) hex digits are the minor number and those left of minor the major number.
<code>number_of_domains</code>	Number of domains in the topology.
<code>number_of_nodes</code>	Number of nodes in the domain, can be used for basic per domain topology discovery.
<code>number_of_ports</code>	Number of ports on the local node

RETURN VALUE

On success, `*mcapi_status` is set to `MCAPI_SUCCESS`. On error, `*mcapi_status` is set to the appropriate error defined below.

ERRORS

<code>MCAPI_ERR_NODE_INITFAILED</code>	The MCAPI environment could not be initialized.
<code>MCAPI_ERR_NODE_INITIALIZED</code>	The MCAPI environment has already been initialized.
<code>MCAPI_ERR_NODE_INVALID</code>	The parameter is not a valid node.
<code>MCAPI_ERR_DOMAIN_INVALID</code>	The parameter is not a valid domain.
<code>MCAPI_ERR_PARAMETER</code>	Incorrect <code>mcapi_parameters</code> or <code>mcapi_info</code> parameter. Note, <code>mcapi_parameters</code> can be a NULL pointer, if the MCAPI version is 1.000, for backwards compatibility purposes.
<code>MCAPI_ERR_GENERAL</code>	Implementation specific error not covered by other status codes. Specifics must be documented.

SEE ALSO

`mcapi_finalize()`

4.1.2 MCAPI_FINALIZE

NAME

mcapi_finalize – Finalize un-initializes an MCAPI node..

SYNOPSIS

```
#include <mcapi.h>

void mcapi_finalize(
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

mcapi_finalize() finalizes the MCAPI environment on the local MCAPI node. It has to be called by each node using MCAPI, to un-initialize the node. It is an error to call mcapi_finalize() without first calling mcapi_initialize(). An MCAPI node can call mcapi_finalize() once for each call to mcapi_initialize(), but it is an error to call mcapi_finalize() multiple times from a given node unless mcapi_initialize() has been called prior to each mcapi_finalize() call. Pending data and outstanding requests will be discarded when mcapi_finalize() is called. Implementations must document other specific functionality of mcapi_finalize(), if applicable.

Comment [ML8]: MCAPI node and domain?

Comment [S9]: The node finalizes itself, so I don't think it is necessary include the domain.

RETURN VALUE

On success, *mcapi_status is set to MCAPI_SUCCESS. On error, *mcapi_status is set to the appropriate error defined below.

ERRORS

MCAPI_ERR_NODE_NOTINIT	The node is not initialized.
MCAPI_ERR_NODE_FINALFAILED	The MCAPI environment could not be finalized.
MCAPI_ERR_GENERAL	Implementation specific error not covered by other status codes. Specifics must be documented.

SEE ALSO

```
mcapi_initialize()
```


4.1.3 MCAPI_DOMAIN_ID_GET

NAME

mcapi_domain_id_get – Returns the domain id associated with the local node.

SYNOPSIS

```
#include <mcapi.h>

mcapi_domain_t mcapi_domain_id_get(
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Returns the domain id associated with the local node.

RETURN VALUE

On success the domain_id is returned, *mcapi_status is set to MCAPI_SUCCESS. On error, *mcapi_status is set to the appropriate error defined below and the return value is set to MCAPI_DOMAIN_INVALID.

ERRORS

MCAPI_ERR_NODE_NOTINIT	The node is not initialized.
MCAPI_ERR_GENERAL	Implementation specific error not covered by other status codes. Specifics must be documented.

NOTE

SEE ALSO

4.1.4 MCAPI_NODE_ID_GET

NAME

mcapi_node_id_get – Returns the node id associated with the local node.

SYNOPSIS

```
#include <mcapi.h>

mcapi_node_t mcapi_node_id_get(

    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Returns the node id associated with the local node.

RETURN VALUE

On success the node_id is returned, *mcapi_status is set to MCAPI_SUCCESS. On error, *mcapi_status is set to the appropriate error defined below and the return value is set to MCAPI_NODE_INVALID.

ERRORS

MCAPI_ERR_NODE_NOTINIT	The node is not initialized.
MCAPI_ERR_GENERAL	Implementation specific error not covered by other status codes. Specifics must be documented.

NOTE

SEE ALSO

4.1.5 MCAPI_NODE_INIT_ATTRIBUTES

NAME

mcapi_node_init_attributes - Initializes node attribute structure.

SYNOPSIS

```
void mcapi_node_init_attributes(  
    MCAPI_OUT mcapi_node_attributes_t* mcapi_node_attributes,  
    MCAPI_OUT mcapi_status_t* mcapi_status  
);
```

DESCRIPTION

This function initializes the values of an mcapi_node_attributes_t structure. For non-default behavior this function should be called prior to calling mcapi_node_set_attribute(). mcapi_node_set_attribute() is then used to change any default values prior to calling mcapi_initialize().

RETURN VALUE

On success *mcapi_status is set to MCAPI_SUCCESS. On error, *mcapi_status is set to the appropriate error defined below.

ERRORS

MCAPI_ERR_PARAMETER	Invalid attributes parameter.
MCAPI_ERR_GENERAL	Implementation specific error not covered by other status codes. Specifics must be documented.

NOTE

SEE ALSO

4.1.6 MCAPI_NODE_SET_ATTRIBUTE

NAME

mcapi_node_set_attribute - Sets the Node attributes to non-default values. The node attributes are passed as a parameter by mcapi_initialize().

SYNOPSIS

```
void mcapi_node_set_attribute(
    MCAPI_OUT mcapi_node_attributes_t* mcapi_node_attributes,
    MCAPI_IN mcapi_uint_t attribute_num,
    MCAPI_IN void* attribute,
    MCAPI_IN size_t attribute_size,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

This function is used to change default values of an mcapi_node_attributes_t data structure prior to calling mcapi_initialize(). This is a blocking function. Calls to this function have no effect on node attributes once the MCAPI has been initialized. The purpose of this function is to define node specific characteristics, and in this MCAPI version has only attribute. The node attributes are expected to be expanded in future versions.

MCAPI pre-defined node attributes:

Attribute num:	Description:	Data type:	Default:
MCAPI_NODE_ATTR_TYPE	Indicates the node type.	mcapi_uint_t	MCAPI_NODE_ATTR_TYPE_REGULAR

RETURN VALUE

On success *mcapi_status is set to MCAPI_SUCCESS. On error, *mcapi_status is set to the appropriate error defined below.

ERRORS

MCAPI_ERR_ATTR_NUM	Unknown attribute number.
MCAPI_ERR_ATTR_VALUE	Incorrect attribute value.
MCAPI_ERR_ATTR_SIZE	Incorrect attribute size.
MCAPI_ERR_ATTR_NOTSUPPORTED	Attribute not supported by the implementation.
MCAPI_ERR_ATTR_READONLY	Attribute cannot be modified.
MCAPI_ERR_PARAMETER	Incorrect mcapi_node_attributes or attribute parameter.

MCAPI_ERR_GENERAL

Implementation specific error not covered by other status codes. Specifics must be documented.

NOTE

SEE ALSO

4.1.7 MCAPI_NODE_GET_ATTRIBUTE

NAME

`mcapi_node_get_attribute` – Get node attributes from a remote node.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_node_get_attribute(
    MCAPI_IN mcapi_domain_t domain_id,
    MCAPI_IN mcapi_node_t node_id,
    MCAPI_IN mcapi_uint_t attribute_num,
    MCAPI_OUT void* attribute,
    MCAPI_IN size_t attribute_size,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

`mcapi_node_get_attribute()` allows the application to query other nodes about their attributes, e.g. type of node, characteristics, etc. `domain_id` and `node_id` specify the node being queried, `attribute_num` indicates which one of the node attributes is being referenced. This is a blocking function. `attribute` points to a structure or scalar to be filled with the value of the attribute specified by `attribute_num`. `attribute_size` is the size in bytes of the attribute. See Section 3.2 and header files for a description of attributes. The `mcapi_node_get_attribute()` function returns the requested attribute value by reference.

MCAPI node attributes:

Attribute num:	Description:	Data type:	Default:
MCAPI_NODE_ATTR_TYPE	Indicates the node type.	<code>mcapi_uint_t</code>	MCAPI_NODE_ATTR_TYPE_REGULAR

RETURN VALUE

On success, `*attribute` is filled with the requested attribute and `*mcapi_status` is set to `MCAPI_SUCCESS`. On error, `*mcapi_status` is set to an error code and `*attribute` is not modified.

ERRORS

MCAPI_ERR_NODE_NOTINIT	The local node is not initialized.
MCAPI_ERR_DOMAIN_INVALID	The parameter is not a valid domain.

MCAPI_ERR_NODE_INVALID	The parameter is not a valid node.
MCAPI_ERR_ATTR_NUM	Unknown attribute number.
MCAPI_ERR_ATTR_SIZE	Incorrect attribute size.
MCAPI_ERR_ATTR_NOTSUPPORTED	Attribute not supported by the implementation.
MCAPI_ERR_PARAMETER	Incorrect attribute parameter.
MCAPI_ERR_GENERAL	Implementation specific error not covered by other status codes. Specifics must be documented.

NOTE

SEE ALSO

4.2 Endpoints

This section describes API functions that create, delete, get and modify endpoints.

4.2.1 MCAPI_ENDPOINT_CREATE

NAME

mcapi_endpoint_create - Create an endpoint.

SYNOPSIS

```
#include <mcapi.h>

mcapi_endpoint_t mcapi_endpoint_create(
    MCAPI_IN mcapi_port_t port_id,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

mcapi_endpoint_create() is used to create an endpoint on the local node with the specified port_id. A port_id of MCAPI_PORT_ANY is used to request the next available endpoint on the local node.

MCAPI supports a simple static naming scheme to create endpoints based on global tuple names, <domain_id, node_id, port_id>. Other nodes can access the created endpoint by calling mcapi_endpoint_get() and specifying the appropriate domain, node and port id.

Static naming allows the programmer to define an MCAPI communication topology at compile time. This facilitates simple initialization. Section 7.17.4 illustrates an example of initialization and bootstrapping using static naming. Creating endpoints using MCAPI_PORT_ANY provides a convenient method to create endpoints without having to specify the port_id.

There are three types of endpoints, message, packet channel and scalar channel. The endpoint type defines certain aspects of the endpoint's behavior. The endpoint type is used to manage avoidance of messages being sent to connected endpoints. The type is set with mcapi_endpoint_set_attribute(). An endpoint is always created as a message type, which is the default type.

RETURN VALUE

On success, an endpoint is returned and *mcapi_status is set to MCAPI_SUCCESS. On error, MCAPI_NULL (or 0) is returned and *mcapi_status is set to the appropriate error defined below. MCAPI_NULL (or 0) could be a valid endpoint value so status has to be checked to ensure correctness.

ERRORS

MCAPI_ERR_PORT_INVALID The parameter is not a valid port. There may be no more available ports or the port may be reserved.

MCAPI_ERR_ENDP_EXISTS The endpoint is already created.

MCAPI_ERR_NODE_NOTINIT The node is not initialized.

MCAPI_ERR_GENERAL

Implementation specific error not covered by other status codes. Specifics must be documented.

NOTE

The node id can only be set using the `mcapi_initialize()` function.

SEE ALSO

`mcapi_initialize()`

4.2.2 MCAPI_ENDPOINT_GET_I

NAME

mcapi_endpoint_get_i – Obtain the endpoint associated with a given tuple.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_endpoint_get_i(
    MCAPI_IN mcapi_domain_t domain_id,
    MCAPI_IN mcapi_node_t node_id,
    MCAPI_IN mcapi_port_t port_id,
    MCAPI_OUT mcapi_endpoint_t* endpoint,
    MCAPI_OUT mcapi_request_t* request,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

mcapi_endpoint_get_i() allows other nodes (“third parties”) to get the endpoint identifier for the endpoint associated with a global tuple name <domain_id, node_id, port_id>. This function is non-blocking and will return immediately. An endpoint can receive messages from a multitude of other endpoints. Message type endpoints can therefore have multiple outstanding endpoint gets from other nodes. Channel connected endpoints on the other hand has a one to one relationship. Channel type endpoints can therefore only have one outstanding get. A second get on a channel type endpoint will result in an error.

RETURN VALUE

On success, *mcapi_status is set to MCAPI_SUCCESS if completed and MCAPI_PENDING if not yet completed. On error, *mcapi_status is set to the appropriate error defined below.

ERRORS

MCAPI_ERR_PORT_INVALID	The parameter is not a valid port. This error also covers endpoints without ports.
MCAPI_ERR_NODE_INVALID	The parameter is not a valid node.
MCAPI_ERR_NODE_NOTINIT	The node is not initialized.
MCAPI_ERR_DOMAIN_INVALID	The parameter is not a valid domain.
MCAPI_ERR_REQUEST_LIMIT	No more request handles available.
MCAPI_ERR_ENDP_GET_LIMIT	The endpoint get reference count is too high. Use service function to check this status.

MCAPI_ERR_PARAMETER

Incorrect endpoint or request parameter.

MCAPI_ERR_GENERAL

Implementation specific error not covered by other status codes. Specifics must be documented.

NOTE

Use the `mcapi_test()`, `mcapi_wait()` and `mcapi_wait_any()` functions to query the status of and `mcapi_cancel()` function to cancel the operation.

SEE ALSO

`mcapi_node_id_get()`

4.2.3 MCAPI_ENDPOINT_GET

NAME

`mcapi_endpoint_get` – Obtain the endpoint associated with a given tuple.

SYNOPSIS

```
#include <mcapi.h>

mcapi_endpoint_t mcapi_endpoint_get(
    MCAPI_IN mcapi_domain_t domain_id,
    MCAPI_IN mcapi_node_t node_id,
    MCAPI_IN mcapi_port_t port_id,
    MCAPI_IN mcapi_timeout_t timeout,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

`mcapi_endpoint_get()` allows other nodes (“third parties”) to get the endpoint identifier for the endpoint associated with a global tuple name `<domain_id, node_id, port_id>`. This function will block until the specified remote endpoint has been created via the `mcapi_endpoint_create()` call or a timeout is reached. An endpoint can receive messages from a multitude of other endpoints. Message type endpoints can therefore have multiple outstanding endpoint gets from other nodes. Channel connected endpoints on the other hand has a one to one relationship. Channel type endpoints can therefore only have one outstanding get. A second get on a channel type endpoint will result in an error. A timeout value of `MCAPI_TIMEOUT_INFINITE` would cause this function to block until completion (success or failure).

RETURN VALUE

On success, an endpoint is returned and `*mcapi_status` is set to `MCAPI_SUCCESS`. On error, `MCAPI_NULL` (or 0) is returned and `*mcapi_status` is set to the appropriate error defined below. `MCAPI_NULL` (or 0) could be a valid endpoint value so status has to be checked to ensure correctness.

ERRORS

<code>MCAPI_ERR_PORT_INVALID</code>	The parameter is not a valid port. This error also covers endpoints without ports (routing nodes).
<code>MCAPI_ERR_NODE_INVALID</code>	The parameter is not a valid node.
<code>MCAPI_ERR_DOMAIN_INVALID</code>	The parameter is not a valid domain.
<code>MCAPI_ERR_NODE_NOTINIT</code>	The node is not initialized.

MCAPI_ERR_ENDP_GET_INVALID	The endpoint get reference count is too high.
MCAPI_TIMEOUT	The operation timed out.
MCAPI_ERR_GENERAL	Implementation specific error not covered by other status codes. Specifics must be documented.

4.2.4 MCAPI_ENDPOINT_DELETE

NAME

`mcapi_endpoint_delete` – delete an endpoint.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_endpoint_delete(
    MCAPI_IN mcapi_endpoint_t endpoint,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Deletes an MCAPI endpoint. Pending messages are discarded. If an endpoint has been connected to a packet or scalar channel, the appropriate close method must be called before deleting the endpoint. Delete is a blocking operation. Since the connection is closed before deleting the endpoint, the delete method does not require any cross-process synchronization and is guaranteed to return in a timely manner (operation will return without having to block on any IPC to any remote nodes). It is an error to attempt to delete an endpoint that has not been closed. Only the node that created an endpoint can delete it. An endpoint that has an open pending but is not yet connected can be deleted.

RETURN VALUE

On success, `*mcapi_status` is set to `MCAPI_SUCCESS`. On error, `*mcapi_status` is set to the appropriate error defined below.

ERRORS

<code>MCAPI_ERR_NODE_NOTINIT</code>	The node is not initialized.
<code>MCAPI_ERR_ENDP_INVALID</code>	Argument is not a valid endpoint descriptor.
<code>MCAPI_ERR_CHAN_CONNECTED</code>	A channel is connected, deletion is not allowed
<code>MCAPI_ERR_ENDP_NOTOWNER</code>	An endpoint can only be deleted by its creator.
<code>MCAPI_ERR_GENERAL</code>	Implementation specific error not covered by other status codes. Specifics must be documented.

SEE ALSO

`mcapi_endpoint_create()`

4.2.5 MCAPI_ENDPOINT_GET_ATTRIBUTE

NAME

`mcapi_endpoint_get_attribute`— Get endpoint attributes.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_endpoint_get_attribute(
    MCAPI_IN mcapi_endpoint_t endpoint,
    MCAPI_IN mcapi_uint_t attribute_num,
    MCAPI_OUT void* attribute,
    MCAPI_IN size_t attribute_size,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

`mcapi_endpoint_get_attribute()` allows the programmer to query endpoint attributes on both node local and remote endpoints. This is a blocking function. `attribute_num` indicates which one of the endpoint attributes is being referenced. `attribute` points to a structure or scalar to be filled with the value of the attribute specified by `attribute_num`. `attribute_size` is the size in bytes of the memory pointed to by "attribute". See Section 3.2 and header files for a description of attributes.

It is an error to attempt a connection between endpoints whose attributes are set in an incompatible way (whether attributes are compatible or not is implementation defined). It is also an error to attempt to change the attributes of endpoints that are connected.

RETURN VALUE

On success, `*attribute` is filled with the requested attribute and `*mcapi_status` is set to `MCAPI_SUCCESS`. On error, `*mcapi_status` is set to an error code and `*attribute` is not modified.

ATTRIBUTES

`mcapi_endp_attr_max_payload_size_t`

This attribute defines the maximum payload size. This is a channel compatibility attribute, meaning that a channel connection requires that this attribute value is the same for both channel endpoints. Default value is implementation defined.

`mcapi_endp_attr_buffer_type_t`

This attribute defines the endpoint buffer type, at the present only FIFO type exists. It means that the order of transmission is FIFO for channels and FIFO per priority level for messages. This is a channel compatibility attribute, meaning that a channel connection requires that this attribute value is the same for both channel endpoints. Default value is implementation defined.

Buffer types: `MCAPI_ENDP_ATTR_FIFO_BUFFER` Default

`mcapi_endp_attr_memory_type_t`

This attribute defines the memory type both the memory's locality, local, shared and remote . This is a channel compatibility attribute, meaning that a channel connection requires that this attribute value is the same for both channel endpoints.

Memory locality:

```
MCAPI_ENDP_ATTR_LOCAL_MEMORY      Default
MCAPI_ENDP_ATTR_SHARED_MEMORY
MCAPI_ENDP_ATTR_REMOTE_MEMORY
```

mcapi_endp_attr_num_priorities_t

This attribute defines the number of endpoint priorities. This is a channel compatibility attribute, meaning that a channel connection requires that this attribute value is the same for both channel endpoints. Default value is implementation defined.

mcapi_endp_attr_priority_t

This attribute defines the endpoint priority, applied to a channel at the time the channel is connected. This is a channel compatibility attribute, meaning that a channel connection requires that this attribute value is the same for both channel endpoints. A lower number means higher priority. A value of MCAPI_MAX_PRIORITY (0) denotes the highest priority.

mcapi_endp_attr_num_send_buffers_t

This attribute contains the number of send buffers at the current endpoint priority level. Default value is implementation defined.

mcapi_endp_attr_num_rcv_buffers_t

This attribute contains the number of receive buffers available. This can for example be used for throttling. Implementation defined default value

mcapi_endp_attr_status_t

This attribute contains endpoint status flags. Flags are used to query the status of an endpoint, e.g. if it is connected and if so what type of channel, direction, etc.

Note: The lower 16 bits are defined in mcapi.h whereas the upper 16 bits are reserved for implementation specific purposes and if used must be defined in implementation_spec.h. It is therefore recommended that the upper 16 bits are masked off at the application level.

```
0x00000000
----- mcapi.h
----- implementation_spec.h
```

Default = 0x00000000

Standard status flags:

```
MCAPI_ENDP_ATTR_STATUS_CONNECTED /* The endpoint is one end of a
                                   connected channel */
MCAPI_ENDP_ATTR_STATUS_OPEN      /* A channels is open on this
                                   endpoint */
MCAPI_ENDP_ATTR_STATUS_OPEN_PENDING /* A channel open is pending */
MCAPI_ENDP_ATTR_STATUS_CLOSE_PENDING /* A channel close is pending */
MCAPI_ENDP_ATTR_STATUS_PKTCHAN    /* Packet channel */
MCAPI_ENDP_ATTR_STATUS_SCLCHAN    /* Scalar channel */
MCAPI_ENDP_ATTR_STATUS_SEND       /* Send side */
MCAPI_ENDP_ATTR_STATUS_RECEIVE    /* Receive side */
```

mcapi_endp_attr_timeout_t

This attribute contains the timeout value for blocking send and receive functions. a value of MCAPI_TIMEOUT_IMMEDIATE (or 0) means that the function will return "immediately", with success or failure. MCAPI_TIMEOUT_INFINITE means that the function will block until it completes with success or failure. Default = MCAPI_TIMEOUT_INFINITE. |

Comment [S11R10]: For tech editor: please organize in a table.

ERRORS

MCAPI_ERR_NODE_NOTINIT	The node is not initialized.
MCAPI_ERR_ENDP_INVALID	Argument is not an endpoint descriptor.
MCAPI_ERR_ATTR_NUM	Unknown attribute number.
MCAPI_ERR_ATTR_SIZE	Incorrect attribute size.
MCAPI_ERR_ATTR_NOTSUPPORTED	Attribute not supported by the implementation.
MCAPI_ERR_PARAMETER	Incorrect attribute parameter.
MCAPI_ERR_GENERAL	Implementation specific error not covered by other status codes. Specifics must be documented.

SEE ALSO

`mcapi_endpoint_set_attribute()`

4.2.6 MCAPI_ENDPOINT_SET_ATTRIBUTE

NAME

`mcapi_endpoint_set_attribute` - Set endpoint attributes.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_endpoint_set_attribute(
    MCAPI_IN mcapi_endpoint_t endpoint,
    MCAPI_IN mcapi_uint_t attribute_num,
    MCAPI_IN const void* attribute,
    MCAPI_IN size_t attribute_size,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

`mcapi_endpoint_set_attribute()` allows the programmer to assign endpoint attributes on the local node. `attribute_num` indicates which one of the endpoint attributes is being referenced. `attribute` points to a structure or scalar to be filled with the value of the attribute specified by `attribute_num`. `attribute_size` is the size in bytes of the memory pointed to by `attribute`. See Section 3.33.3 and `mcapi.h` for a description of attributes. Endpoint attributes can be set only on the local node.

It is an error to attempt a connection between endpoints whose attributes are set in an incompatible way (whether attributes are compatible or not is implementation defined). It is also an error to attempt to change the attributes of endpoints that are connected.

RETURN VALUE

On success, `*mcapi_status` is set to `MCAPI_SUCCESS`. On error, `*mcapi_status` is set to the appropriate error defined below.

ATTRIBUTES

`mcapi_endp_attr_max_payload_size_t`

This attribute defines the maximum payload size. This is a channel compatibility attribute, meaning that a channel connection requires that this attribute value is the same for both channel endpoints. Default value is implementation defined.

`mcapi_endp_attr_buffer_type_t`

This attribute defines the endpoint buffer type, at the present only FIFO type exists. It means that the order of transmission is FIFO for channels and FIFO per priority level for messages. This is a channel compatibility attribute, meaning that a channel connection requires that this attribute value is the same for both channel endpoints. Default value is implementation defined.

Buffer types: `MCAPI_ENDP_ATTR_FIFO_BUFFER` Default

`mcapi_endp_attr_memory_type_t`

This attribute defines the memory type both the memory's locality, local, shared and remote. This is a channel compatibility attribute, meaning that a channel connection requires that this attribute value is the same for both channel endpoints.

Formatted: Font: (Default) Arial

Memory locality:	
MCAPI_ENDP_ATTR_LOCAL_MEMORY	Default
MCAPI_ENDP_ATTR_SHARED_MEMORY	
MCAPI_ENDP_ATTR_REMOTE_MEMORY	

mcapi_endp_attr_num_priorities_t

This attribute defines the number of endpoint priorities. This is a channel compatibility attribute, meaning that a channel connection requires that this attribute value is the same for both channel endpoints. Default value is implementation defined.

mcapi_endp_attr_priority_t

This attribute defines the endpoint priority, applied to a channel at the time the channel is connected. This is a channel compatibility attribute, meaning that a channel connection requires that this attribute value is the same for both channel endpoints. A lower number means higher priority. A value of MCAPI_MAX_PRORITY (0) denotes the highest priority.

mcapi_endp_attr_num_send_buffers_t

This attribute contains the number of send buffers at the current endpoint priority level. Default value is implementation defined.

mcapi_endp_attr_num_rcv_buffers_t

This attribute contains the number of receive buffers available. This can for example be used for throttling. Implementation defined default value

mcapi_endp_attr_status_t

This attribute contains endpoint status flags. Flags are used to query the status of an endpoint, e.g. if it is connected and if so what type of channel, direction, etc.

Note: The lower 16 bits are defined in mcapi.h whereas the upper 16 bits are reserved for implementation specific purposes and if used must be defined in implementation_spec.h. It is therefore recommended that the upper 16 bits are masked off at the application level.

0x00000000	
-----	mcapi.h
-----	implementation_spec.h

Default = 0x00000000

Standard status flags:

MCAPI_ENDP_ATTR_STATUS_CONNECTED	/* The endpoint is one end of a connected channel*/
MCAPI_ENDP_ATTR_STATUS_OPEN	/* A channels is open on this endpoint*/
MCAPI_ENDP_ATTR_STATUS_OPEN_PENDING	/* A channel open is pending */
MCAPI_ENDP_ATTR_STATUS_CLOSE_PENDING	/* A channel close is pending */
MCAPI_ENDP_ATTR_STATUS_PKTCHAN	/* Packet channel */
MCAPI_ENDP_ATTR_STATUS_SCLCHAN	/* Scalar channel */
MCAPI_ENDP_ATTR_STATUS_SEND	/* Send side */
MCAPI_ENDP_ATTR_STATUS_RECEIVE	/* Receive side */

mcapi_endp_attr_timeout_t

This attribute contains the timeout value for blocking send and receive functions. a value of MCAPI_TIMEOUT_IMMEDIATE (or 0) means that the function will return "immediately", with success or failure. MCAPI_TIMEOUT_INFINITE means that the function will block until it completes with success or failure. Default = MCAPI_TIMEOUT_INFINITE.

Comment [S12]: For tech editor: please organize in a table.

ERRORS

MCAPI_ERR_NODE_NOTINIT	The node is not initialized.
MCAPI_ERR_ENDP_INVALID	Argument is not an endpoint descriptor.
MCAPI_ERR_ATTR_NUM	Unknown attribute number.
MCAPI_ERR_ATTR_SIZE	Incorrect attribute size.
MCAPI_ERR_ATTR_VALUE	Incorrect attribute value. Attribute value ranges are implementation defined.
MCAPI_ERR_CHAN_CONNECTED	Attribute changes not allowed on connected endpoints.
MCAPI_ERR_ATTR_READONLY	Attribute cannot be modified.
MCAPI_ERR_ENDP_REMOTE	Attributes can only be set on local endpoints.
MCAPI_ERR_ATTR_NOTSUPPORTED	Attribute not supported by the implementation.
MCAPI_ERR_PARAMETER	Incorrect attribute parameter (NULL pointer).
MCAPI_ERR_GENERAL	Implementation specific error not covered by other status codes. Specifics must be documented.

SEE ALSO

`mcapi_endpoint_set_attribute()`

4.3 MCAPI Messages

MCAPI Messages facilitate connectionless transfer of data buffers. The messaging API provides a “user-specified buffer” communications interface – the programmer specifies a buffer of data to be sent on the send side, and the user specifies an empty buffer to be filled with incoming data on the receive side. The implementation must be able to transfer messages to and from any buffer the programmer specifies, although the implementation may use extra buffering internally to queue up data between the sender and receiver.

4.3.1 MCAPI_MSG_SEND_I

NAME

`mcapi_msg_send_i` – Sends a (connectionless) message from a send endpoint to a receive endpoint.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_msg_send_i(
    MCAPI_IN mcapi_endpoint_t send_endpoint,
    MCAPI_IN mcapi_endpoint_t receive_endpoint,
    MCAPI_IN void* buffer,
    MCAPI_IN size_t buffer_size,
    MCAPI_IN mcapi_priority_t priority,
    MCAPI_OUT mcapi_request_t* request,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Sends a (connectionless) message from a send endpoint to a receive endpoint. It is a non-blocking function, and returns immediately. `send_endpoint`, is a local endpoint identifying the send endpoint, `receive_endpoint` identifies a receive endpoint. `buffer` is the application provided buffer, `buffer_size` is the buffer size in bytes. `priority` determines the message priority with a value of 0 being the highest priority and `request` is the identifier used to determine if the send operation has completed on the sending endpoint and the buffer can be reused by the application. Furthermore, this method will abandon the send and return `MCAPI_ERR_MEM_LIMIT` if sufficient memory space is not available. This function cannot be used to send a message to a connected endpoint. Implementations may chose to prevent messages from being sent to connected endpoint or to leave it up to the application to manage this. Functionality for this may be added in a future version of MCAPI in it is therefore recommended that implementations preventing messages from being sent to connected endpoint use `MCAPI_ERR_GENERAL` to report an error. The behavior should be documented.

RETURN VALUE

On success, `*mcapi_status` is set to `MCAPI_SUCCESS` if completed and `MCAPI_PENDING` if not yet completed. On error, `*mcapi_status` is set to the appropriate error defined below.

ERRORS

<code>MCAPI_ERR_NODE_NOTINIT</code>	The node is not initialized.
<code>MCAPI_ERR_ENDP_INVALID</code>	One or both endpoints are invalid.
<code>MCAPI_ERR_MSG_SIZE</code>	The message size exceeds the maximum size allowed by the MCAPI implementation.
<code>MCAPI_ERR_REQUEST_LIMIT</code>	No more request handles available.

MCAPI_ERR_MEM_LIMIT	No memory available.
MCAPI_ERR_PRIORITY	Incorrect priority level.MCAPI_ERR_TRANSMISSION Transmission failure. This error code is optional, and if supported by an implementation, its functionality shall be described.
MCAPI_ERR_PARAMETER NULL and buffer_size > 0) parameter.	Incorrect request or buffer (applies if buffer =
MCAPI_ERR_GENERAL	Implementation specific error not covered by other status codes. Specifics must be documented.

NOTE

Use the `mcapi_test()`, `mcapi_wait()` and `mcapi_wait_any()` functions to query the status of and `mcapi_cancel()` function to cancel the operation.

4.3.2 MCAPI_MSG_SEND

NAME

`mcapi_msg_send` – sends a (connectionless) message from a send endpoint to a receive endpoint.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_msg_send(
    MCAPI_IN mcapi_endpoint_t  send_endpoint,
    MCAPI_IN mcapi_endpoint_t  receive_endpoint,
    MCAPI_IN void*  buffer,
    MCAPI_IN size_t  buffer_size,
    MCAPI_IN mcapi_priority_t  priority,
    MCAPI_OUT mcapi_status_t*  mcapi_status
);
```

DESCRIPTION

Sends a (connectionless) message from a send endpoint to a receive endpoint. It is a blocking function, and returns once the buffer can be reused by the application. `send_endpoint` is a local endpoint identifying the send endpoint and `receive_endpoint` identifies a receive endpoint. `buffer` is the application provided buffer and `buffer_size` is the buffer size in bytes. `priority` determines the message priority with a value of 0 being the highest priority. This method will block if there is insufficient memory space available. When sufficient space becomes available, the function will complete. This function cannot be used allowed to send a message to a connected endpoint. Implementations may chose to prevent messages from being sent to connected endpoint or to leave it up to the application to manage this. Functionality for this may be added in a future version of MCAPI in it is therefore recommended that implementations preventing messages from being sent to connected endpoint use `MCAPI_ERR_GENERAL` to report an error. The behavior should be documented.

RETURN VALUE

On success, `*mcapi_status` is set to `MCAPI_SUCCESS`. On error, `*mcapi_status` is set to the appropriate error defined below. Success means that the entire buffer has been sent.

ERRORS

<code>MCAPI_ERR_NODE_NOTINIT</code>	The node is not initialized.
<code>MCAPI_ERR_ENDP_INVALID</code>	One or both endpoints are invalid.
<code>MCAPI_ERR_MSG_SIZE</code>	The message size exceeds the maximum size allowed by the MCAPI implementation.
<code>MCAPI_ERR_PRIORITY</code>	Incorrect priority level.

MCAPI_ERR_TRANSMISSION	Transmission failure. This error code is optional, and if supported by an implementation, its functionality shall be described.
MCAPI_ERR_PARAMETER	Incorrect buffer (applies if buffer = NULL and buffer_size > 0) parameter.
MCAPI_TIMEOUT	The operation timed out. Implementations can optionally support timeout for this function. The timeout value is set with endpoint attributes.
MCAPI_ERR_GENERAL	Implementation specific error not covered by other status codes. Specifics must be documented.

4.3.3 MCAPI_MSG_RECV_I

NAME

`mcapi_msg_rcv_i` – receives a (connectionless) message from a receive endpoint.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_msg_rcv_i(
    MCAPI_IN mcapi_endpoint_t  receive_endpoint,
    MCAPI_OUT void*  buffer,
    MCAPI_IN size_t  buffer_size,
    MCAPI_OUT mcapi_request_t* request,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Receives a (connectionless) message from a receive endpoint. It is a non-blocking function, and returns immediately. `receive_endpoint` is a local endpoint identifying the receive endpoint. `buffer` is the application provided buffer, and `buffer_size` is the buffer size in bytes. `request` is the identifier used to determine if the receive operation has completed (all the data is in the buffer). Furthermore, this method will abandon the receive and return `MCAPI_ERR_MEM_LIMIT` if the system cannot either wait for sufficient memory to become available or allocate enough memory.

RETURN VALUE

On success, `*mcapi_status` is set to `MCAPI_SUCCESS` if completed and `MCAPI_PENDING` if not yet completed. On error, `*mcapi_status` is set to the appropriate error defined below.

ERRORS

<code>MCAPI_ERR_NODE_NOTINIT</code>	The node is not initialized.
<code>MCAPI_ERR_ENDP_INVALID</code>	Argument is not a valid local endpoint descriptor.
<code>MCAPI_ERR_MSG_TRUNCATED</code>	The message size exceeds the <code>buffer_size</code> .
<code>MCAPI_ERR_TRANSMISSION</code>	Transmission failure. This error code is optional, and if supported by an implementation, its functionality shall be described.
<code>MCAPI_ERR_REQUEST_LIMIT</code>	No more request handles available.
<code>MCAPI_ERR_MEM_LIMIT</code>	No memory available.
<code>MCAPI_ERR_PARAMETER</code>	Incorrect <code>buffer</code> and/or <code>request</code> parameter.

MCAPI_ERR_GENERAL

Implementation specific error not covered by other status codes. Specifics must be documented.

NOTE

Use the `mcapi_test()`, `mcapi_wait()` and `mcapi_wait_any()` functions to query the status of and `mcapi_cancel()` function to cancel the operation.

4.3.4 MCAPI_MSG_RECV

NAME

`mcapi_msg_recv` – receives a (connectionless) message from a receive endpoint.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_msg_recv(
    MCAPI_IN mcapi_endpoint_t  receive_endpoint,
    MCAPI_OUT void*  buffer,
    MCAPI_IN size_t  buffer_size,
    MCAPI_OUT size_t* received_size,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Receives a (connectionless) message from a receive endpoint. It is a blocking function, and returns once a message is available and the received data filled into the buffer. `receive_endpoint` is a local endpoint identifying the receive endpoint. `buffer` is the application provided buffer, and `buffer_size` is the buffer size in bytes. The `received_size` parameter is filled with the actual size of the received message. This method will block if there is insufficient memory space available. When sufficient space becomes available, the function will complete.

RETURN VALUE

On success, `*mcapi_status` is set to `MCAPI_SUCCESS`. On error, `*mcapi_status` is set to the appropriate error defined below.

ERRORS

<code>MCAPI_ERR_NODE_NOTINIT</code>	The node is not initialized.
<code>MCAPI_ERR_ENDP_INVALID</code>	Argument is not a valid local endpoint descriptor.
<code>MCAPI_ERR_MSG_TRUNCATED</code>	The message size exceeds the <code>buffer_size</code> .
<code>MCAPI_ERR_TRANSMISSION</code>	Transmission failure. This error code is optional, and if supported by an implementation, its functionality shall be described.
<code>MCAPI_ERR_PARAMETER</code>	Incorrect <code>buffer</code> and/or <code>received_size</code> parameter.
<code>MCAPI_TIMEOUT</code>	The operation timed out. Implementations can optionally support timeout for this function. The timeout value is set with endpoint attributes.

MCAPI_ERR_GENERAL

Implementation specific error not covered by other status codes. Specifics must be documented.

4.3.5 MCAPI_MSG_AVAILABLE

NAME

`mcapi_msg_available` – checks if messages are available on a receive endpoint.

SYNOPSIS

```
#include <mcapi.h>

mcapi_uint_t mcapi_msg_available(
    MCAPI_IN mcapi_endpoint_t receive_endpoint,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Checks if messages are available on a local receive endpoint. The function returns in a timely fashion. The number of “available” incoming messages is defined as the number of `mcapi_msg_rcv()` operations that are guaranteed to not block waiting for incoming data. `receive_endpoint` is a local identifier for the receive endpoint. The call only checks the availability of messages and does not dequeue them. `mcapi_msg_available()` can only be used to check availability on endpoints on the node local to the caller.

RETURN VALUE

On success, the number of available messages is returned and `*mcapi_status` is set to `MCAPI_SUCCESS`. On error, `MCAPI_NULL` (or 0) is returned and `*mcapi_status` is set to the appropriate error defined below. `MCAPI_NULL` (or 0) could be a valid number of available messages, so status has to be checked to ensure correctness.

ERRORS

<code>MCAPI_ERR_NODE_NOTINIT</code>	The node is not initialized.
<code>MCAPI_ERR_ENDP_INVALID</code>	Argument is not a valid local endpoint descriptor.
<code>MCAPI_ERR_GENERAL</code>	Implementation specific error not covered by other status codes. Specifics must be documented.

NOTE

The status code must be checked to distinguish between no messages and an error condition.

4.4 MCAP Packet Channels

MCAP packet channels transfer data packets between a pair of connected endpoints. A connection between two endpoints is established via a two-phase process. First, some node in the system calls `mcapi_pktchan_connect_i()` to define a connection between two endpoints. This function returns immediately. In the second phase, both sender and receiver open their end of the channel by invoking `mcapi_pktchan_send_open_i()` and `mcapi_pktchan_recv_open_i()`, respectively. The connection is synchronized when both the sender and receiver open functions have completed. In order to avoid deadlock situations, the open functions are non-blocking.

This two-phased binding approach has several important benefits. The "connect" call can be made by any node in the system, which allows the programmer to define the entire channel topology in a single piece of code. This code could even be auto-generated by some static connection tool. This makes it easy to change the channel topology without having to modify multiple source files. This approach also allows the sender and receiver to do their work without any knowledge of what remote nodes they are connected to. This allows for better modularity and application scaling.

Packet channels provide a "system-specified buffer" interface. The programmer specifies the address of a buffer of data to be sent on the send side, but the receiver's "recv" function returns a buffer of data at an address chosen by the system. This is different from the "user-specified buffer" interface use by MCAP messaging – with messages the programmer chooses the buffer in which data is received, and with packet channels the system chooses the buffer.

It is not allowed to send messages to connected endpoints.

4.4.1 MCAPI_PKTCHAN_CONNECT_I

NAME

`mcapi_pktchan_connect_i` – Connects send and receive side endpoints.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_pktchan_connect_i(
    MCAPI_IN mcapi_endpoint_t  send_endpoint,
    MCAPI_IN mcapi_endpoint_t  receive_endpoint,
    MCAPI_OUT mcapi_request_t* request,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Connects a pair of endpoints into a unidirectional FIFO channel. The connect operation can be performed by the sender, the receiver, or by a third party. The connect can happen at the start of the program, or dynamically at run time.

Connect is a non-blocking function. Synchronization to ensure the channel has been created is provided by the open call discussed later.

Channel type endpoints can be connected. A message type endpoint with a get reference count of 0 can also be connected, in which case its type will change to channel for the duration of the connection. If the endpoint get reference count is 0 its reference count will be raised to 1 for the duration of the connection and if so reset to 0 upon close. Attempts to connect message type endpoints with get reference counts > 0 will result in an error.

Attempts to make multiple connections to a single endpoint will be detected as errors. The type of channel connected to an endpoint must match the type of open call invoked by that endpoint; the open function will return an error if the opened channel type does not match the connected channel type, or direction.

It is an error to attempt a connection between endpoints whose attributes are set in an incompatible way (whether attributes are compatible or not is implementation defined). It is also an error to attempt to change the attributes of endpoints that are connected. It is an error to connect an endpoint with itself. A previously connected endpoint can't be connected until the previous channel is disconnected (implies both sides closed).

RETURN VALUE

On success `*mcapi_status` is set to `MCAPI_SUCCESS` if completed and `MCAPI_PENDING` if not yet completed. On error, `*mcapi_status` is set to the appropriate error defined below.

ERRORS

<code>MCAPI_ERR_NODE_NOTINIT</code>	The node is not initialized.
-------------------------------------	------------------------------

MCAPI_ERR_ENDP_INVALID	Argument is not a valid endpoint descriptor, or both endpoints are the same.
MCAPI_ERR_CHAN_CONNECTED	A channel connection has already been established for one or both of the specified endpoints.
MCAPI_ERR_CHAN_CLOSEPENDING	A close request is pending.
MCAPI_ERR_REQUEST_LIMIT	No more request handles available.
MCAPI_ERR_ENDP_GET_LIMIT	The endpoint get reference count is too high.
MCAPI_ERR_ATTR_INCOMPATIBLE	Connection of endpoints with incompatible attributes not allowed.
MCAPI_ERR_PARAMETER	Incorrect request parameter.
MCAPI_ERR_GENERAL	Implementation specific error not covered by other status codes. Specifics must be documented.

NOTE

Use the `mcapi_test()`, `mcapi_wait()` and `mcapi_wait_any()` functions to query the status and `mcapi_cancel()` function to cancel the operation.

4.4.2 MCAPI_PKTCHAN_RECV_OPEN_I

NAME

mcapi_pktchan_recv_open_i – Creates a typed and directional, local representation of the channel. It also provides synchronization for channel creation between two endpoints. Opens are required on both receive and send endpoints.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_pktchan_recv_open_i(
    MCAPI_OUT mcapi_pktchan_recv_hdl_t* receive_handle,
    MCAPI_IN mcapi_endpoint_t receive_endpoint,
    MCAPI_OUT mcapi_request_t* request,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Opens the receive end of a packet channel. The corresponding calls are required on both sides for synchronization to ensure that the channel has been created. It is a non-blocking function, and the `receive_handle` is filled in upon successful completion. No specific ordering of calls between sender and receiver is required since the call is non-blocking. `receive_endpoint` is the local endpoint associated with the channel. The open call returns a typed, local handle for the connected channel that is used for channel receive operations. An endpoint with a previously open channel can't be opened until the previous channel is disconnected (implies both sides closed).

RETURN VALUE

On success, meaning that both sides of the channel are successfully open, `*mcapi_status` is set to `MCAPI_SUCCESS` if completed and `MCAPI_PENDING` if not yet completed. On error `*mcapi_status` is set to the appropriate error defined below.

ERRORS

MCAPI_ERR_NODE_NOTINIT	The node is not initialized.
MCAPI_ERR_ENDP_INVALID	Argument is not a valid local endpoint descriptor.
MCAPI_ERR_CHAN_TYPE	Attempt to open a packet channel on an endpoint that has been connected with a different channel type.
MCAPI_ERR_CHAN_DIRECTION	Attempt to open a receive handle on an endpoint that was connected as a sender.
MCAPI_ERR_REQUEST_LIMIT	No more request handles available.

MCAPI_ERR_CHAN_OPENPENDING	An open request is pending.
MCAPI_ERR_CHAN_CLOSEPENDING	A close request is pending.
MCAPI_ERR_CHAN_OPEN	The channel is already open.
MCAPI_ERR_ENDP_DELETED	The endpoint has been deleted.
MCAPI_ERR_PARAMETER	Incorrect handle or <code>request</code> parameter.
MCAPI_ERR_GENERAL	Implementation specific error not covered by other status codes. Specifics must be documented.

NOTE

Use the `mcapi_test()`, `mcapi_wait()` and `mcapi_wait_any()` functions to query the status and `mcapi_cancel()` function to cancel the operation.

4.4.3 MCAPI_PKTCHAN_SEND_OPEN_I

NAME

`mcapi_pktchan_send_open_i` – Creates a typed and directional, local representation of the channel. It also provides synchronization for channel creation between two endpoints. Opens are required on both receive and send endpoints.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_pktchan_send_open_i(
    MCAPI_OUT mcapi_pktchan_send_hdl_t* send_handle,
    MCAPI_IN mcapi_endpoint_t send_endpoint,
    MCAPI_OUT mcapi_request_t* request,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Opens the send end of a packet channel. The corresponding calls are required on both sides for synchronization to ensure that the channel has been created. It is a non-blocking function, and the `send_handle` is filled in upon successful completion. No specific ordering of calls between sender and receiver is required since the call is non-blocking. `send_endpoint` is the local endpoint associated with the channel. The open call returns a typed, local handle for the connected endpoint that is used by channel send operations. An endpoint with a previously open channel can't be opened until the previous channel is disconnected (implies both sides closed).

RETURN VALUE

On success, meaning that both sides of the channel are successfully open, `*mcapi_status` is set to `MCAPI_SUCCESS` if completed and `MCAPI_PENDING` if not yet completed. On error, `*mcapi_status` is set to the appropriate error defined below.

ERRORS

<code>MCAPI_ERR_NODE_NOTINIT</code>	The node is not initialized.
<code>MCAPI_ERR_ENDP_INVALID</code>	Argument is not a valid local endpoint descriptor.
<code>MCAPI_ERR_CHAN_TYPE</code>	Attempt to open a packet channel on an endpoint that has been connected with a different channel type.
<code>MCAPI_ERR_CHAN_DIRECTION</code>	Attempt to open a send handle on a port that was connected as a receiver.

MCAPI_ERR_REQUEST_LIMIT	No more request handles available.
MCAPI_ERR_CHAN_OPENPENDING	An open request is pending.
MCAPI_ERR_CHAN_CLOSEPENDING	A close request is pending.
MCAPI_ERR_CHAN_OPEN	The channel is already open.
MCAPI_ERR_ENDP_DELETED	The endpoint has been deleted.
MCAPI_ERR_PARAMETER	Incorrect handle or <code>request</code> parameter.
MCAPI_ERR_GENERAL	Implementation specific error not covered by other status codes. Specifics must be documented.

NOTE

Use the `mcapi_test()`, `mcapi_wait()` and `mcapi_wait_any()` functions to query the status and `mcapi_cancel()` function to cancel the operation.

4.4.4 MCAPI_PKTCHAN_SEND_I

NAME

`mcapi_pktchan_send_i` – sends a (connected) packet on a channel.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_pktchan_send_i(
    MCAPI_IN mcapi_pktchan_send_hdl_t send_handle,
    MCAPI_IN void* buffer,
    MCAPI_IN size_t size,
    MCAPI_OUT mcapi_request_t* request,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Sends a packet on a connected channel. It is a non-blocking function, and returns immediately. `buffer` is the application provided buffer and `size` is the buffer size. `request` is the identifier used to determine if the send operation has completed on the sending endpoint and the buffer can be reused. While this method returns immediately, data transfer will not complete until the packet has been transmitted. The definition of transmission in this context is implementation defined and may include blocking until the send buffer is available for reuse. Alternatively the buffer's availability for reuse can be tested with the `mcapi_pktchan_release_test()` function. The behavior must be documented by the implementation. Furthermore, this method will abandon the send and return `MCAPI_ERR_MEM_LIMIT` if the system cannot either wait for sufficient memory to become available or allocate enough memory.

RETURN VALUE

On success, `*mcapi_status` is set to `MCAPI_SUCCESS` if completed and `MCAPI_PENDING` if not yet completed. On error, `*mcapi_status` is set to the appropriate error defined below.

ERRORS

<code>MCAPI_ERR_NODE_NOTINIT</code>	The node is not initialized.
<code>MCAPI_ERR_CHAN_INVALID</code>	Argument is not a valid channel handle.
<code>MCAPI_ERR_PKT_SIZE</code>	The packet size exceeds the maximum size allowed by the MCAPI implementation.
<code>MCAPI_ERR_REQUEST_LIMIT</code>	No more request handles available.
<code>MCAPI_ERR_MEM_LIMIT</code>	No memory available.
<code>MCAPI_ERR_TRANSMISSION</code>	Transmission failure. This error code is optional, and if supported by an implementation, its functionality shall be described.

MCAPI_ERR_PARAMETER

Incorrect `request` or `buffer` (applies if `buffer = 0` and `buffer_size > 0`) parameter.

MCAPI_ERR_GENERAL

Implementation specific error not covered by other status codes. Specifics must be documented.

NOTE

Use the `mcapi_test()`, `mcapi_wait()` and `mcapi_wait_any()` functions to query the status and `mcapi_cancel()` function to cancel the operation.

4.4.5 MCAPI_PKTCHAN_SEND

NAME

`mcapi_pktchan_send` – sends a (connected) packet on a channel.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_pktchan_send(
    MCAPI_IN mcapi_pktchan_send_hdl_t send_handle,
    MCAPI_IN void* buffer,
    MCAPI_IN size_t size,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Sends a packet on a connected channel. This method will block until the packet has been transmitted. The definition of transmission in this context is implementation defined and may include blocking until the send buffer is available for reuse. Alternatively the buffer's availability for reuse can be tested with the `mcapi_pktchan_release_test()` function. The behavior must be documented by the implementation. `send_handle` is the local send handle which represents the send endpoint associated with the channel. `buffer` is the application provided buffer and `size` is the buffer size. Since channels behave like FIFOs, by default this method will block if the packet can't be transmitted because of lack of memory space. When sufficient space becomes available, the function will complete. By default this method will block if there is insufficient memory space available. When sufficient space becomes available, the function will complete.

RETURN VALUE

On success, `*mcapi_status` is set to `MCAPI_SUCCESS`. On error, `*mcapi_status` is set to the appropriate error defined below. Success means that the entire buffer has been sent.

ERRORS

<code>MCAPI_ERR_NODE_NOTINIT</code>	The node is not initialized.
<code>MCAPI_ERR_CHAN_INVALID</code>	Argument is not a valid channel handle.
<code>MCAPI_ERR_PKT_SIZE</code>	The message size exceeds the maximum size allowed by the MCAPI implementation.
<code>MCAPI_ERR_TRANSMISSION</code>	Transmission failure. This error code is optional, and if supported by an implementation, its functionality shall be described.
<code>MCAPI_ERR_PARAMETER</code>	Incorrect buffer (applies if <code>buffer = 0</code> and <code>buffer_size > 0</code>) parameter.

MCAPI_TIMEOUT

The operation timed out. Implementations can optionally support timeout for this function. The timeout value is set with endpoint attributes.

MCAPI_ERR_GENERAL

Implementation specific error not covered by other status codes. Specifics must be documented.

4.4.6 MCAPI_PKTCHAN_RECV_I

NAME

mcapi_pktchan_recv_i – receives a (connected) packet on a channel.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_pktchan_recv_i(
    MCAPI_IN mcapi_pktchan_recv_hndl_t receive_handle,
    MCAPI_OUT void** buffer,
    MCAPI_OUT mcapi_request_t* request,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Receives a packet on a connected channel. It is a non-blocking function, and returns immediately. `receive_handle` is the local representation of the handle used to receive packets. When the receive operation completes, the `buffer` parameter is filled with the address of a system-supplied buffer containing the received packet. After the receive request has completed and the application is finished with `buffer`, `buffer` must be returned to the system by calling `mcapi_pktchan_release()`. `request` is the identifier used to determine if the receive operation has completed and `buffer` is ready for use; the `mcapi_test()`, `mcapi_wait()` or `mcapi_wait_any()` function will return the actual size of the received packet. Furthermore, this method will abandon the receive and return `MCAPI_ERR_MEM_LIMIT` if sufficient memory space is not available.

RETURN VALUE

On success, `*mcapi_status` is set to `MCAPI_SUCCESS` if completed and `MCAPI_PENDING` if not yet completed. On error, `*mcapi_status` is set to the appropriate error defined below.

ERRORS

MCAPI_ERR_NODE_NOTINIT	The node is not initialized.
MCAPI_ERR_CHAN_INVALID	Argument is not a valid channel handle.
MCAPI_ERR_MEM_LIMIT	No memory available
MCAPI_ERR_REQUEST_LIMIT	No more request handles available.
MCAPI_ERR_TRANSMISSION	Transmission failure. This error code is optional, and if supported by an implementation, its functionality shall be described.
MCAPI_ERR_PARAMETER	Incorrect <code>buffer</code> or <code>request</code> parameter.

MCAPI_ERR_GENERAL

Implementation specific error not covered by other status codes. Specifics must be documented.

NOTE

Use the `mcapi_test()`, `mcapi_wait()` and `mcapi_wait_any()` functions to query the status of and `mcapi_cancel()` function to cancel the operation.

4.4.7 MCAPI_PKTCHAN_RECV

NAME

`mcapi_pktchan_recv` – receives a data packet on a (connected) channel.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_pktchan_recv(
    MCAPI_IN mcapi_pktchan_recv_hndl_t receive_handle,
    MCAPI_OUT void** buffer,
    MCAPI_OUT size_t* received_size,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Receives a packet on a connected channel. It is a blocking function, and returns when the data has been written to the buffer. `receive_handle` is the local representation of the receive endpoint associated with the channel. When the receive operation completes, the `buffer` parameter is filled with the address of a system-supplied buffer containing the received packet. and `received_size` is filled with the size of the packet in that buffer. When the application finishes with buffer, it must return it to the system by calling `mcapi_pktchan_release()`. By default this method will block if there is insufficient memory space available. When sufficient space becomes available, the function will complete.

RETURN VALUE

On success, `*mcapi_status` is set to `MCAPI_SUCCESS`. On error, `*mcapi_status` is set to the appropriate error defined below.

ERRORS

<code>MCAPI_ERR_NODE_NOTINIT</code>	The node is not initialized.
<code>MCAPI_ERR_CHAN_INVALID</code>	Argument is not a valid channel handle.
<code>MCAPI_ERR_TRANSMISSION</code>	Transmission failure. This error code is optional, and if supported by an implementation, it's functionality shall be described.
<code>MCAPI_ERR_PARAMETER</code>	Incorrect <code>buffer</code> or <code>received_size</code> parameter.
<code>MCAPI_TIMEOUT</code>	The operation timed out. Implementations can optionally support timeout for this function. The timeout value is set with endpoint attributes.

MCAPI_ERR_GENERAL

Implementation specific error not covered by other status codes. Specifics must be documented.

4.4.8 MCAPI_PKTCHAN_AVAILABLE

NAME

`mcapi_pktchan_available` – checks if packets are available on a receive endpoint.

SYNOPSIS

```
#include <mcapi.h>

mcapi_uint_t mcapi_pktchan_available(
    MCAPI_IN mcapi_pktchan_rcv_hdl_t receive_handle,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Checks if packets are available on a receive endpoint. This function returns in a timely fashion. The number of available packets is defined as the number of receive operations that could be performed without blocking to wait for incoming data. `receive_handle` is the local handle for the packet channel. The call only checks the availability of packets and does not dequeue them.

RETURN VALUE

On success, the number of available packets are returned and `*mcapi_status` is set to `MCAPI_SUCCESS`. On error, `MCAPI_NULL` (or 0) is returned and `*mcapi_status` is set to the appropriate error defined below. `MCAPI_NULL` (or 0) could be a valid number of available packets, so status has to be checked to ensure correctness.

ERRORS

<code>MCAPI_ERR_NODE_NOTINIT</code>	The node is not initialized.
<code>MCAPI_ERR_CHAN_INVALID</code>	Argument is not a channel handle.
<code>MCAPI_ERR_GENERAL</code>	Implementation specific error not covered by other status codes. Specifics must be documented.

NOTE

The status code must be checked to distinguish between no messages and an error condition.

4.4.9 MCAPI_PKTCHAN_RELEASE

NAME

`mcapi_pktchan_release` – releases a packet buffer obtained from a `mcapi_pktchan_rcv()` call.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_pktchan_release(
    MCAPI_IN void* buffer,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

When a user is finished with a packet buffer obtained from `mcapi_pktchan_rcv_i()` or `mcapi_pktchan_rcv()`, they must invoke this function to return the buffer to the system. Buffers can be released in any order. This function is guaranteed to return in a timely fashion. Buffers can be released by the receiving node only.

RETURN VALUE

On success `*mcapi_status` is set to `MCAPI_SUCCESS`. On error, `*mcapi_status` is set to the appropriate error defined below.

ERRORS

<code>MCAPI_ERR_NODE_NOTINIT</code>	The node is not initialized.
<code>MCAPI_ERR_BUF_INVALID</code>	Argument is not a valid buffer descriptor.
<code>MCAPI_ERR_GENERAL</code>	Implementation specific error not covered by other status codes. Specifics must be documented.

SEE ALSO

```
mcapi_pktchan_rcv(), mcapi_pktchan_rcv_i()
```

4.4.10 MCAPI_PKTCHAN_RELEASE_TEST

NAME

`mcapi_pktchan_release_test` – Tests if a packet buffer obtained from a `mcapi_pktchan_rcv()` has been released (with `mcapi_pktchan_release`).

SYNOPSIS

```
#include <mcapi.h>

mcapi_boolean_t mcapi_pktchan_release_test(
    MCAPI_IN void* buffer,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Checks if a packet buffer has been released by the receiver. This function is blocking and returns in a timely fashion.

RETURN VALUE

On success, `MCAPI_TRUE` is returned and `*mcapi_status` is set to `MCAPI_SUCCESS`. If the operation has not completed `MCAPI_FALSE` is returned and `*mcapi_status` is set to `MCAPI_PENDING`. On error `MCAPI_FALSE` is returned and `*mcapi_status` is set to the appropriate error defined below.

ERRORS

<code>MCAPI_ERR_NODE_NOTINIT</code>	The node is not initialized.
<code>MCAPI_ERR_BUF_INVALID</code>	Argument is not a valid buffer descriptor.
<code>MCAPI_ERR_PARAMETER</code>	Incorrect <code>buffer</code> or parameter.
<code>MCAPI_ERR_GENERAL</code>	Implementation specific error not covered by other status codes. Specifics must be documented.

SEE ALSO

```
mcapi_pktchan_release()
```

4.4.11 MCAPI_PKTCHAN_RECV_CLOSE_I

NAME

mcapi_pktchan_recv_close_i – Closes the receive side of the channel.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_pktchan_recv_close_i(
    MCAPI_IN mcapi_pktchan_recv_hndl_t receive_handle,
    MCAPI_OUT mcapi_request_t* request,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Closes the receive side of a channel. The sender makes the send-side call and the receiver makes the receive-side call. The corresponding calls are required on both sides to ensure that the channel has been properly closed. It is a non-blocking function, and returns immediately. *receive_handle* is the local representation of the handle used to receive packets. All pending packets are discarded, and any attempt to send more packets will give an error. A packet channel is disconnected when, both sides have issued close calls and the last (second) close operation is performed. If the endpoint type was changed from message to channel and/or the get reference count was increased to 1 by the connect, the type will be reset to message and the get reference count to 0.

RETURN VALUE

On success, meaning that both sides of the channel are successfully closed, **mcapi_status* is set to *MCAPI_SUCCESS* if completed and *MCAPI_PENDING* if not yet completed. On error **mcapi_status* is set to the appropriate error defined below.

ERRORS

MCAPI_ERR_NODE_NOTINIT	The node is not initialized.
MCAPI_ERR_CHAN_INVALID	Argument is not a valid channel handle.
MCAPI_ERR_CHAN_TYPE	Attempt to close a packet channel on an endpoint that has been connected with a different channel type.
MCAPI_ERR_CHAN_DIRECTION	Attempt to close a send handle on a port that was connected as a receiver, or vice versa.
MCAPI_ERR_CHAN_NOTOPEN	The channel is not open.
MCAPI_ERR_REQUEST_LIMIT	No more request handles available.

MCAPI_ERR_PARAMETER

Incorrect request parameter.

MCAPI_ERR_GENERAL

Implementation specific error not covered by other status codes. Specifics must be documented.

NOTE

Use the `mcapi_test()`, `mcapi_wait()` and `mcapi_wait_any()` functions to query the status of and `mcapi_cancel()` function to cancel the operation.

4.4.12 MCAPI_PKTCHAN_SEND_CLOSE_I

NAME

`mcapi_pktchan_send_close_i` – closes channel on a send endpoint.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_pktchan_send_close_i(
    MCAPI_IN mcapi_pktchan_send_hndl_t send_handle,
    MCAPI_OUT mcapi_request_t* request,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Closes the send side of a channel. The sender makes the send-side call and the receiver makes the receive-side call. The corresponding calls are required on both sides to ensure that the channel has been properly closed. It is a non-blocking function, and returns immediately. `send_handle` is the local representation of the handle used to send packets. Pending packets at the receiver are not discarded. A packet channel is disconnected when, both sides have issued close calls and the last (second) close operation is performed. If the endpoint type was changed from message to channel and/or the get reference count was increased to 1 by the connect, the type will be reset to message and the get reference count to 0.

RETURN VALUE

On success, meaning that both sides of the channel are successfully closed, `*mcapi_status` is set to `MCAPI_SUCCESS` if completed and `MCAPI_PENDING` if not yet completed. On error `*mcapi_status` is set to the appropriate error defined below.

ERRORS

<code>MCAPI_ERR_NODE_NOTINIT</code>	The node is not initialized.
<code>MCAPI_ERR_CHAN_INVALID</code>	Argument is not a channel handle.
<code>MCAPI_ERR_CHAN_TYPE</code>	Attempt to close a packet channel on an endpoint that has been connected with a different channel type.
<code>MCAPI_ERR_CHAN_DIRECTION</code>	Attempt to close a send handle on a port that was connected as a receiver, or vice versa.
<code>MCAPI_ERR_CHAN_NOTOPEN</code>	The channel is not open.
<code>MCAPI_ERR_REQUEST_LIMIT</code>	No more request handles available.

MCAPI_ERR_PARAMETER

Incorrect request parameter.

MCAPI_ERR_GENERAL

Implementation specific error not covered by other status codes. Specifics must be documented.

NOTE

Use the `mcapi_test()`, `mcapi_wait()` and `mcapi_wait_any()` functions to query the status of and `mcapi_cancel()` function to cancel the operation.

4.5 MCAPI Scalar Channels

MCAPI scalar channels are used to transfer 8-bit, 16-bit, 32-bit and 64-bit scalars on a connected channel. The connection process for scalar channels uses the same two-phase mechanism as packet channels. See the packet channel section for a detailed description of the connection process.

The MCAPI scalar channels API provides only blocking send and receive methods. Scalar channels are intended to provide a very low overhead interface for moving a stream of values. In fact, some embedded systems may be able to implement a scalar channel as a hardware FIFO. The sort of streaming algorithms that take advantage of scalar channels should not require a non-blocking send or receive method; each process should simply receive a value to work on, do its work, send the result out on a channel, and repeat. Applications that require non-blocking semantics should use packet channels instead of scalar channels. The scalar functions only support communication of same size scalars on both sides, i.e. an 8-bit send must be read by an 8-bit receive, etc.

4.5.1 MCAPI_SCLCHAN_CONNECT_I

NAME

`mcapi_sclchan_connect_i` – connects a pair of scalar channel endpoints.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_sclchan_connect_i(
    MCAPI_IN mcapi_endpoint_t send_endpoint,
    MCAPI_IN mcapi_endpoint_t receive_endpoint,
    MCAPI_OUT mcapi_request_t* request,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Connects a pair of endpoints into a unidirectional FIFO channel. The connect operation can be performed by the sender, the receiver, or by a third party. The connect can happen once at the start of the program or dynamically at run time.

`mcapi_sclchan_connect_i()` is a non-blocking function. Synchronization to ensure the channel has been created is provided by the open call discussed later.

Note that this function behaves like the packet channel connect call.

Channel type endpoints can be connected. A message type endpoint with a get reference count of 0 can also be connected, in which case its type will change to channel for the duration of the connection. If the endpoint get reference count is 0 its reference count will be raised to 1 for the duration of the connection and if so reset to 0 upon close. Attempts to connect message type endpoints with get reference counts > 0 will result in an error.

Attempts to make multiple connections to a single endpoint will be detected as errors. The type of channel connected to an endpoint must match the type of open call invoked by that endpoint; the open function will return an error if the opened channel type does not match the connected channel type, or direction.

It is an error to attempt a connection between endpoints whose attributes are set in an incompatible way (whether attributes are compatible or not is implementation defined). It is also an error to attempt to change the attributes of endpoints that are connected. It is an error to connect an endpoint with itself. A previously connected endpoint can't be connected until the previous channels is disconnected (implies both sides closed).

RETURN VALUE

On success, `*mcapi_status` is set to `MCAPI_SUCCESS` if completed and `MCAPI_PENDING` if not yet completed. On error `*mcapi_status` is set to the appropriate error defined below.

ERRORS

MCAPI_ERR_NODE_NOTINIT	The node is not initialized.
MCAPI_ERR_ENDP_INVALID	Argument is not a valid endpoint descriptor, or both endpoints are the same.
MCAPI_ERR_CHAN_CONNECTED	A channel connection has already been established for one or both of the specified endpoints.
MCAPI_ERR_CHAN_CLOSEPENDING	A close request is pending.
MCAPI_ERR_ATTR_INCOMPATIBLE	Connection of endpoints with incompatible attributes not allowed.
MCAPI_ERR_REQUEST_LIMIT	No more request handles available.
MCAPI_ERR_ENDP_GET_LIMIT	The endpoint get reference count is too high.
MCAPI_ERR_PARAMETER	Incorrect request parameter.
MCAPI_ERR_GENERAL	Implementation specific error not covered by other status codes. Specifics must be documented.

NOTE

Use the `mcapi_test()`, `mcapi_wait()` and `mcapi_wait_any()` functions to query the status of and `mcapi_cancel()` function to cancel the operation.

4.5.2 MCAPI_SCLCHAN_RECV_OPEN_I

NAME

`mcapi_sclchan_recv_open_i` – Creates a typed, local representation of a scalar channel.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_sclchan_recv_open_i(
    MCAPI_OUT mcapi_sclchan_recv_hndl_t* receive_handle,
    MCAPI_IN mcapi_endpoint_t receive_endpoint,
    MCAPI_OUT mcapi_request_t* request,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Opens the receive end of a scalar channel. It also provides synchronization for channel creation between two endpoints. The corresponding calls are required on both sides to synchronize the endpoints. It is a non-blocking function, and the `recv_handle` is filled in upon successful completion. No specific ordering of calls between sender and receiver is required since the call is non-blocking. `receive_endpoint` is the local endpoint identifier. The call returns a local handle for the connected channel. An endpoint with a previously open channel can't be opened until the previous channel is disconnected (implies both sides closed).

RETURN VALUE

On success, meaning that both sides of the channel are successfully open, `*mcapi_status` is set to `MCAPI_SUCCESS` if completed and `MCAPI_PENDING` if not yet completed. On error, `*mcapi_status` is set to the appropriate error defined below.

ERRORS

<code>MCAPI_ERR_NODE_NOTINIT</code>	The node is not initialized.
<code>MCAPI_ERR_ENDP_INVALID</code>	Argument is not an endpoint descriptor. A remote endpoint is also invalid for this function.
<code>MCAPI_ERR_CHAN_TYPE</code>	Attempt to open a packet channel on an endpoint that has been connected with a different channel type.
<code>MCAPI_ERR_CHAN_DIRECTION</code>	Attempt to open a send handle on a port that was connected as a receiver, or vice versa.
<code>MCAPI_ERR_CHAN_OPENPENDING</code>	An open request is pending.

<code>MCAPI_ERR_CHAN_CLOSEPENDING</code>	A close request is pending.
<code>MCAPI_ERR_CHAN_OPEN</code>	The channel is already open.
<code>MCAPI_ERR_REQUEST_LIMIT</code>	No more request handles available.
<code>MCAPI_ERR_ENDP_DELETED</code>	The endpoint has been deleted.
<code>MCAPI_ERR_PARAMETER</code>	Incorrect handle or <code>request</code> parameter.
<code>MCAPI_ERR_GENERAL</code>	Implementation specific error not covered by other status codes. Specifics must be documented.

NOTE

Use the `mcapi_test()`, `mcapi_wait()` and `mcapi_wait_any()` functions to query the status and `mcapi_cancel()` function to cancel the operation.

4.5.3 MCAPI_SCLCHAN_SEND_OPEN_I

NAME

`mcapi_sclchan_send_open_i` – Creates a typed, local representation of a scalar channel.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_sclchan_send_open_i(
    MCAPI_OUT mcapi_sclchan_send_hndl_t* send_handle,
    MCAPI_IN mcapi_endpoint_t send_endpoint,
    MCAPI_OUT mcapi_request_t* request,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Opens the send end of a scalar channel. . It also provides synchronization for channel creation between two endpoints. The corresponding calls are required on both sides to synchronize the endpoints. It is a non-blocking function, and the `send_handle` is filled in upon successful completion. No specific ordering of calls between sender and receiver is required since the call is non-blocking. `send_endpoint` is the local endpoint identifier. The call returns a local handle for connected channel. An endpoint with a previously open channel can't be opened until the previous channel is disconnected (implies both sides closed).

RETURN VALUE

On success, meaning that both sides of the channel are successfully open, `*mcapi_status` is set to `MCAPI_SUCCESS` if completed and `MCAPI_PENDING` if not yet completed. On error, `*mcapi_status` is set to the appropriate error defined below.

ERRORS

<code>MCAPI_ERR_NODE_NOTINIT</code>	The node is not initialized.
<code>MCAPI_ERR_ENDP_INVALID</code>	Argument is not an endpoint descriptor. A remote endpoint is also invalid for this function.
<code>MCAPI_ERR_CHAN_TYPE</code>	Attempt to open a packet channel on an endpoint that has been connected with a different channel type.
<code>MCAPI_ERR_CHAN_DIRECTION</code>	Attempt to open a send handle on a port that was connected as a receiver, or vice versa.

MCAPI_ERR_CHAN_OPENPENDING	An open request is pending.
MCAPI_ERR_CHAN_CLOSEPENDING	A close request is pending.
MCAPI_ERR_CHAN_OPEN	The channel is already open.
MCAPI_ERR_REQUEST_LIMIT	No more request handles available.
MCAPI_ERR_ENDP_DELETED	The endpoint has been deleted.
MCAPI_ERR_PARAMETER	Incorrect handle or <code>request</code> parameter.
MCAPI_ERR_GENERAL	Implementation specific error not covered by other status codes. Specifics must be documented.

NOTE

Use the `mcapi_test()`, `mcapi_wait()` and `mcapi_wait_any()` functions to query the status and `mcapi_cancel()` function to cancel the operation.

4.5.4 MCAPI_SCLCHAN_SEND_UINT64

NAME

mcapi_sclchan_send_uint64 – sends a (connected) 64-bit scalar on a channel.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_sclchan_send_uint64(
    MCAPI_IN mcapi_sclchan_send_hndl_t send_handle,
    MCAPI_IN mcapi_uint64_t dataword,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Sends a scalar on a connected channel. It is a blocking function, and returns immediately unless the buffer is full. send_handle is the send endpoint identifier. dataword is the scalar. Since channels behave like FIFOs, this method will block if there is no free space in the channel's receive buffer. When sufficient space becomes available (due to receive calls), the function will complete.

RETURN VALUE

On success, *mcapi_status is set to MCAPI_SUCCESS. On error *mcapi_status is set to the appropriate error defined below. Optionally, implementations may choose to always set *mcapi_status to MCAPI_SUCCESS for performance reasons.

ERRORS

MCAPI_ERR_NODE_NOTINIT	The node is not initialized.
MCAPI_ERR_CHAN_INVALID	Argument is not a channel handle.
MCAPI_ERR_GENERAL	Implementation specific error not covered by other status codes. Specifics must be documented.

4.5.5 MCAPI_SCLCHAN_SEND_UINT32

NAME

`mcapi_sclchan_send_uint32` – sends a (connected) 32-bit scalar on a channel.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_sclchan_send_uint32(
    MCAPI_IN mcapi_sclchan_send_hndl_t send_handle,
    MCAPI_IN mcapi_uint32_t dataword,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Sends a scalar on a connected channel. It is a blocking function, and returns immediately unless the buffer is full. `send_handle` is the send endpoint identifier. `dataword` is the scalar. Since channels behave like FIFOs, this method will block if there is no free space in the channel's receive buffer. When sufficient space becomes available (due to receive calls), the function will complete.

RETURN VALUE

On success, `*mcapi_status` is set to `MCAPI_SUCCESS`. On error `*mcapi_status` is set to the appropriate error defined below. Optionally, implementations may choose to always set `*mcapi_status` to `MCAPI_SUCCESS` for performance reasons.

ERRORS

<code>MCAPI_ERR_NODE_NOTINIT</code>	The node is not initialized.
<code>MCAPI_ERR_CHAN_INVALID</code>	Argument is not a channel handle.
<code>MCAPI_ERR_GENERAL</code>	Implementation specific error not covered by other status codes. Specifics must be documented.

4.5.6 MCAPI_SCLCHAN_SEND_UINT16

NAME

`mcapi_sclchan_send_uint16` – sends a (connected) 16-bit scalar on a channel.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_sclchan_send_uint16(
    MCAPI_IN mcapi_sclchan_send_hndl_t send_handle,
    MCAPI_IN mcapi_uint16_t dataword,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Sends a scalar on a connected channel. It is a blocking function, and returns immediately unless the buffer is full. `send_handle` is the send endpoint identifier. `dataword` is the scalar. Since channels behave like FIFOs, this method will block if there is no free space in the channel's receive buffer. When sufficient space becomes available (due to receive calls), the function will complete.

RETURN VALUE

On success, `*mcapi_status` is set to `MCAPI_SUCCESS`. On error `*mcapi_status` is set to the appropriate error defined below. Optionally, implementations may choose to always set `*mcapi_status` to `MCAPI_SUCCESS` for performance reasons.

ERRORS

<code>MCAPI_ERR_NODE_NOTINIT</code>	The node is not initialized.
<code>MCAPI_ERR_CHAN_INVALID</code>	Argument is not a channel handle.
<code>MCAPI_ERR_GENERAL</code>	Implementation specific error not covered by other status codes. Specifics must be documented.

4.5.7 MCAPI_SCLCHAN_SEND_UINT8

NAME

mcapi_sclchan_send_uint8 – sends a (connected) 8-bit scalar on a channel.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_sclchan_send_uint8(
    MCAPI_IN mcapi_sclchan_send_hndl_t send_handle,
    MCAPI_IN mcapi_uint8_t dataword,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Sends a scalar on a connected channel. It is a blocking function, and returns immediately unless the buffer is full. `send_handle` is the send endpoint identifier. `dataword` is the scalar. Since channels behave like FIFOs, this method will block if there is no free space in the channel's receive buffer. When sufficient space becomes available (due to receive calls), the function will complete.

RETURN VALUE

On success, `*mcapi_status` is set to `MCAPI_SUCCESS`. On error `*mcapi_status` is set to the appropriate error defined below. Optionally, implementations may choose to always set `*mcapi_status` to `MCAPI_SUCCESS` for performance reasons.

ERRORS

MCAPI_ERR_NODE_NOTINIT	The node is not initialized.
MCAPI_ERR_CHAN_INVALID	Argument is not a channel handle.
MCAPI_ERR_GENERAL	Implementation specific error not covered by other status codes. Specifics must be documented.

4.5.8 MCAPI_SCLCHAN_RECV_UINT64

NAME

mcapi_sclchan_recv_uint64 – receives a (connected) 64-bit scalar on a channel.

SYNOPSIS

```
#include <mcapi.h>

mcapi_uint64_t mcapi_sclchan_recv_uint64(
    MCAPI_IN mcapi_sclchan_recv_hndl_t receive_handle,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Receives a scalar on a connected channel. It is a blocking function, and returns when a scalar is available. `receive_handle` is the receive endpoint identifier.

RETURN VALUE

On success, a value of type `uint64_t` is returned and `*mcapi_status` is set to `MCAPI_SUCCESS`. On error, the return value is undefined and `*mcapi_status` is set to the appropriate error defined below. Optionally, implementations may choose to always set `*mcapi_status` to `MCAPI_SUCCESS` for performance reasons.

ERRORS

MCAPI_ERR_NODE_NOTINIT	The node is not initialized.
MCAPI_ERR_CHAN_INVALID	Argument is not a channel handle.
MCAPI_ERR_GENERAL	Implementation specific error not covered by other status codes. Specifics must be documented.

NOTE

The receive scalar size must match the send size.

4.5.9 MCAPI_SCLCHAN_RECV_UINT32

NAME

mcapi_sclchan_recv_uint32 – receives a 32-bit scalar on a (connected) channel.

SYNOPSIS

```
#include <mcapi.h>

mcapi_uint32_t mcapi_sclchan_recv_uint32(
    MCAPI_IN mcapi_sclchan_recv_hdl_t receive_handle,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Receives a scalar on a connected channel. It is a blocking function, and returns when a scalar is available. `receive_handle` is the receive endpoint identifier.

RETURN VALUE

On success, a value of type `uint32_t` is returned and `*mcapi_status` is set to `MCAPI_SUCCESS`. On error, the return value is undefined and `*mcapi_status` is set to the appropriate error defined below. Optionally, implementations may choose to always set `*mcapi_status` to `MCAPI_SUCCESS` for performance reasons.

ERRORS

MCAPI_ERR_NODE_NOTINIT	The node is not initialized.
MCAPI_ERR_CHAN_INVALID	Argument is not a channel handle.
MCAPI_ERR_GENERAL	Implementation specific error not covered by other status codes. Specifics must be documented.

NOTE

The receive scalar size must match the send size.

4.5.10 MCAPI_ SCLCHAN_RECV_UINT16

NAME

mcapi_sclchan_recv_uint16 – receives a 16-bit scalar on a (connected) channel.

SYNOPSIS

```
#include <mcapi.h>

mcapi_uint16_t mcapi_sclchan_recv_uint16(
    MCAPI_IN mcapi_sclchan_recv_hdl_t receive_handle,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Receives a scalar on a connected channel. It is a blocking function, and returns when a scalar is available. receive_handle is the receive endpoint identifier.

RETURN VALUE

On success, a value of type uint16_t is returned and *mcapi_status is set to MCAPI_SUCCESS. On error, the return value is undefined and *mcapi_status is set to the appropriate error defined below. Optionally, implementations may choose to always set *mcapi_status to MCAPI_SUCCESS for performance reasons.

ERRORS

MCAPI_ERR_NODE_NOTINIT	The node is not initialized.
MCAPI_ERR_CHAN_INVALID	Argument is not a channel handle.
MCAPI_ERR_GENERAL	Implementation specific error not covered by other status codes. Specifics must be documented.

NOTE

The receive scalar size must match the send size.

4.5.11 MCAPI_SCLCHAN_RECV_UINT8

NAME

mcapi_sclchan_recv_uint8 – receives a (connected) 8-bit scalar on a channel.

SYNOPSIS

```
#include <mcapi.h>

mcapi_uint8_t mcapi_sclchan_recv_uint8(
    MCAPI_IN mcapi_sclchan_recv_hdl_t receive_handle,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Receives a scalar on a connected channel. It is a blocking function, and returns when a scalar is available. receive_handle is the receive endpoint identifier.

RETURN VALUE

On success, a value of type uint8_t is returned and *mcapi_status is set to MCAPI_SUCCESS. On error, the return value is undefined and *mcapi_status is set to the appropriate error defined below. Optionally, implementations may choose to always set *mcapi_status to MCAPI_SUCCESS for performance reasons.

ERRORS

MCAPI_ERR_NODE_NOTINIT	The node is not initialized.
MCAPI_ERR_CHAN_INVALID	Argument is not a channel handle.
MCAPI_ERR_GENERAL	Implementation specific error not covered by other status codes. Specifics must be documented.

NOTE

The receive scalar size must match the send size.

4.5.12 MCAPI_SCLCHAN_AVAILABLE

NAME

`mcapi_sclchan_available` – checks if scalars are available on a receive endpoint.

SYNOPSIS

```
#include <mcapi.h>

mcapi_uint_t mcapi_sclchan_available(
    MCAPI_IN mcapi_sclchan_rcv_hndl_t receive_handle,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Checks if scalars are available on a receive endpoint. The function returns immediately. `receive_endpoint` is the receive endpoint identifier. The call only checks the availability of messages does not dequeue them.

RETURN VALUE

On success, the number of available scalars is returned and `*mcapi_status` is set to `MCAPI_SUCCESS`. On error, `MCAPI_NULL` (or 0) is returned and `*mcapi_status` is set to the appropriate error defined below. `MCAPI_NULL` (or 0) could be a valid number of available scalars, so status has to be checked to ensure correctness.

ERRORS

<code>MCAPI_ERR_NODE_NOTINIT</code>	The node is not initialized.
<code>MCAPI_ERR_CHAN_INVALID</code>	Argument is not a channel handle.
<code>MCAPI_ERR_GENERAL</code>	Implementation specific error not covered by other status codes. Specifics must be documented.

NOTE

The status code must be checked to distinguish between no messages and an error condition.

4.5.13 MCAPI_ SCLCHAN_RECV_CLOSE_I

NAME

mcapi_sclchan_recv_close_i – closes channel on a receive endpoint.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_sclchan_recv_close_i(
    MCAPI_IN mcapi_sclchan_recv_hdl_t receive_handle,
    MCAPI_OUT mcapi_request_t* request,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Closes the receive side of a channel. The corresponding calls are required on both send and receive sides to ensure that the channel is properly closed. It is a non-blocking function, and returns immediately. `receive_handle` is the local representation of the handle used to receive packets. All pending scalars are discarded, and any attempt to send more scalars will give an error. A scalar channel is disconnected when, both sides have issued close calls and the last (second) close operation is performed. If the endpoint type was changed from message to channel and/or the get reference count was increased to 1 by the connect, the type will be reset to message and the get reference count to 0.

RETURN VALUE

On success, meaning that both sides of the channel are successfully closed, `*mcapi_status` is set to `MCAPI_SUCCESS` if completed and `MCAPI_PENDING` if not yet completed. On error `*mcapi_status` is set to the appropriate error defined below.

ERRORS

MCAPI_ERR_NODE_NOTINIT	The node is not initialized.
MCAPI_ERR_CHAN_INVALID	Argument is not a channel handle.
MCAPI_ERR_CHAN_TYPE	Attempt to close a packet channel on an endpoint that has been connected with a different channel type.
MCAPI_ERR_CHAN_DIRECTION	Attempt to close a send handle on a port that was connected as a receiver, or vice versa.
MCAPI_ERR_CHAN_NOTOPEN	The channel is not open.
MCAPI_ERR_REQUEST_LIMIT	No more request handles available.

MCAPI_ERR_PARAMETER

Incorrect request parameter.

MCAPI_ERR_GENERAL

Implementation specific error not covered by other status codes. Specifics must be documented.

NOTE

Use the `mcapi_test()`, `mcapi_wait()` and `mcapi_wait_any()` functions to query the status of and `mcapi_cancel()` function to cancel the operation.

4.5.14 MCAPI_SCLCHAN_SEND_CLOSE_I

NAME

mcapi_sclchan_send_close_i – closes channel on a send endpoint.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_sclchan_send_close_i(
    MCAPI_IN mcapi_sclchan_send_hdl_t send_handle,
    MCAPI_OUT mcapi_request_t* request,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Closes the send side of a channel. The corresponding calls are required on both send and receive sides to ensure that the channel is properly closed. It is a non-blocking function, and returns immediately. send_handle is the local representation of the handle used to send packets. Pending scalars at the receiver are not discarded. A scalar channel is disconnected when, both sides have issued close calls and the last (second) close operation is performed. If the endpoint type was changed from message to channel and/or the get reference count was increased to 1 by the connect, the type will be reset to message and the get reference count to 0.

RETURN VALUE

On success, meaning that both sides of the channel are successfully closed, *mcapi_status is set to MCAPI_SUCCESS if completed and MCAPI_PENDING if not yet completed. On error *mcapi_status is set to the appropriate error defined below.

ERRORS

MCAPI_ERR_NODE_NOTINIT	The node is not initialized.
MCAPI_ERR_CHAN_INVALID	Argument is not a channel handle.
MCAPI_ERR_CHAN_TYPE	Attempt to close a packet channel on an endpoint that has been connected with a different channel type.
MCAPI_ERR_CHAN_DIRECTION	Attempt to close a send handle on a port that was connected as a receiver, or vice versa.
MCAPI_ERR_CHAN_NOTOPEN	The channel is not open.
MCAPI_ERR_REQUEST_LIMIT	No more request handles available.

MCAPI_ERR_PARAMETER

Incorrect request parameter.

MCAPI_ERR_GENERAL

Implementation specific error not covered by other status codes. Specifics must be documented.

NOTE

Use the `mcapi_test()`, `mcapi_wait()` and `mcapi_wait_any()` functions to query the status of and `mcapi_cancel()` function to cancel the operation.

4.6 MCAPi non-blocking operations

The connectionless message and packet channel API functions have both blocking and non-blocking variants. The non-blocking versions fill in a `mcapi_request_t` or `mca_request_t` object and return control to the user before the communication operation is completed. The `mcapi_test()`, `mcapi_wait()`, `mcapi_wait_any()` and `mcapi_cancel()` functions are used to query the status of the non-blocking operation.

4.6.1 MCAPI_TEST

NAME

`mcapi_test` – Tests if a non-blocking operation has completed.

SYNOPSIS

```
#include <mcapi.h>

mcapi_boolean_t mcapi_test(
    MCAPI_IN mcapi_request_t* request,
    MCAPI_OUT size_t* size,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Checks if a non-blocking operation has completed. The function returns in a timely fashion. `request` is the identifier for the non-blocking operation. The call only checks the completion of an operation and doesn't affect any messages/packets/scalars and does not release the request. If the specified request completes and the pending operation was a send or receive operation, the `size` parameter is set to the number of bytes that were either sent or received by the non-blocking transaction. The request is released with either of `mcapi_wait()`, `mcapi_wait_any()`, or `mcapi_cancel()`.

RETURN VALUE

On success, `MCAPI_TRUE` is returned and `*mcapi_status` is set to `MCAPI_SUCCESS`. If the operation has not completed `MCAPI_FALSE` is returned and `*mcapi_status` is set to `MCAPI_PENDING`. On error `MCAPI_FALSE` is returned and `*mcapi_status` is set to the appropriate error defined below or an error code for the corresponding non-blocking routine associated with the request structure.

ERRORS

<code>MCAPI_ERR_NODE_NOTINIT</code>	The node is not initialized.
<code>MCAPI_ERR_REQUEST_INVALID</code>	Argument is not a valid request handle. This also applies if the request has been canceled.
<code>MCAPI_ERR_PARAMETER</code>	Incorrect size parameter.
<code>MCAPI_ERR_GENERAL</code>	Implementation specific error not covered by other status codes. Specifics must be documented.

SEE ALSO

```
mcapi_endpoint_t mcapi_endpoint_getmcapi_endpoint_get_i(),
mcapi_msg_send_i(), mcapi_msg_rcv_i(), mcapi_pktchan_connect_i(),
```

```
mcapi_pktchan_rcv_open_i(), mcapi_pktchan_send_open_i(),  
mcapi_pktchan_send_i(), mcapi_pktchan_rcv_i(),  
mcapi_pktchan_rcv_close_i(), mcapi_pktchan_send_close_i(),  
mcapi_sclchan_connect_i(), mcapi_sclchan_rcv_open_i(),  
mcapi_sclchan_send_open_i(), mcapi_sclchan_rcv_close_i(),  
mcapi_sclchan_send_close_i()
```

NOTE

It is important to release the request with either of `mcapi_wait()`, `mcapi_wait_any()`, or `mcapi_cancel()`, in order to manage system resources.

4.6.2 MCAPI_WAIT

NAME

`mcapi_wait` – waits for a non-blocking operation to complete.

SYNOPSIS

```
#include <mcapi.h>

mcapi_boolean_t mcapi_wait(
    MCAPI_IN mcapi_request_t* request,
    MCAPI_OUT size_t* size,
    MCAPI_IN mcapi_timeout_t timeout,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Wait until a non-blocking operation has completed. It is a blocking function and returns when the operation has completed, has been canceled, or a timeout has occurred. `request` is the identifier for the non-blocking operation. The call only waits for the completion of an operation (all buffers referenced in the operation have been filled or consumed and can now be safely accessed by the application) and doesn't affect any messages/packets/scalars. The `size` parameter is set to the number of bytes that were either sent or received by the non-blocking transaction that completed (size is irrelevant for non-blocking connect and close calls). The `mcapi_wait()` call will return if the request is cancelled by a call to `mcapi_cancel()`, and the returned `mcapi_status` will indicate that the request was cancelled. The units for `timeout` are implementation defined. If a timeout occurs the returned status will indicate that the timeout occurred. A value of 0 for the `timeout` parameter indicates that the function will return without blocking with failure or success and a value of `MCAPI_INFINITE` for the `timeout` parameter indicates no timeout is requested, i.e. the function will block until it is unblocked because of failure or success. In regards to timeouts `mcapi_wait` is non destructive, i.e. the request is not cleared in the case of a timeout and can be waited upon multiple times. Multiple threads attempting to wait on the same request will result in an error.

RETURN VALUE

On success, `MCAPI_TRUE` is returned and `*mcapi_status` is set to `MCAPI_SUCCESS`. On error or timeout `MCAPI_FALSE` is returned and `*mcapi_status` is set to the appropriate error/status defined below for parameter errors for the `mcapi_wait()` call or an error code for the corresponding non-blocking routine associated with the request structure.

ERRORS

<code>MCAPI_ERR_NODE_NOTINIT</code>	The node is not initialized.
<code>MCAPI_ERR_REQUEST_INVALID</code>	Argument is not a valid request handle.
<code>MCAPI_ERR_REQUEST_CANCELLED</code>	The request was canceled, by another thread (during the waiting).

MCAPI_ERR_WAIT_PENDING	A wait is pending.
MCAPI_TIMEOUT	The operation timed out.
MCAPI_ERR_PARAMETER	Incorrect request or size parameter.
MCAPI_ERR_GENERAL	Implementation specific error not covered by other status codes. Specifics must be documented.

SEE ALSO

```
mcapi_endpoint_t mcapi_endpoint_getmcapi_endpoint_get_i(),
mcapi_msg_send_i(), mcapi_msg_recv_i(), mcapi_pktchan_connect_i(),
mcapi_pktchan_recv_open_i(), mcapi_pktchan_send_open_i(),
mcapi_pktchan_send_i(), mcapi_pktchan_recv_i(),
mcapi_pktchan_recv_close_i(), mcapi_pktchan_send_close_i(),
mcapi_sclchan_connect_i(), mcapi_sclchan_recv_open_i(),
mcapi_sclchan_send_open_i(), mcapi_sclchan_recv_close_i(),
mcapi_sclchan_send_close_i()
```

4.6.3 MCAPI_WAIT_ANY

NAME

`mcapi_wait_any` – waits for any non-blocking operation in a list to complete.

SYNOPSIS

```
#include <mcapi.h>

mcapi_uint_t mcapi_wait_any(
    MCAPI_IN size_t number,
    MCAPI_IN mcapi_request_t* requests,
    MCAPI_OUT size_t* size,
    MCAPI_IN mcapi_timeout_t timeout,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Wait until any non-blocking operation of a list has completed. It is a blocking function and returns the index into the `requests` array (starting from 0) indicating which of any outstanding operations has completed (successfully or with an error). `number` is the number of requests in the array. `requests` is the array of `mcapi_request_t` identifiers for the non-blocking operations. The call only waits for the completion of an operation and doesn't affect any messages/packets/scalars. The `size` parameter is set to number of bytes that were either sent or received by the non-blocking transaction that completed (`size` is irrelevant for non-blocking connect and close calls). The `mcapi_wait_any()` call will return `MCAPI_RETURN_VALUE_INVALID` if any of the requests are cancelled by calls to `mcapi_cancel()` (during the waiting). 0 could be a valid index, so status has to be checked to ensure correctness. The returned status will indicate that a request was cancelled. The units for `timeout` are implementation defined. If a timeout occurs the `mcapi_status` parameter will indicate that a timeout occurred. A value of 0 for the `timeout` parameter indicates that the function will return without blocking with failure or success and a value of `MCAPI_INFINITE` for the `timeout` parameter indicates no timeout is requested, i.e. the function will block until it is unblocked because of failure or success. In regards to timeouts `mcapi_wait_any()` is non destructive, i.e. the request is not cleared in the case of a timeout and can be waited upon multiple times. Multiple threads attempting to wait on the same request will result in an error.

RETURN VALUE

On success, the index into the `requests` array of the `mcapi_request_t` identifier that has completed or has been canceled is returned and `*mcapi_status` is set to `MCAPI_SUCCESS`. On error or timeout `MCAPI_RETURN_VALUE_INVALID` is returned and `*mcapi_status` is set to the appropriate error/status defined below for parameter errors for the `mcapi_wait_any()` call or errors from the requesting functions, as defined in those respective functions.

ERRORS

<code>MCAPI_ERR_NODE_NOTINIT</code>	The node is not initialized.
-------------------------------------	------------------------------

MCAPI_ERR_REQUEST_INVALID	Argument is not a valid request handle.
MCAPI_ERR_REQUEST_CANCELLED	One of the requests was canceled, by another thread (during the waiting).
MCAPI_ERR_WAIT_PENDING	A wait is pending.
MCAPI_TIMEOUT	The operation timed out.
MCAPI_ERR_PARAMETER	Incorrect number (if = 0), requests or size parameter.
MCAPI_ERR_GENERAL	Implementation specific error not covered by other status codes. Specifics must be documented.

SEE ALSO

```
mcapi_endpoint_t mcapi_endpoint_getmcapi_endpoint_get_i(),
mcapi_msg_send_i(), mcapi_msg_recv_i(), mcapi_pktchan_connect_i(),
mcapi_pktchan_recv_open_i(), mcapi_pktchan_send_open_i(),
mcapi_pktchan_send_i(), mcapi_pktchan_recv_i(),
mcapi_pktchan_recv_close_i(), mcapi_pktchan_send_close_i(),
mcapi_sclchan_connect_i(), mcapi_sclchan_recv_open_i(),
mcapi_sclchan_send_open_i(), mcapi_sclchan_recv_close_i(),
mcapi_sclchan_send_close_i()
```

4.6.4 MCAPI_CANCEL

NAME

`mcapi_cancel` – cancels an outstanding non-blocking operation.

SYNOPSIS

```
#include <mcapi.h>

void mcapi_cancel(
    MCAPI_IN mcapi_request_t* request,
    MCAPI_OUT mcapi_status_t* mcapi_status
);
```

DESCRIPTION

Cancels an outstanding non-blocking operation. It is a blocking function and returns when the operation has been canceled. `request` is the identifier for the non-blocking operation. Any pending calls to `mcapi_wait()` or `mcapi_wait_any()` for this request will also be cancelled. The returned status of a canceled `mcapi_wait()` or `mcapi_wait_any()` call will indicate that the request was cancelled.

RETURN VALUE

On success, `*mcapi_status` is set to `MCAPI_SUCCESS`. On error `*mcapi_status` is set to the appropriate error defined below.

ERRORS

<code>MCAPI_ERR_NODE_NOTINIT</code>	The node is not initialized.
<code>MCAPI_ERR_REQUEST_INVALID</code>	Argument is not a valid request handle (the operation and <code>mcapi_wait()</code> or <code>mcapi_wait_any()</code> may have completed).
<code>MCAPI_ERR_PARAMETER</code>	Incorrect <code>request</code> parameter.
<code>MCAPI_ERR_GENERAL</code>	Implementation specific error not covered by other status codes. Specifics must be documented.

SEE ALSO

```
mcapi_endpoint_t mcapi_endpoint_getmcapi_endpoint_get_i(),
mcapi_msg_send_i(), mcapi_msg_recv_i(), mcapi_pktchan_connect_i(),
mcapi_pktchan_recv_open_i(), mcapi_pktchan_send_open_i(),
mcapi_pktchan_send_i(), mcapi_pktchan_recv_i(),
mcapi_pktchan_recv_close_i(), mcapi_pktchan_send_close_i(),
mcapi_sclchan_connect_i(), mcapi_sclchan_recv_open_i(),
mcapi_sclchan_send_open_i(), mcapi_sclchan_recv_close_i(),
mcapi_sclchan_send_close_i()
```

4.7 MCAP_I support functions

Support functions are included to simplify the use of MCAP_I.

4.7.1 MCAPI_DISPLAY_STATUS

NAME

`mcapi_display_status` – Displays a status message in text format.

SYNOPSIS

```
#include <mcapi.h>

char* mcapi_display_status(
    MCAPI_IN mcapi_status_t mcapi_status,
    MCAPI_OUT char* status_message,
    MCAPI_IN size_t size
);
```

DESCRIPTION

Display an MCAPI status message in text format. It is a blocking function and returns in a timely fashion. The `status_message` parameter should point to an application provided char buffer of size `size` (recommended value for `size` = `MCAPI_MAX_STATUS_MSG_LEN`). `mcapi_status` is the status code. The status message string is NULL terminated.

RETURN VALUE

The status code in text format is returned. If the status code is invalid a NULL string is returned.

ERRORS

None.

SEE ALSO

5 Header files

mca.h - includes the MCA common definitions, mcapi.h - includes the MCAPI common definitions, mca_impl_spec.h- includes required implementation specific definitions, mcapi_impl_spec.h- includes required implementation specific definitions, mcapi_v1000_to_v2000.h - includes definitions to ease going from version 1.0 to 2.000, mcapi_1000_initialize.c- includes a 1.0 to 2.000 conversion function.

HEADER FILES ARE ATTACHED AS SEPARATE DOCUMENTS

6 FAQ

Q: Is a reference implementation available? What is the intended purpose of the reference implementation?

A: An example implementation is available for download at the Multicore Association website. The example implementation models the functionality of the specification and does not intend to be a high performance implementation.

Q: Does MCAPi provide binary interoperability?

A: MCAPi provides a set of communications primitives that supports software source code portability. Interoperability is not addressed at this stage and is left to the respective MCAPi implementations. Implementers should however be encouraged to ensure implementation interoperability, as much as possible. We may consider an interoperable messaging protocol in future MCAPi versions.

Q: What is the rationale behind MCAPi's use of unidirectional channels as opposed to bidirectional channels?

A: MCAPi is meant to serve as an extremely lightweight and thin communications API for use by applications and by other messaging systems (such as a light-weight sockets) layered on top of it. MCAPi is expected to run on both multicore processors running real operating systems and running extremely thin run-time environments in bare-metal mode. MCAPi is also meant largely for on-chip multicore environments where communications is reliable.

MCAPi chose to go with unidirectional channels because they use less resources on both the send and receive sides. A bidirectional channel would need to consume twice the resources on both the send and receive sides (for example, hardware queues).

Many of our use cases require only unidirectional channels. Although unidirectional channels could be implemented using a bidirectional channel, it would waste half the resources. Conversely, bidirectional channels can be implemented on top of unidirectional channels when desired. Locking mechanisms on the send and receive sides provided by the underlying system on top of which MCAPi is built (or provided by a layer such as MRAPi) can be further used to achieve atomicity between sends and receives on each side.

Bidirectional channels are particularly effective in error-prone environments. MCAPi's target is multicore processors, where the communications is expected to be reliable.

Q: How does MCAPi compare to other existing communication protocols for multicore?

A: MCAPi is a low level communication protocol specifically targeting statically configured heterogeneous multicore processors. The static property allows implementations to have lower overhead than other communication protocols such as sockets or MPI.

Q: Regarding message sends, when can the sender reuse the buffer?

A: With blocking calls, the send can reuse once the call has returned. With non-blocking calls, it is the responsibility of the sender to verify using API calls that the buffer can be reused.

Q: Can MCAPi support a software-controlled implementation of global power management? If so, how?

A: MCAPi provides communication functionality between different cores in a multicore multiprocessor and can function at user level and system level. If the global power management scheme requires data to be passed between different cores, an MCAPi implementation could be used to provide the communications functionality.

Q: Will MCAPi require modification of existing operating systems?

A: MCAPi is a communication API and does not define requirements for implementation. It is possible to implement MCAPi on top of a commercial off-the-shelf OS with no OS modifications, however an efficient implementation of MCAPi may require changes to the operating system.

Q: Can MCAPi be implemented in hardware?

A: The goal of MCAPi is to make possible efficient implementation in hardware.

Q: If MCAPi is oriented towards static connection topologies, why support methods like `mcapi_endpoint_delete()`?

A: The ability to release any resources allocated by an endpoint may be particularly important in hardware implementations where resources are strictly limited. Supporting the deletion of endpoints allows the application programmer to dispose of endpoints that are only necessary for a portion of the application's life.

Q: What facilities are provided for debugging MCAPi programs?

A: The MCAPi API is designed to be implemented as a C library, using standard tools and enabling the use of standard debuggers and profilers. The MCAPi error codes should also help give some visibility into what sort of program errors have occurred. The MCAPi channel-connection mechanism is designed to work with static layout tools, so that a channel topology can be defined and debugged outside of the program itself. We also expect that some implementations will provide tools for dynamically collecting information like which channels have been connected and how many packets have been transferred.

Q: Why doesn't the MCAPi API support one to multiple, multiple to one and barrier type communications?

A: To keep the MCAPi API simple and allow for efficient implementations more complex functions that can be layered on top of MCAPi are not part of the MCAPi specification.

Q: Why doesn't MCAPi provide name service functionality?

A: The MCAPi API is meant for multicore chips or multi-chip boards which are static environments in the sense that the number of cores is always the same. A naming service is therefore not needed and would add unnecessary complexity. The initialization functionality added in version 2.000 provides for basic topology discovery. A naming service could be implemented on top of MCAPi.

Q: What happens if my cores have different byte order (endian-ness)?

A: The MCAPi API specification doesn't state how to implement an MCAPi API. The specific implementation would have to address different byte order.

Q: How do I use my MCAPi calls so that the source code is portable between implementations with different node numbering policies?

A: To make code portable between implementations that allow only one node number (and `mcapi_initialize` call) per process, and those that allow one node number (and `mcapi_initialize` call) per thread, the code should only make one `mcapi_initialize` call per process. If there are multiple threads within the process, they should use different port numbers (endpoints) for communication.

Similarly, to make code portable between implementations that allow only one node number (and `mcapi_init` call) per processor, and those that allow one node number (and `mcapi_init` call) per task, the code should only make one `mcapi_init` call per processor. If there are multiple tasks running on the processor, they should use different port numbers (endpoints) for communication.

Q: Does the MCAP API spec provide for multiple readers or multiple writers on an endpoint?

A: No the spec does not specifically provide for this. However, it does not prevent it either. An MCAP API implementation can choose to provide a mechanism whereby multiple endpoints resolve to essentially the same destination. For example, if an SoC contained a hardware device which managed multiple hardware queues used for moving messages around the SoC, an MCAP API implementation could allow multiple writers to a single hardware queue by simply mapping a set of distinct send endpoints to that single hardware queue. The implementation might choose to do this by assigning a range of `port_ids` to map to the single hardware device. For example, assume the MCAP API implementation identifies the hardware device that provided the hardware managed queues as `node_id` 10, and the range of ports 0..16 is assigned to map to hardware queue zero within that device. Then a send to endpoint `<0,10,0>` as well as a send to endpoint `<0,10,1>` would both map to a send to queue zero in the hardware device. Different nodes within the system wishing to send to the queue would have to allocate a unique send endpoint, but the sends would all end up in queue zero in the hardware accelerator. Similarly, an MCAP API implementation could map multiple receive endpoints to a single hardware queue. In this example any reads from multiple endpoints mapped to the single hardware queue would be serviced on a first come, first served basis. This would be suitable for a load balancing application, but would not be sufficient to serve as a multicast operation. Support for multicast operations will be considered in future MCAP API specification releases.

Q: How can I implement callback capability with MCAP API?

A: Although MCAP API does not directly provide a callback capability, MCAP API is designed to have many kinds of application services layered on top of it, and callbacks can be done this way. This approach would require you to use the MCAP API non-blocking send or receive functions, and to write a new function which would take the `mcapi_request_t` returned from calling these MCAP API functions along with a pointer to your callback function as parameters. This new function could use `mcapi_test`, `mcapi_wait`, or `mcapi_wait_any` to determine when the outstanding MCAP API request has completed and then call your callback function. This might be implemented by having a separate thread monitoring a list of outstanding requests and providing callbacks on completion. If you cannot create multithreaded applications your callback solution may be more difficult to implement.

Q: I'd like to have a name service for looking up MCAP API endpoints. How can I achieve this functionality?

The MCAP API standard defines a communication layer that allows easy creation of a name server. The application can choose a statically determined (domain, node, port) tuple where the name service will run. For example, domain, node 0, port 0 might be a good choice. The name server process can be started by calling `mcapi_initialize(0,0, ...)` to bind itself to domain 0, node 0 and then calling `mcapi_endpoint_create(0, &status)`. The name server can then use message send and receive operations to implement its service. Clients can discover the name server by invoking `mcapi_endpoint_get(0, 0, 0, &status)` to get the its `endpoint_t`. Given the `endpoint_t`, each client can send and receive messages as required by your name service protocol.

7 Use Cases

7.1 Example Usage of Static Naming for Initialization

MCAPI's static tuple based naming mechanism makes it straightforward to implement a simple initialization scheme, including third party set up of static connections. This example describes how a packet channel can be created and used using the static tuple based naming scheme.

```
#define SENDER_DOMAIN 0
#define SENDER_NODE 0
#define SEND_PORT_ID 17

#define RECEIVER_DOMAIN 0
#define RECEIVER_NODE 1
#define RECV_PORT_ID 37
```

Sender Process:

```
mcapi_endpoint_t send_endpoint = mcapi_endpoint_create(SEND_PORT_ID,
&status);
```

Receiver Process:

```
mcapi_endpoint_t receive_endpoint = mcapi_endpoint_create(RECV_PORT_ID,
&status);
```

The connection can now be established by the receiver, the sender, or a third party process. We will use the example of a third party process.

Third party process:

```
mcapi_endpoint_t send_endpoint = mcapi_endpoint_get(SENDER_DOMAIN,
SENDER_NODE, SEND_PORT_ID, &status);
mcapi_endpoint_t receive_endpoint = mcapi_endpoint_get(RECEIVER_DOMAIN,
RECEIVER_NODE,
RECV_PORT_ID, &status);
mcapi_pktchan_connect_i(send_endpoint, receive_endpoint, ...)
```

Thus, the performance impact of a global name-service lookup is entirely in the `mcapi_endpoint_get()` call.

7.2 Example Initialization (Discovery and Bootstrapping) of Dynamic Endpoints

This example describes how MCAPI performs discovery and bootstrapping when the static naming based on tuples is not being used, rather the dynamic endpoint scheme is used.

Whether communication is by messages or by channels, MCAP requires a sending node to have available an endpoint on the receiving node in order to communicate. Thus there is the following discovery and bootstrapping issue in MCAP when static naming is not in use (and in any dynamic communication API for that matter): How to transfer the first endpoint from a receiver to a sender? This section proposes a discovery model, which addresses the issue of how an MCAP sender can bootstrap itself by finding the endpoint for a receiver, essentially before any user-level communication has been established.

The basic idea is to allow the MCAP runtime system to facilitate the creation of a root endpoint that is visible to all the nodes in the system. An endpoint can “publish” itself to the communications layer as the “root” or first-level name server of the namespace using statically known domain and node numbers. For example, the system could have exactly one statically numbered node, domain 0, node 0. All other nodes in the system would then register with node 0 via asynchronous messaging, sending it messages to let it know that they exist and what endpoints they have dynamically created. Similarly, dynamic nodes and endpoints can discover each other by sending messages to the ‘root node’. The format of the message is user defined.

7.3 Automotive Use Case

7.3.1 Characteristics

7.3.1.1 Sensors

Tens to hundreds of sensor inputs read on periodic basis. Each sensor is read and its data is processed by a scheduled task.

7.3.1.2 Control Task

A control task takes sensor values and computes values to apply to various actuators in the engine.

7.3.1.3 Lost Data

Lost data is not desirable, but old data quickly becomes irrelevant, most recent sample is most important.

7.3.1.4 Types of Tasks

Consists of both control and signal processing, especially FFT.

7.3.1.5 Load Balance

The load balance changes as engine speed increases. The frequency at which the control task must be run is determined by the RPM of the engine.

7.3.1.6 Message Size and Frequency

Messages are expected to be small and message frequency is high.

7.3.1.7 Synchronization

Synchronization between control and data tasks should be minimal to avoid negative impacts on latency of the control task; if shared memory is used it will always be one task writing and the other reading, so synch is not required or desirable; deadlock will not occur but old data may be used if an update is not ready.

7.3.2 Key functionality requirements

7.3.2.1 Control Task

There must be a control task collecting all data and calculating updates; this task must update engine parameters continuously; Updates to engine parameters must occur when the engine crankshaft is at a particular angle, so the faster the engine is running, the more frequently this task must run.

7.3.2.2 Angle Task

There must be a data task to monitor engine RPM and schedule the control task.

7.3.2.3 Data Tasks

There must be a set of tens to hundreds of task to poll sensors; the task must communicate this data to the control task.

7.3.3 Context/Constraints

7.3.3.1 Operating System

Often there is no commercial operating system involved, although the notion of time critical tasks and task scheduling must be supported by some type of executive; however this may be changing; possible candidates are OSEK, or other RTOS.

7.3.3.2 Polling/Interrupts

Sensor inputs may be polled and/or associated with interrupts.

7.3.3.3 Reliability

Sensors assumed to be reliable; interconnect assumed to be reliable; task completion within scheduled deadline is assumed to be reliable for the control task, and less reliable for the data tasks.

7.3.4 Metrics

7.3.4.1 Latency of the control task

Dependent on engine RPM. At 1800 RPM the task must complete every 33.33ms, and at 9000 RPM the task must complete every 6.667ms.

7.3.4.2 # dropped sensor readings

Ideally zero.

7.3.4.3 Latencies of data tasks

Ideally the sum of the latencies plus message send/receive times should be < latency of the control loop given current engine RPM. In general, individual tasks are expected to complete in times varying from 1ms up to 1600ms, depending on the nature of the sensor and the type of processing required for its data.

7.3.4.4 Code size

Automotive customers expect their code to fit into on chip SRAM. Current generation of chips often has 1Mb of SRAM, with 2Mb on the near horizon.

7.3.5 Possible Factorings

- 1 gp core for control, 1 gp core for data
- 1 gp core for control/data, dedicated SIMD core for signal processing, other SP cores for remainder of data processing
- 1 core per cylinder, or 1 core per group of cylinders

7.3.6 MCAPI Requirements Implications

- The control task must be able to determine if data is available from each data task and if not the task should be able to proceed in a non-blocking fashion using the last data from the sensor in question; this implies some form of non-blocking msg test or select mechanism [Automotive:7.3.2.1]
- It would be desirable for the control task to use MCAPI methods to send updates to actuators as 'messages', this implies some form of underlying efficient driver should be implemented; the send should be non-blocking [Automotive:7.3.2.1]

7.3.7 Mental Models

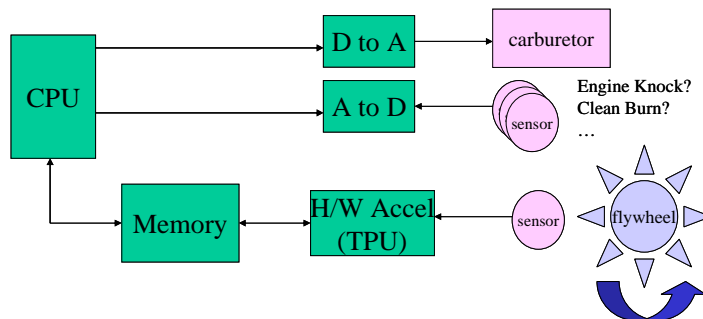


Figure 1 – example hardware

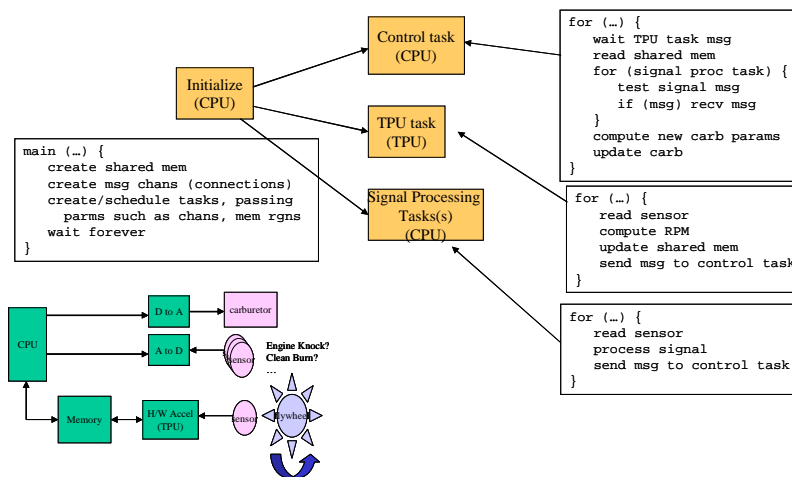


Figure 2 – A possible mapping

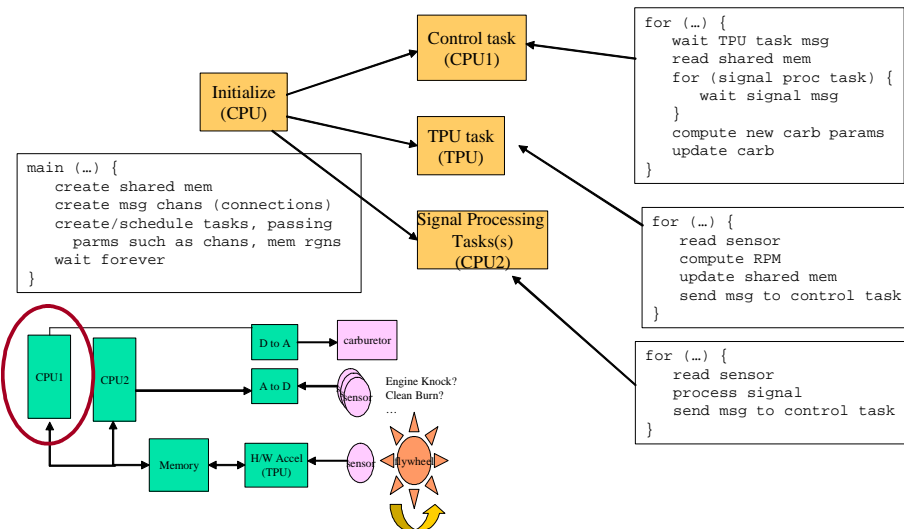


Figure 3 – Alternative Hardware

7.3.8 Applying the API draft to the pseudo-code

7.3.8.1 Initial Mapping

```
/*
 * MCAPI 2.000 Automotive Use Case
 * control_task.c
 *
 */

#include "mcapi.h"
#include "automotive.h"
extern char* shmget(int size);

////////////////////////////////////
// The control task
////////////////////////////////////
void Control_Task(void) {
    char* sMem;
    mcapi_uint8_t tFlag;
    mcapi_endpoint_t tpu_endpt, tpu_remote_endpt;
    mcapi_endpoint_t sig_endpt, sig_remote_endpt;
    mcapi_endpoint_t tmp_endpt;
    mcapi_sclchan_rcv_hdl_t tpu_chan;
    mcapi_pktchan_rcv_hdl_t sig_chan;
    SIG_DATA sDat;
    size_t tSize;
    mcapi_request_t r1, r2;
    mcapi_status_t err;
    mcapi_priority_t priority = PRIO_HIGH;
    mcapi_param_t mcapi_parameters;
    mcapi_info_t mcapi_info;

    // init the system
```

```

    mcapi_initialize(ENGINE_DOMAIN, CNTRL_NODE, &mcapi_parameters,
&mcapi_info, &err);
    CHECK_STATUS(err);

    // first create two local endpoints
    tpu_endpt = mcapi_endpoint_create(CNTRL_PORT_TPU, &err);
    CHECK_STATUS(err);

    sig_endpt = mcapi_endpoint_create(CNTRL_PORT_SIG, &err);
    CHECK_STATUS(err);

    // now we get two remote endpoints
    mcapi_endpoint_get_i(ENGINE_DOMAIN, TPU_NODE, TPU_PORT_CNTRL,
&tpu_remote_endpt, &r1, &err);
    CHECK_STATUS(err);

    mcapi_endpoint_get_i(ENGINE_DOMAIN, SIG_NODE, SIG_PORT_CNTRL,
&sig_remote_endpt, &r2, &err);
    CHECK_STATUS(err);

    // wait on the endpoints
    while(!((mcapi_test(&r1, &tSize, &err))) && (mcapi_test(&r2, &tSize,
&err))) {
        // KEEP WAITING
    }

    // allocate shared memory and send the ptr to TPU task
    sMem = shmemget(32);

    tmp_endpt = mcapi_endpoint_create(MCAPI_PORT_ANY, &err);
    CHECK_STATUS(err);

    mcapi_msg_send(tmp_endpt, tpu_remote_endpt, sMem, SHMEM_SIZE, priority,
&err);
    CHECK_STATUS(err);

    // connect the channels
    mcapi_sclchan_connect_i(tpu_endpt, tpu_remote_endpt, &r1, &err);
    CHECK_STATUS(err);

    mcapi_pktchan_connect_i(sig_endpt, sig_remote_endpt, &r2, &err);
    CHECK_STATUS(err);

    // wait on the connections
    while(!((mcapi_test(&r1, &tSize, &err))) && (mcapi_test(&r2, &tSize,
&err))) {
        // KEEP WAITING
    }

    // now open the channels
    mcapi_sclchan_rcv_open_i(&tpu_chan, tpu_endpt, &r1, &err);
    CHECK_STATUS(err);

    mcapi_pktchan_rcv_open_i(&sig_chan, sig_endpt, &r2, &err);
    CHECK_STATUS(err);

    // wait on the channels
    while(!((mcapi_test(&r1, &tSize, &err))) && (mcapi_test(&r2, &tSize,
&err))) {
        // KEEP WAITING
    }

```

```

// now ALL of the bootstrapping is finished
// we move to the processing phase below

while (1) {
    // wait for TPU to indicate it has updated
    // shared memory, indicated by receipt of a flag
    tFlag = mcapi_sclchan_rcv_uint8(tpu_chan, &err);
    CHECK_STATUS(err);

    // read the shared memory
    if (sMem[0] != 0) {
        // process the shared memory data
    } else {
        // PANIC -- there was an error with the shared mem
    }

    // now get data from the signal processing task
    // would be a loop if there were multiple sig tasks
    mcapi_pktchan_rcv(sig_chan, (void **) &sDat, &tSize, &err);
    CHECK_STATUS(err);

    // Compute new carb params & update carb
}
}

```

```

/*
 * MCAPI 2.000 Automotive Use Case
 * tpu_task.c
 *
 */

#include "mcapi.h"
#include "automotive.h"

////////////////////////////////////////
// The TPU task
////////////////////////////////////////
void TPU_Task() {
    char* sMem;
    size_t msgSize;
    size_t nSize;
    mcapi_endpoint_t cntrl_endpt, cntrl_remote_endpt;
    mcapi_sclchan_send_hdl_t cntrl_chan;
    mcapi_request_t rl;
    mcapi_status_t err;
    mcapi_timeout_t timeout = 500;
    mcapi_param_t mcapi_parameters;
    mcapi_info_t mcapi_info;

    // init the system
    mcapi_initialize(ENGINE_DOMAIN, TPU_NODE, &mcapi_parameters, &mcapi_info,
&err);
    CHECK_STATUS(err);

    cntrl_endpt = mcapi_endpoint_create(TPU_PORT_CNTRL, &err);
    CHECK_STATUS(err);

    mcapi_endpoint_get_i(ENGINE_DOMAIN, CNTRL_NODE, CNTRL_PORT_TPU,
&cntrl_remote_endpt, &rl, &err);
    CHECK_STATUS(err);

    // wait on the remote endpoint
    mcapi_wait(&rl, &nSize, &err, timeout);
    CHECK_STATUS(err);

    // now get the shared mem ptr
    mcapi_msg_rcv(cntrl_endpt, &sMem, SHMEM_SIZE, &msgSize, &err);
    CHECK_MEM(sMem);
    CHECK_STATUS(err);

    // NOTE - connection handled by control task
    // open the channel
    mcapi_sclchan_send_open_i(&cntrl_chan, cntrl_endpt, &rl, &err);
    CHECK_STATUS(err);

    // wait on the open
    mcapi_wait(&rl, NULL, &err, timeout);
    CHECK_STATUS(err);

    // ALL bootstrapping is finished, begin processing
    while (1) {
        // do something that updates shared mem
        sMem[0] = 1;

        // send a scalar flag to cntrl process

```



```

    // indicating sMem has been updated
    mcapi_sclchan_send_uint8(cntrl_chan, (mcapi_uint8_t) 1, &err);
    CHECK_STATUS(err);
}

}

/*
 * MCAPI 2.000 Automotive Use Case
 * sig_task.c
 *
 */

#include "mcapi.h"
#include "automotive.h"

////////////////////////////////////////
// The SIG Processing Task
////////////////////////////////////////
void SIG_task() {
    mcapi_endpoint_t cntrl_endpt, cntrl_remote_endpt;
    mcapi_pktchan_send_hdl_t cntrl_chan;
    mcapi_request_t r1;
    mcapi_status_t err;
    size_t size;
    mcapi_param_t mcapi_parameters;
    mcapi_info_t mcapi_info;
    mcapi_timeout_t timeout = 500;

    // init the system
    mcapi_initialize(ENGINE_DOMAIN, SIG_NODE, &mcapi_parameters, &mcapi_info,
&err);
    CHECK_STATUS(err);

    cntrl_endpt = mcapi_endpoint_create(SIG_PORT_CNTRL, &err);
    CHECK_STATUS(err);

    mcapi_endpoint_get_i(ENGINE_DOMAIN, CNTRL_NODE, CNTRL_PORT_SIG,
&cntrl_remote_endpt, &r1, &err);
    CHECK_STATUS(err);

    // wait on the remote endpoint
    mcapi_wait(&r1, &size, &err, timeout);
    CHECK_STATUS(err);

    // NOTE - connection handled by control task
    // open the channel
    mcapi_pktchan_send_open_i(&cntrl_chan, cntrl_endpt, &r1, &err);
    CHECK_STATUS(err);

    // wait on the open
    mcapi_wait(&r1, &size, &err, timeout);
    CHECK_STATUS(err);

    // All bootstrap is finished, now begin processing
    while (1) {
        // Read sensor & process signal
        SIG_DATA sDat; // populate this with results

        // send the data to the control process

```

```
    mcapi_pktchan_send(cntrl_chan, &sDat, sizeof(sDat), &err);  
    CHECK_STATUS(err);  
}  
}
```

7.3.8.2 Changes required to port to new multicore device

To map this code to additional CPUs, the only change required to this code is in the constant definitions for node and port numbers in the creation of endpoints.

7.3.8.3 Changes required to port to new multicore device

To map this code to additional CPUs, the only change required to this code is in the constant definitions for node and port numbers in the creation of endpoints.

7.4 Multimedia Processing Use Cases

Multimedia processing includes coding and decoding between various audio and video formats. Applications range from low-power mobile devices such as cell phones, with limited resolution and audio quality, to set-top-boxes and HDTV with extremely demanding performance requirements.

The following will review some of the use cases and communications characteristics for the multimedia application domain.

7.4.1 Characteristics

7.4.1.1 Simple Scenario

Figure 1 is a simple illustration of a multicore multimedia architecture. In this scenario, a multicore processor is executing a multimedia application, which is accelerated by a DSP integrated in the multicore device. The application has some data (for example, a video frame encoded in the MPEG4 format), and uses the DSP to decode into an image suitable for display. The data is moved (at least conceptually) to the DSP, the DSP processing runs, and the data is moved back to the general purpose processor (GPP).




Figure 3 - Simple Multimedia Scenario

Despite the simplicity of the above scenario, a number of characteristics are illustrated:

Heterogeneous Processors

The DSP and GPP have different instruction sets and potentially different data representations. The accelerators are typically 16, 24 or 32 bit special purpose devices. The accelerators may have limited code and data spaces. For example, total DSP code space for the application, operating system and communication infrastructure may be under 64K instruction words. In order to save room for applications, the communication infrastructure code footprint is ideally quite small (e.g., less than 1K VLIW instruction words). The communication data footprint is of the same magnitude.

There seems to be a trend towards 32 bit processing and less constrained acceleration devices in the future, although this is not altogether clear, in particular for low-power mobile devices.

In the example, the application is executing in user mode on a general purpose operating system such as Linux. Other operating systems such as QNX Neutrino, VxWorks, WinCE, etc may be used. The DSP is running potentially another standard operating system, or a “home grown” operating system, or perhaps in a “bare metal” environment, with no operating services to speak of.

Heterogeneous Communication Infrastructure

Figure 1 shows a network on chip connection between the GPP (in this example an ARM processor) and DSP. This may range from a simple bus to a well-designed network on chip infrastructure. However, (and perhaps more typically today), this could be an ad-hoc collection of buses and bridges, with different DMA engines and other mechanisms for moving data around the system.

Heterogeneous Application Environments

Multimedia applications are very complex, and have a number of frameworks such as gstreamer, DirectShow, MDF, OpenMax, etc. The communication infrastructure should “play well” with these environments.

For example, Figure 2 shows the OpenMax multimedia framework from the OpenMax specification³. Host-side communication APIs may already be defined by the framework. The MCAPI communication infrastructure is more likely suited to the tunneled communication between acceleration components.

³ See <http://www.khronos.org/openmax/>

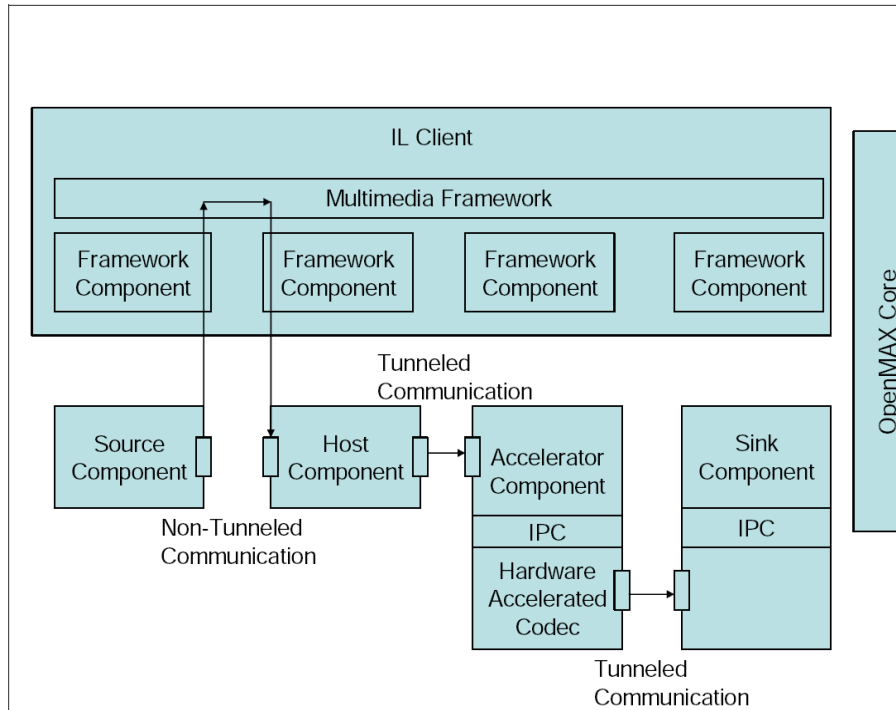


Figure 4 - OpenMax Communication

Application Non-Resident Code Size

Due to the large number of standards, encoding formats, processing variants, etc, the non-resident multimedia code size is quite large (millions of lines of code). This is only the code running on the accelerators, and does not include the code on the general purpose processors. Of course, depending on the operating mode, only a small subset of this code will be configured and running at a particular point in time.

Client/Server Communication Pattern

In simple configurations as illustrated in Figure 1, the communication follows a client/server pattern. However, as we will see shortly, this is not typically the case.

In the simple case, communication may be considered coarse grained, in that a large chunk of data is passed to the accelerator (for example, a video frame), and the DSP runs for a relatively long period of time to compute the results. Again, this coarse grained behavior is not always the case, as is described below.

Coarse Grain Data Rates and Latencies

For the coarse grained scenario, data rates and latencies are relatively conservative. Order of magnitude data sizes are 1K bytes, and rates under 100 Hz.

Potential for Zero Copy

Performance and power are key constraints. Unnecessary data movement and/or processing should be avoided where possible. Therefore, although the figure shows distinct data buffers on the GPP and DSP side, in practice the data may reside in a common shared memory. No data is actually copied, when processing control is transferred from the GPP to the DSP (and back again).

However, in many cases, the architecture does not feature a shared memory capability, in which case the data must be moved. Ideally, the communication infrastructure should mask these platform-dependant constraints, and allow zero copy where supportable by the architecture. This permits portable application code.

7.4.1.2 More Complex Scenarios

Figure 3 is a conceptual representation of a more complex scenario. This is motivated by applications such as H.264 encoding/decoding.

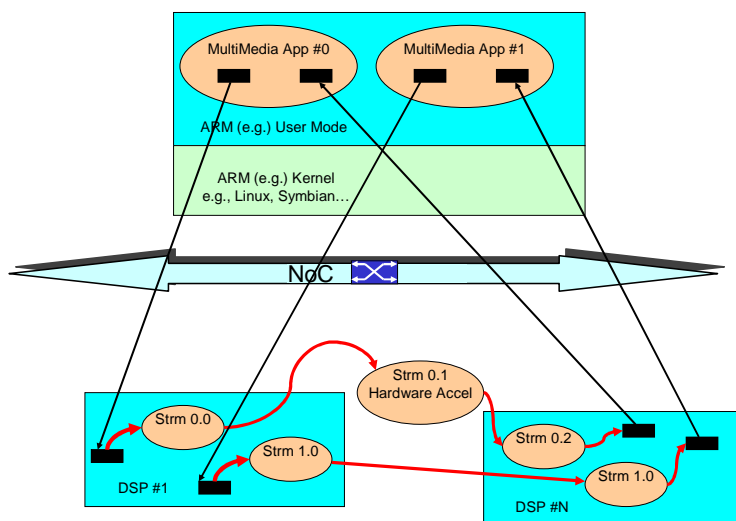


Figure 5 - More Complex Scenarios

High Processing Rates

The system computation requirements may exceed 100 GOPS/sec. The following characteristic may be viewed as a corollary:

Multi-Stage Processing

Due to extreme compute requirements, no one core can perform all of the processing – processing must be split into several stages. Currently, multimedia devices with order 10 cores are common.

Hardware Components

Even when split into multiple stages, computation for a particular stage may exceed the capabilities of a programmable device. In this case, the processing stage may be implemented in hardware. Therefore, the communication infrastructure must support communication to/from hardware devices.

Trend to finer-grained processing and communication

Several factors are influencing a trend to finer-grained processing and communication. One factor is the complexity, change, and customization options of the emerging standards. This requires great flexibility in the implementation, which in turn favors an approach which decomposes the problem into many parts, with the option of quickly changing the implementation of the parts. It seems that a natural fallout of this finer-grained decomposition is the requirement for finer-grained communication sizes and rates.

For example: a processor sends order 100 bytes to a hardware block for processing that takes order 100 cycles.

However, many different variants and approaches are in active investigation, the optimal strategy is not clear.

Streamed Processing

In the multi-stage configuration, a dataflow or streaming style of computation may be more appropriate. Depending on the granularity of the decomposition, the “tokens” exchanged between stages may be large (e.g., video frames), down to individual data values.

Multiple Flows

In many cases, multiple flows must be supported. For example, one or more audio streams, the main video stream, picture-in-picture streams, etc.

Dynamic Operation

Multimedia devices are increasingly dynamic. This includes many operating modes (voice calls, digital still camera, video reception and transmission, audio playback and record, etc). Arbitrary subsets of these operating modes may be selected for simultaneous operation by the user at fairly fine time-scales (e.g., in the order of seconds). Within a mode, operation may be dynamic at even finer time scales. For example, packets from a wireless connection may arrive in bursts.

Non-deterministic communication rates

In many cases, the amount of data consumed and produced by a processing stage may not be deterministic, resulting in variable communication rates.

Low Power Operation

Power is a central concern for mobile devices. Components must be switched off or moved to lower voltage/operating frequency depending on system load. Computation occasionally must be redistributed over the available resources in order to optimize power consumption.

Emergent System-Level Properties & Problems

Due to the complexity of the device and the unintended coupling of components, multimedia devices may exhibit difficult to understand “emergent” properties. For example, the system may deadlock, overflow buffers, miss deadlines, or provide “bursty” output, etc. Mechanisms to deal with these problems are very desirable.

Quality of Service

Device operation must meet quality of service constraints (processing latency, throughput, jitter, etc). Ideally, the communications infrastructure should support this. Due to high performance requirements, parts of this infrastructure may be in hardware.

Security

Multimedia flows have a strong security requirement. It may be required to prevent access to decoded data (for example), even if the operating system(s) have been compromised.

7.4.1.3 Metrics

The following table summarizes some of the quantitative characteristics of the multimedia application domain.

Attribute	Value (ranges)	Notes
Communication code size footprint on accelerators	8 K bytes	Becoming less constrained in the future
Communication data size footprint	8 K bytes	Becoming less constrained in the future
Non-resident code size	millions of lines of code	Increasing in the future
Number of processors	order 10	
Hardware accelerators	under 10	Different accelerators in high-end systems
Communication sizes	1K to under 100 bytes	Trend to finer-grain
Communication latency	10 mSec to under 1 uSec	Trend to finer-grain in future

7.4.2 Key Functionality Requirements

Table 1 summarizes the key requirements of multimedia devices and comments upon the ability MCAPI to meet the requirement.

Table 1 - Multimedia Derived Requirements

Description	Derived From	MCAPI
Low footprint on constrained devices (code and data size)	Characteristic 1	Yes
Framework for heterogeneous data types	Characteristic 1	Yes
Framework for complex transport configuration, quality of service, message delivery semantics and queuing policies (FIFO, most recent only, etc)	Characteristic 2	Some, others can be built on top.
Minimal assumptions about operating environment, and services required	Characteristic 3	Yes
No unnecessary data movement if architecture has supporting features (e.g., shared memory)	Characteristic 7	Yes, zero copy functionality
Allow direct binding to hardware, allowing end-to-end message delivery rates in the order of 100 cycles	Characteristic 8 Characteristic 9 Characteristic 10 Characteristic 11	Yes, no MCAPI feature prevents it
Framework for word-by-word message send/receive	Characteristic 8 Characteristic 9 Characteristic 10 Characteristic 11 Characteristic 12	Yes
Allow blocking and non-blocking operations	Characteristic 3	Yes
Provide message matching framework, to allow selection of messages of particular type, priority, deadline, etc	Characteristic 18	No, can be built on top of MCAPI.
Allow multi-drop messages. E.g., potential for >1 senders to queuing point, >1 readers	Characteristic 16 Characteristic 18	No, can be built on top of MCAPI.

7.4.3 Pseudo-Code Example:

The pseudo-code scenario contains the execution of a DSP algorithm such as a FIR filter on a multicore processor system containing a general purpose processor and DSP. The high level steps are as follows:

1. Initialize communication channels between processors.
2. GPP sends data for processing to DSP.
3. GPP sends code to execute to DSP
4. GPP signals for the DSP to begin execution.
5. GPP retrieves processed data from DSP

```

/*
 * MCAPI 2.000 Multimedia Use Case
 * shared.h
 *
 */

#ifndef SHARED_H_
#define SHARED_H_

#define CHECK_STATUS(status)
/* Shared Header File, shared.h */
enum {

```

```

        DSP_READY,
        DSP_DATA,
        DSP_CODE,
        DSP_TERMINATE,
        DSP_EXECUTE,
        COMPLETED
};

#define DOMAIN_0          0

#define PORT_COMMAND      0
#define PORT_DATA         1
#define PORT_DATA_RECV    2

#endif /* SHARED_H_ */

/*
 * MCAPI 2.000 Multimedia Use Case
 * gp_proc.c
 */

/* General Purpose Processor Code */
#include "shared.h"
#include "mcapi.h"

/* Predefined Node numbering */
#define GPP_1 0 /* 0-63 reserved for homogenous MC */
#define DSP_1 64 /* 64+ used for DSP nodes */

/* Predefined Port numbering */
#define PORT_COMMAND 0
#define PORT_DATA 1
#define PORT_DATA_REC 2

mcapi_endpoint_t command_endpoint;
mcapi_endpoint_t data_endpoint;
mcapi_endpoint_t remote_command_endpoint;
mcapi_endpoint_t remote_data_endpoint;
mcapi_endpoint_t data_rcv_endpoint;
mcapi_endpoint_t remote_data_rcv_endpoint ;

mcapi_pktchan_send_hdl_t data_chan;
mcapi_sclchan_send_hdl_t command_chan;
mcapi_pktchan_rcv_hdl_t data_rcv_chan;

void *data;
int data_size;
void *code;
int code_size;

void initialize_comms()
{
    mcapi_node_t gp_node = 0;
    mcapi_param_t mcapi_parameters;
    mcapi_info_t mcapi_info;
    mcapi_status_t status;
    mcapi_request_t request;

    mcapi_initialize(DOMAIN_0, gp_node, &mcapi_parameters, &mcapi_info,
&status);

```

```

CHECK_STATUS(status);

command_endpoint = mcapi_endpoint_create(PORT_COMMAND, & status);
CHECK_STATUS(status);

data_endpoint = mcapi_endpoint_create(PORT_DATA, & status);
CHECK_STATUS(status);

data_recv_endpoint = mcapi_endpoint_create(PORT_DATA_RECV, & status);
CHECK_STATUS(status);

remote_command_endpoint = mcapi_endpoint_get(DOMAIN_0, DSP_1,
PORT_COMMAND, &status);
CHECK_STATUS(status);

remote_data_endpoint = mcapi_endpoint_get(DOMAIN_0, DSP_1, PORT_DATA,
&status);
CHECK_STATUS(status);

remote_data_recv_endpoint = mcapi_endpoint_get(DOMAIN_0, DSP_1,
PORT_DATA_RECV,
&status);
CHECK_STATUS(status);

mcapi_pktchan_connect_i(data_endpoint, remote_data_endpoint, &request,
&status);
CHECK_STATUS(status);

mcapi_sclchan_connect_i(command_endpoint, remote_command_endpoint,
&request, &status)
CHECK_STATUS(status);

mcapi_pktchan_connect_i(data_recv_endpoint, remote_data_recv_endpoint,
&request, &status);
CHECK_STATUS(status);

mcapi_pktchan_send_open_i(&data_chan, data_endpoint, &request,
&status);
CHECK_STATUS(status);

mcapi_sclchan_send_open_i(&command_chan, command_endpoint, &request,
&status);
CHECK_STATUS(status);

mcapi_pktchan_recv_open_i(&data_recv_chan, data_recv_endpoint,
&request,
&status);
CHECK_STATUS(status);
}

void send_data(void *data, int size)
{
    mcapi_status_t status;
    mcapi_pktchan_send(data_chan, data, size, &status);
    CHECK_STATUS(status);
}

void send_dsp_cmd(int cmd)
{
    mcapi_status_t status;

```

```

        mcapi_sclchan_send_uint32(command_chan, cmd, &status);
        CHECK_STATUS(status);
    }

    void read_data(void **dst, int *size)
    {
        mcapi_status_t status;
        mcapi_pktchan_rcv(data_rcv_chan, dst, size, &status);
        CHECK_STATUS(status);
    }

    void shutdown_comms()
    {
        mcapi_status_t status;
        mcapi_request_t request;

        mcapi_pktchan_rcv_close_i(data_rcv_chan, &request, &status);
        CHECK_STATUS(status);
        mcapi_pktchan_send_close_i(data_chan, &request, &status);
        CHECK_STATUS(status);
        mcapi_sclchan_send_close_i(command_chan, &request, &status);
        CHECK_STATUS(status);

        mcapi_endpoint_delete(command_endpoint, &status);
        CHECK_STATUS(status);

        mcapi_endpoint_delete(data_endpoint, &status);
        CHECK_STATUS(status);

        mcapi_endpoint_delete(data_rcv_endpoint, &status);
        CHECK_STATUS(status);
    }

    void *allocate(int size) {
        /* Routine to allocate memory */
    }

    int perform_multimedia_function(void *code, int code_size, void *data, int
data_size)
    {
        void* result;
        int size = 0;
        initialize_comms();
        send_dsp_cmd(DSP_DATA);
        send_data(data, data_size);
        send_dsp_cmd(DSP_CODE);
        send_data(code, code_size);
        send_dsp_cmd(DSP_EXECUTE);
        read_data(&result, &size);
        send_dsp_cmd(DSP_TERMINATE);
        shutdown_comms();
        return COMPLETED;
    }

    /*
     * MCAPI 2.000 Multimedia Use Case
     * dsp.c
     *
     */

    /* DSP Code */

```

```

#include "shared.h"
#include "mcapi.h"

mcapi_endpoint_t dsp_command_endpoint;
mcapi_endpoint_t dsp_data_endpoint;
mcapi_endpoint_t dsp_data_send_endpoint;

mcapi_pktchan_rcv_hdl_t data_chan;
mcapi_sclchan_rcv_hdl_t command_chan;
mcapi_pktchan_send_hdl_t data_send_chan;

void *buffer;
void *code_buffer;
void *data_buffer;
int code_size;
int data_size;
void *result_buffer;
int result_size;

int dsp_initialize()
{
    mcapi_domain_t domain = 0;
    mcapi_node_t dsp_node = 0;
    mcapi_param_t mcapi_parameters;
    mcapi_info_t mcapi_info;
    mcapi_status_t status;

    mcapi_request_t request;

    mcapi_initialize(domain, dsp_node, &mcapi_parameters, &mcapi_info,
&status);

    CHECK_STATUS(status);

    dsp_command_endpoint = mcapi_endpoint_create(PORT_COMMAND, &status);
    CHECK_STATUS(status);

    dsp_data_endpoint = mcapi_endpoint_create(PORT_DATA, &status);
    CHECK_STATUS(status);
    dsp_data_send_endpoint = mcapi_endpoint_create(PORT_DATA_RECV,
&status);
    CHECK_STATUS(status);

    mcapi_pktchan_rcv_open_i(&data_chan, dsp_data_endpoint, &request,
&status);
    CHECK_STATUS(status);

    mcapi_sclchan_rcv_open_i(&command_chan, dsp_command_endpoint,
&request,
&status);
    CHECK_STATUS(status);

    mcapi_pktchan_send_open_i(&data_send_chan, dsp_data_send_endpoint,
&request,
&status);
    CHECK_STATUS(status);

    return status;
}

```

```

int dsp_shutdown()
{
    mcapi_status_t status;
    mcapi_request_t request;

    mcapi_pktchan_release(data_buffer, &status);

    mcapi_pktchan_release(code_buffer, &status);

    mcapi_pktchan_rcv_close_i(data_chan, &request, &status);
    CHECK_STATUS(status);

    mcapi_sclchan_rcv_close_i(command_chan, &request, &status);
    CHECK_STATUS(status);

    mcapi_pktchan_send_close_i(data_send_chan, &request, &status);
    CHECK_STATUS(status);

    mcapi_endpoint_delete(dsp_command_endpoint, &status);
    CHECK_STATUS(status);

    mcapi_endpoint_delete(dsp_data_endpoint, &status);
    CHECK_STATUS(status);

    mcapi_endpoint_delete(dsp_data_send_endpoint, &status);
    CHECK_STATUS(status);

    return status;
}

int receive_dsp_cmd()
{
    mcapi_uint32_t cmd;
    mcapi_status_t status;

    cmd = mcapi_sclchan_rcv_uint32(command_chan, &status);
    CHECK_STATUS(status);
}

int dsp_command_loop()
{
    int command = DSP_READY;
    mcapi_status_t status;

    dsp_initialize();

    while (command != DSP_TERMINATE) {
        switch (command) {
            case DSP_DATA:
                mcapi_pktchan_rcv(data_chan, &data_buffer, &data_size, &status);
                CHECK_STATUS(status);
                break;
            case DSP_CODE:
                mcapi_pktchan_rcv(data_chan, &code_buffer, &code_size, &status);
                CHECK_STATUS(status);
                /* Copy code to from buffer to local execute memory */
                break;
            case DSP_EXECUTE:
                /* Tell DSP to Execute - Assume code writes to result_buffer */

```

```

        mcapi_pktchan_send(data_send_chan, result_buffer, result_size,
                           &status);
        CHECK_STATUS(status);
        break;
    case DSP_TERMINATE:
        dsp_shutdown();
        break;
    default:
        break;
}
command = receive_dsp_cmd();
}
}

```

7.5 Packet Processing

This example presents the typical startup and inner loop of a packet processing application. There are two source files: `load_balancer.c` and `worker.c`. The main entrypoint is in `load_balancer.c`.

The program begins in the load balancer, which spawns a set of worker processes and binds channels to them. Each worker has two channel connections to the load balancer: a packet channel for work requests and a scalar channel for acks. When work arrives on the packet channel from the load balancer, the worker processes it and then sends back an ack to the load balancer. The load balancer will not send new work to a worker unless an “ack” word is available.

This example uses both packet channels and scalar channels. Scalar channels are used for acks from the workers to the load balancer, since each ack is a single word and performance will be better without the packetization overhead. Packet channels are used to send work requests to the workers, since each work request is a structure containing multiple fields. The example also shows how to use both statically named and dynamic (anonymous) endpoints; it uses the endpoint creation functions `mcapi_endpoint_create()`, `mcapi_create_anonymous_endpoint()`, and `mcapi_endpoint_get()`.

On architectures with hardware support for scalars it would make more sense to use scalar channels for the work requests as well. Work requests are usually composed of only a few words, so they can be reasonably sent as individual words. More importantly, the load balancer is the rate-limiting step in this program; we can always add more workers but they’re useless if the load balancer can’t feed them. Since the load balancer is limited by communication overhead, hardware accelerated scalar channels could potentially improve overall performance if the overhead of a packet channel send is greater than the comparable scalar channel sends.

7.5.1 Packet Processing Code

7.5.1.1 common.h

```
/*
 * MCAPI 2.000 Packet Processing Use Case
 * common.h
 *
 */

// Header file containing a couple of declarations that are shared
// between the load balancer and the worker.

#ifndef COMMON_H_
#define COMMON_H_

#include <stdio.h>
#include <stdlib.h>

extern char* mcapi_strerror(int s);

enum {
    WORKER_REQUEST_PORT_ID,
    WORKER_ACK_PORT_ID
};

#define DOMAIN_0 0
#define NODE_LOAD_BALANCER 0

#define CHECK_STATUS(S) do { \
    if (S != MCAPI_SUCCESS) { \
        printf("Failed with MCAPI status %d (%s)\n", S, mcapi_strerror(S)); \
        exit(1); \
    } \
} while (0)

typedef struct{
    unsigned char is_valid;
    int worker;
} PacketInfo;

#endif /* COMMON_H_ */
```

7.5.1.2 load_balancer.c

```
/*
 * MCAPI 2.000 Packet Processing Use Case
 * load_balancer.c
 *
 */

#include "common.h"
#include "mcapi.h"

#include <stdbool.h>

extern void spawn_new_thread(void* funct, int worker_id);
extern void worker_spawn_function(int worker_num);
extern int get_next_packet(PacketInfo* packet_info);
extern void drop_packet(PacketInfo* packet_info);
```



```

// We will have four worker processes
#define NUM_WORKERS 4

// An array of packet channel send ports for sending work descriptors
// to each of the worker processes.
mcapi_port_t work_requests_out_port[NUM_WORKERS];

// An array of scalar channel receive ports for getting acks back from
// the workers.
mcapi_port_t acks_in_port[NUM_WORKERS];

mcapi_pktchan_send_hndl_t work_requests_out_hndl[NUM_WORKERS];
mcapi_pktchan_recv_hndl_t acks_in_hndl[NUM_WORKERS];

// function declarations
void create_and_init_workers(void);
void dispatch_packets(void);
void shutdown_lb(void);

// The entrypoint for this packet processing application.
int main(void)
{
    mcapi_param_t mcapi_parameters;
    mcapi_info_t mcapi_info;
    mcapi_status_t status;

    mcapi_initialize(DOMAIN_0, NODE_LOAD_BALANCER, &mcapi_parameters,
&mcapi_info, &status);

    create_and_init_workers();
    dispatch_packets();
    shutdown_lb();

    return 0;
}

void create_and_init_workers()
{
    int i;
    mcapi_request_t request;

    for (i = 0; i < NUM_WORKERS; i++)
    {
        mcapi_status_t status;

        // Spawn a new thread; pass parameters so the new thread will execute
        // worker_spawn_function(bootstrap_endpoint)
        spawn_new_thread(&worker_spawn_function, i);
        int worker_node = i + 1;

        // Create a send endpoint to send packets to the worker; get the
        // worker's receive endpoint via mcapi_endpoint_get()
        mcapi_endpoint_t work_request_out_endpoint =
mcapi_endpoint_create(MCAPI_PORT_ANY, &status);
        CHECK_STATUS(status);
        mcapi_endpoint_t work_request_remote_endpoint =
mcapi_endpoint_get(DOMAIN_0, worker_node, WORKER_REQUEST_PORT_ID, &status);
        CHECK_STATUS(status);
    }
}

```

```

    // Bind the channel and open our local send port.
    mcapi_pktchan_connect_i(work_request_out_endpoint,
work_request_remote_endpoint, &request, &status);
    CHECK_STATUS(status);
    mcapi_pktchan_send_open_i(&work_requests_out_hndl[i],
work_request_out_endpoint, &request, &status);
    CHECK_STATUS(status);

    // Repeat the process to create an ack scalar channel from the
    // worker back to us.
    mcapi_endpoint_t ack_in_endpoint = mcapi_endpoint_create(MCAPI_PORT_ANY,
&status);
    CHECK_STATUS(status);
    mcapi_endpoint_t ack_remote_endpoint = mcapi_endpoint_get(DOMAIN_0,
worker_node, WORKER_ACK_PORT_ID, &status);
    CHECK_STATUS(status);

    mcapi_sclchan_connect_i(ack_remote_endpoint, ack_in_endpoint, &request,
&status);
    CHECK_STATUS(status);

    mcapi_sclchan_rcv_open_i(&acks_in_hndl[i], ack_in_endpoint, &request,
&status);
    CHECK_STATUS(status);
}
}

void dispatch_packets()
{
    PacketInfo packet_info;
    mcapi_status_t status;

    while(get_next_packet(&packet_info))
    {
        // Because we maintain "session state" across packets, each
        // incoming packet is associated with a particular worker.
        int worker = packet_info.worker;

        // Each worker sends back acks when it is ready for more work.
        if (mcapi_sclchan_available(acks_in_hndl[worker], &status))
        {
            // An ack is available; pull it off and send more work
            int unused_result = mcapi_sclchan_rcv_uint8(acks_in_hndl[worker],
&status);
            mcapi_pktchan_send(work_requests_out_hndl[worker], &packet_info,
sizeof(packet_info), &status);
            CHECK_STATUS(status);
        }
        else
        {
            // No ack available; drop the packet (or queue, or do some other
            // form of exception processing) and move on.
            drop_packet(&packet_info);
        }
    }
}

```

```

void shutdown_lb()
{
    mcapi_status_t status;
    mcapi_request_t request;
    int i;

    // Shutdown each worker in turn
    for(i = 0; i < NUM_WORKERS; i++)
    {
        // Send an "invalid" packet to trigger worker shutdown
        PacketInfo invalid_work;
        invalid_work.is_valid = false;
        mcapi_pktchan_send(work_requests_out_hdl[i], &invalid_work,
sizeof(invalid_work), &status);
        CHECK_STATUS(status);

        // Close our ports; don't worry about errors (what would we do?)
        mcapi_pktchan_send_close_i(work_requests_out_hdl[i], &request,
&status);
        mcapi_sclchan_rcv_close_i(acks_in_hdl[i], &request, &status);
    }
}

```

7.5.1.3 worker.c

```

/*
 * MCAPI 2.000 Packet Processing Use Case
 * worker.c
 *
 */

#include "common.h"
#include "mcapi.h"
#include <stdbool.h>

void some_function_or_another(PacketInfo* packet_info);

// Local port for incoming work requests from the load balancer. We
// use a packet channel because each work request is a structure with
// multiple fields specifying the work to be done.
mcapi_port_t work_request_in;

// Local port for outgoing acks to the load balancer. We use a scalar
// channel because each ack is a single word indicating the number of
// work items that have been completed.
mcapi_port_t ack_out;

// a local buffer for use by the incoming packet channel
char work_channel_buffer[1024];

// function declarations
void bind_channels(void);
void do_work(void);
void shutdown(void);

// The function called within each new worker thread. This function
// binds takes an endpoint in the load balancer as its parameter so
// that it can communicate with the load balancer and create channels
// between the worker and the load balancer.

```

```

void worker_spawn_function(int worker_num)
{
    mcapi_param_t mcapi_parameters;
    mcapi_info_t mcapi_info;
    mcapi_status_t status;

    mcapi_initialize(DOMAIN_0, (1 + worker_num), &mcapi_parameters,
&mcapi_info, &status);

    bind_channels();
    do_work();
    shutdown();
}

void bind_channels()
{
    mcapi_status_t status;
    mcapi_request_t request;
    mcapi_pktchan_rcv_hndl_t work_request_in;
    mcapi_sclchan_send_hndl_t ack_out;

    // Create a packet receive endpoint for incoming work requests; use
    // a static port number so that load balancer can look it up via
    // mcapi_endpoint_get().
    mcapi_endpoint_t work_request_in_endpoint =
mcapi_endpoint_create(WORKER_REQUEST_PORT_ID, &status);
    CHECK_STATUS(status);

    // The load balancer will bind a channel to that endpoint; we use
    // the open function to synchronize and initialize our local packet
    // receive port.
    mcapi_pktchan_rcv_open_i(&work_request_in, work_request_in_endpoint,
&request, &status);
    CHECK_STATUS(status);

    // Repeat the process to create a scalar channel from us back to the
    // load balancer.
    mcapi_endpoint_t ack_out_endpoint =
mcapi_endpoint_create(WORKER_ACK_PORT_ID, &status);
    CHECK_STATUS(status);

    mcapi_sclchan_send_open_i(&ack_out, ack_out_endpoint, &request, &status);
    CHECK_STATUS(status);
}

void do_work()
{
    mcapi_status_t status;
    size_t size;

    // The first thing we do is send an ack so that the load balancer
    // knows we're ready for work.
    mcapi_sclchan_send_uint8(ack_out, 1 /* Any value would do. */, &status);

    while (1)
    {
        // receive a work request from the load balancer
        PacketInfo* work_info;
        mcapi_pktchan_rcv(work_request_in, (void*)&work_info, &size, &status);
        CHECK_STATUS(status);
    }
}

```

```

    // if the work request is marked "invalid", we're done
    if (work_info->is_valid == false)
    {
        break;
    }

    // do the work
    some_function_or_another(work_info);

    // we're done; return the buffer and send an ack
    mcapi_pktchan_release(work_info, &status);
    mcapi_sclchan_send_uint8(ack_out, 1 /* Any value would do. */, &status);
}

void shutdown()
{
    mcapi_status_t status;
    mcapi_request_t request;

    // close our channels
    mcapi_pktchan_rcv_close_i(work_request_in, &request, &status);
    mcapi_sclchan_send_close_i(ack_out, &request, &status);
}

```

Acknowledgements

The MCAPI working group would like to acknowledge the significant contributions of the following people in the creation of this API specification:

Working Group

Ajay Kamalvanshi, Nokia Siemens Networks
Anant Agarwal, Tilera
Sven Brehmer, PolyCore Software
Max Domeika, Intel
Peter Flake, Imperas
Steven Furr, Freescale Semiconductor, Inc.
Masaki Gondo, eSOL
Patrick Griffin, Tilera
Jim Holt, Freescale Semiconductor, Inc.
Rob Jackson, Imperas
Kevin Kissell, MIPS
Maarten Koenig, Wind River
Markus Levy, Multicore Association
Charles Pilkington, ST Micro
Andrew Richards, Codeplay
Colin Riley, Codeplay
Frank Schirrmeister, Imperas

The MCAPI working group also would like to thank the external reviewers who provided input and helped us to improve the specification below is a partial list of the external reviewers (some preferred to not be mentioned).

Reviewers

David Addison, ST Micro
Sterling Augustine, Tensilica
Badrinath Dorairajan, Aricent
Marc Gauthier, Tensilica
David Heine, Tensilica
Dominic Herity, Redplain Technology
Arun Joseph, IBM
Tammy Leino, Mentor Graphics
Michael Kardonik, Freescale Semiconductor, Inc.
Anjna Khanna, Cadence Design Systems
Grant Martin, Tensilica
Tim Mattson, Intel
Dror Maydan, Tensilica
Girish Mundada, Micronas
Karl Mörner, Enea
Emeka Nwafor, Zeligsoft
Don Padgett, Padgett Computing
Ola Redell, Enea
Michele Reese, Freescale Semiconductor, Inc.
Greg Regnier, Intel
Vedvyas Shanbhogue, Intel
Patrik Strömblad, Enea